

Where Deep Learning Losses Come From

Likelihood, cross-entropy, and regularization from probability

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

Three familiar losses, one probabilistic origin

You have already met these losses:

Three familiar losses, one probabilistic origin

You have already met these losses:

$$\mathcal{L}_{\text{reg}} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2 \quad (\text{regression})$$

Three familiar losses, one probabilistic origin

You have already met these losses:

$$\mathcal{L}_{\text{reg}} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2 \quad (\text{regression})$$

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_i [y_i \log p_i + (1 - y_i) \log(1 - p_i)] \quad (\text{binary classification})$$

Three familiar losses, one probabilistic origin

You have already met these losses:

$$\mathcal{L}_{\text{reg}} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2 \quad (\text{regression})$$

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_i [y_i \log p_i + (1 - y_i) \log(1 - p_i)] \quad (\text{binary classification})$$

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{data}} + \lambda \|\boldsymbol{\theta}\|_2^2 \quad (\text{regularized})$$

Three familiar losses, one probabilistic origin

You have already met these losses:

$$\mathcal{L}_{\text{reg}} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2 \quad (\text{regression})$$

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_i [y_i \log p_i + (1 - y_i) \log(1 - p_i)] \quad (\text{binary classification})$$

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{data}} + \lambda \|\boldsymbol{\theta}\|_2^2 \quad (\text{regularized})$$

Objective. Derive each term from an explicit likelihood or prior.

Two probabilistic assumptions determine the two halves of the objective.

The labels below are signposts, not assumed vocabulary. We proceed in order: predictive distribution and likelihood;

i.i.d. data; log-likelihood and NLL; then the prior $p(\theta)$ and MAP.

Two probabilistic assumptions determine the two halves of the objective.

The labels below are signposts, not assumed vocabulary. We proceed in order: predictive distribution and likelihood;

i.i.d. data; log-likelihood and NLL; then the prior $p(\theta)$ and MAP.



Probability primer


The **support** $\mathcal{S}_Y$ is the set of values that Y can take.

Probability mass function (PMF)

$$p_{Y(y)} = P(Y = y), \quad y \in \mathcal{S}_Y$$

$$p_{Y(y)} \geq 0, \quad \sum_{y \in \mathcal{S}_Y} p_{Y(y)} = 1$$

Bernoulli(0.75) PMF



each bar is an outcome probability

Examples. Bernoulli: $\{0, 1\}$ · categorical: $\{1, \dots, K\}$ · Poisson: $\{0, 1, 2, \dots\}$

Discrete distributions assign probability to outcomes


The **support** $\mathcal{S}_Y$ is the set of values that Y can take.

Probability mass function (PMF)

$$p_{Y(y)} = P(Y = y), \quad y \in \mathcal{S}_Y$$

$$p_{Y(y)} \geq 0, \quad \sum_{y \in \mathcal{S}_Y} p_{Y(y)} = 1$$

Bernoulli(0.75) PMF



each bar is an outcome probability

Examples. Bernoulli: $\{0, 1\}$ · categorical: $\{1, \dots, K\}$ · Poisson: $\{0, 1, 2, \dots\}$

Exact values can have nonzero probability: here $P(Y = 1) = 0.75$.

Continuous distributions assign density across a support


The **support** may be bounded or unbounded:

$$\mathcal{S}_Y = [a, b]$$

for a Uniform variable; $\mathcal{S}_Y = \mathbb{R}$ for a Gaussian.

Probability density function (PDF)

$$p_{Y(y)} \geq 0, \quad \int_{\mathcal{S}_Y} p_{Y(y)} \, dy = 1$$



the total area under a PDF is 1

Continuous distributions assign density across a support


The **support** may be bounded or unbounded:

$$\mathcal{S}_Y = [a, b]$$

for a Uniform variable; $\mathcal{S}_Y = \mathbb{R}$ for a Gaussian.

Probability density function (PDF)

$$p_{Y(y)} \geq 0, \quad \int_{\mathcal{S}_Y} p_{Y(y)} \, dy = 1$$



the total area under a PDF is 1

For continuous Y , probability comes from **area under the density.**

For a continuous Y , probability comes from **area**:


$$P(a \leq Y \leq b) = \int_a^b p_{Y(y)} \, dy$$

Density is not probability

For a continuous Y , probability comes from **area**:

$$P(a \leq Y \leq b) = \int_a^b p_{Y(y)} dy$$


$p_{Y(y)} \neq P(Y = y)$ — a density can **exceed 1**; only *integrals* are probabilities.



the shaded area is $P(a \leq Y \leq b)$

$$Y \sim \mathcal{U}(a, b) \quad p_{Y(y)} = \begin{cases} 1/(b-a) & a \leq y \leq b \\ 0 & \text{otherwise} \end{cases}$$

The support matters: outside $[a, b]$, the density is exactly 0.




$$Y \sim \mathcal{N}(\mu, \sigma^2) \quad p(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$

μ is the mean; $\sigma > 0$ is the standard deviation (scale); σ^2 is the variance.

$$Y \sim \mathcal{N}(\mu, \sigma^2) \quad p(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$

μ is the mean; $\sigma > 0$ is the standard deviation (scale); σ^2 is the variance.

— $\mu = 0, \sigma = 1$ — $\mu = 0, \sigma = 1.8$ — $\mu = -1.5, \sigma = 1$



Only two knobs: **location** μ and **scale** σ .

From probability to likelihood

Read the model notation one symbol at a time

Q QUESTION

$$p_{\theta}(y \mid x)$$

Read the model notation one symbol at a time

$$p_{\theta}(y \mid x)$$

X / x	X is the random input vector; x is one observed input
Y / y	Y is the random target; y is one possible or observed value
θ	the parameter vector — the model's learned weights and biases
$p_{\theta}(y \mid x)$	the probability (discrete Y) or density (continuous Y) assigned to y , given x

Read the model notation one symbol at a time

$$p_{\theta}(y \mid x)$$

X / x	X is the random input vector; x is one observed input
Y / y	Y is the random target; y is one possible or observed value
θ	the parameter vector — the model's learned weights and biases
$p_{\theta}(y \mid x)$	the probability (discrete Y) or density (continuous Y) assigned to y , given x

Read aloud: “model θ ’s probability or density for y , given x .”

A point prediction summarizes a distribution


REGRESSION • $Y \in \mathbb{R}$

Point prediction: mean of $Y \mid \mathbf{X} = \mathbf{x}$

$$\hat{y}(\mathbf{x}) = \mathbb{E}_{\boldsymbol{\theta}}[Y \mid \mathbf{X} = \mathbf{x}] = \mu_{\boldsymbol{\theta}}(\mathbf{x})$$

Full predictive distribution

$$Y \mid \mathbf{X} = \mathbf{x} \\ \sim \mathcal{N}(\mu_{\boldsymbol{\theta}}(\mathbf{x}), \sigma_{\boldsymbol{\theta}}^2(\mathbf{x}))$$



A point prediction summarizes a distribution


REGRESSION • $Y \in \mathbb{R}$

Point prediction: mean of $Y \mid \mathbf{X} = \mathbf{x}$

$$\hat{y}(\mathbf{x}) = \mathbb{E}_{\theta}[Y \mid \mathbf{X} = \mathbf{x}] = \mu_{\theta}(\mathbf{x})$$

Full predictive distribution

$$Y \mid \mathbf{X} = \mathbf{x} \\ \sim \mathcal{N}(\mu_{\theta}(\mathbf{x}), \sigma_{\theta}^2(\mathbf{x}))$$



BINARY CLASSIFICATION • $Y \in \{0, 1\}$


Point prediction: most probable outcome (mode)

$$\hat{y}(\mathbf{x}) = \mathbf{1}[p_{\theta}(\mathbf{x}) \geq 0.5]$$

Full predictive distribution

$$Y \mid \mathbf{X} = \mathbf{x} \\ \sim \text{Bernoulli}(p_{\theta}(\mathbf{x}))$$

example: $p_{\theta}(\mathbf{x}) = 0.75$



A point prediction summarizes a distribution


REGRESSION • $Y \in \mathbb{R}$

Point prediction: mean of $Y \mid \mathbf{X} = x$

$$\hat{y}(x) = \mathbb{E}_{\theta}[Y \mid \mathbf{X} = x] = \mu_{\theta}(x)$$

Full predictive distribution

$$Y \mid \mathbf{X} = x \\ \sim \mathcal{N}(\mu_{\theta}(x), \sigma_{\theta}^2(x))$$



BINARY CLASSIFICATION • $Y \in \{0, 1\}$


Point prediction: most probable outcome (mode)

$$\hat{y}(x) = \mathbf{1}[p_{\theta}(x) \geq 0.5]$$

Full predictive distribution

$$Y \mid \mathbf{X} = x \\ \sim \text{Bernoulli}(p_{\theta}(x))$$

example: $p_{\theta}(x) = 0.75$



Neural network: $x \rightarrow$ distribution parameters. Prediction: $\hat{y} =$ mean or mode.

Prediction fixes the model and varies the outcome

Example: $Y \in \{0, 1\}$. Fix the input x^* and one parameter setting θ_A .

Example: $Y \in \{0, 1\}$. Fix the input x^* and one parameter setting θ_A .


$$p_{\theta_A}(Y = 0 \mid x^*) = 0.20$$

$$p_{\theta_A}(Y = 1 \mid x^*) = 0.80$$

$$0.20 + 0.80 = 1$$

We vary the possible value of Y while keeping θ_A and x^* fixed.

predictive distribution



Prediction fixes the model and varies the outcome

Example: $Y \in \{0, 1\}$. Fix the input x^* and one parameter setting θ_A .


$$p_{\theta_A}(Y = 0 \mid x^*) = 0.20$$

$$p_{\theta_A}(Y = 1 \mid x^*) = 0.80$$

$$0.20 + 0.80 = 1$$

We vary the possible value of Y while keeping θ_A and x^* fixed.

predictive distribution



$$\hat{y}(x^*) = \arg \max_y p_{\theta_A}(y \mid x^*) = 1$$

One observed label gives one likelihood score

D DERIVATION

Now observe $y^* = 1$ for the input x^* , so $\mathcal{D} = \{(x^*, 1)\}$.

$$L(\theta; \mathcal{D}) = p_{\theta}(y^* = 1 \mid x^*)$$

One observed label gives one likelihood score

Now observe $y^* = 1$ for the input x^* , so $\mathcal{D} = \{(x^*, 1)\}$.

$$L(\theta; \mathcal{D}) = p_{\theta}(y^* = 1 \mid x^*)$$


Candidate A: $L(\theta_A; \mathcal{D}) = 0.80$

Candidate B: $L(\theta_B; \mathcal{D}) = 0.30$

$$L(\theta_A; \mathcal{D}) > L(\theta_B; \mathcal{D})$$

θ_A makes the observed label $y^* = 1$ less surprising.

likelihood of the fixed observation



One observed label gives one likelihood score

Now observe $y^* = 1$ for the input x^* , so $\mathcal{D} = \{(x^*, 1)\}$.

$$L(\theta; \mathcal{D}) = p_{\theta}(y^* = 1 \mid x^*)$$


Candidate A: $L(\theta_A; \mathcal{D}) = 0.80$

Candidate B: $L(\theta_B; \mathcal{D}) = 0.30$

$$L(\theta_A; \mathcal{D}) > L(\theta_B; \mathcal{D})$$

θ_A makes the observed label $y^* = 1$ less surprising.

likelihood of the fixed observation



Likelihood is a **score as a function of θ** ; the scores need not sum to 1 over parameter settings.

One continuous observation gives a likelihood too

Suppose $Y \mid \mu \sim \mathcal{N}(\mu, 1)$ and we observe the fixed value $y^* = 1$.

$$L(\mu; y^* = 1) = [p_{Y(y \mid \mu)}]_{y=1} = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(1 - \mu)^2}{2}\right).$$

One continuous observation gives a likelihood too


Suppose $Y \mid \mu \sim \mathcal{N}(\mu, 1)$ and we observe the fixed value $y^* = 1$.

$$L(\mu; y^* = 1) = \left[p_{Y(y \mid \mu)} \right]_{y=1} = \frac{1}{\sqrt{2\pi}} \exp \left(-\frac{(1 - \mu)^2}{2} \right).$$

Candidate $\mu = 1$: $L(1) \approx 0.399$

Candidate $\mu = 0$: $L(0) \approx 0.242$

$\mu = 1$ centres the Gaussian at $y^* = 1$
and gives the largest **density**.



One continuous observation gives a likelihood too


Suppose $Y \mid \mu \sim \mathcal{N}(\mu, 1)$ and we observe the fixed value $y^* = 1$.

$$L(\mu; y^* = 1) = \left[p_{Y(y \mid \mu)} \right]_{y=1} = \frac{1}{\sqrt{2\pi}} \exp \left(-\frac{(1 - \mu)^2}{2} \right).$$

Candidate $\mu = 1$: $L(1) \approx 0.399$

Candidate $\mu = 0$: $L(0) \approx 0.242$

$\mu = 1$ centres the Gaussian at $y^* = 1$ and gives the largest **density**.



For continuous Y , this is a density value—not $P(Y = 1)$. Fix the observation, then compare parameter values.

A dataset needs a joint likelihood

For two observations, the likelihood starts as a **joint** probability or density:

$$L(\boldsymbol{\theta}; \mathcal{D}) = p_{\boldsymbol{\theta}}(y_1, y_2 \mid \boldsymbol{x}_1, \boldsymbol{x}_2).$$

For two observations, the likelihood starts as a **joint** probability or density:

$$L(\boldsymbol{\theta}; \mathcal{D}) = p_{\boldsymbol{\theta}}(y_1, y_2 \mid \boldsymbol{x}_1, \boldsymbol{x}_2).$$

We may write a product only after making an assumption about how the observations relate.

Question: when is $p(y_1, y_2 \mid \boldsymbol{\theta}) = p(y_1 \mid \boldsymbol{\theta})p(y_2 \mid \boldsymbol{\theta})$ valid?

A dataset needs a joint likelihood

For two observations, the likelihood starts as a **joint** probability or density:

$$L(\boldsymbol{\theta}; \mathcal{D}) = p_{\boldsymbol{\theta}}(y_1, y_2 \mid \boldsymbol{x}_1, \boldsymbol{x}_2).$$

We may write a product only after making an assumption about how the observations relate.

Question: when is $p(y_1, y_2 \mid \boldsymbol{\theta}) = p(y_1 \mid \boldsymbol{\theta})p(y_2 \mid \boldsymbol{\theta})$ valid?

The product is a consequence of **conditional independence**; it is not automatic.

Independent means “no extra information”

D DERIVATION

After fixing θ , observing one outcome does not change the distribution of another:

$$p(y_2 \mid y_1, \theta) = p(y_2 \mid \theta).$$

Independent means “no extra information”

After fixing θ , observing one outcome does not change the distribution of another:

$$p(y_2 \mid y_1, \theta) = p(y_2 \mid \theta).$$

The chain rule then becomes

$$p(y_1, y_2 \mid \theta) = p(y_1 \mid \theta)p(y_2 \mid y_1, \theta) = p(y_1 \mid \theta)p(y_2 \mid \theta).$$

Independent means “no extra information”

After fixing θ , observing one outcome does not change the distribution of another:

$$p(y_2 \mid y_1, \theta) = p(y_2 \mid \theta).$$

The chain rule then becomes

$$p(y_1, y_2 \mid \theta) = p(y_1 \mid \theta)p(y_2 \mid y_1, \theta) = p(y_1 \mid \theta)p(y_2 \mid \theta).$$

Coin-flip example. If $P(H \mid \theta) = \theta$ and the flips are independent, then $P(H, H \mid \theta) = \theta^2$.

Independent means “no extra information”

After fixing θ , observing one outcome does not change the distribution of another:

$$p(y_2 \mid y_1, \theta) = p(y_2 \mid \theta).$$

The chain rule then becomes

$$p(y_1, y_2 \mid \theta) = p(y_1 \mid \theta)p(y_2 \mid y_1, \theta) = p(y_1 \mid \theta)p(y_2 \mid \theta).$$

Coin-flip example. If $P(H \mid \theta) = \theta$ and the flips are independent, then $P(H, H \mid \theta) = \theta^2$.

Independence is what licenses multiplication.

Identically distributed means “the same sampling rule”

D DERIVATION

Coin flips

$Y_i \mid \theta \sim \text{Bernoulli}(\theta)$ for every i .
Each flip uses the same bias θ .

Supervised learning

The pairs $(\mathbf{X}_i, Y_i)$ come from the same population, and one shared θ defines $p_\theta(y_i \mid \mathbf{x}_i)$.

Identically distributed means “the same sampling rule”

D DERIVATION

Coin flips

$Y_i \mid \theta \sim \text{Bernoulli}(\theta)$ for every i .
Each flip uses the same bias θ .

Supervised learning

The pairs $(\mathbf{X}_i, Y_i)$ come from the same population, and one shared θ defines $p_\theta(y_i \mid \mathbf{x}_i)$.

“Independent” and “identically distributed” are two separate claims; one can hold without the other.

Identically distributed means “the same sampling rule”

D DERIVATION

Coin flips

$Y_i \mid \theta \sim \text{Bernoulli}(\theta)$ for every i .
Each flip uses the same bias θ .

Supervised learning

The pairs $(\mathbf{X}_i, Y_i)$ come from the same population, and one shared θ defines $p_\theta(y_i \mid \mathbf{x}_i)$.

“Independent” and “identically distributed” are two separate claims; one can hold without the other.

i.i.d. = independent observations + one common data-generating rule.

Not independent: daily weather

Tomorrow's weather depends on today's. A first-order Markov model writes

$$p(y_{1:T} \mid \theta) = p(y_1 \mid \theta) \prod_{t=2}^T p(y_t \mid y_{t-1}, \theta).$$

The factors are **conditional on the previous day**, not independent marginals.

Not independent: daily weather

Tomorrow's weather depends on today's. A first-order Markov model writes

$$p(y_{1:T} \mid \theta) = p(y_1 \mid \theta) \prod_{t=2}^T p(y_t \mid y_{t-1}, \theta).$$

The factors are **conditional on the previous day**, not independent marginals.

Independent but not identical: two sensors

Two sensors measure the same temperature μ independently, but have different noise levels:

$$Y_i \mid \mu \sim \mathcal{N}(\mu, \sigma_i^2), \quad \sigma_1 \neq \sigma_2.$$

The likelihood is a product, but its factors use different σ_i .

Not independent: daily weather

Tomorrow's weather depends on today's. A first-order Markov model writes

$$p(y_{1:T} \mid \theta) = p(y_1 \mid \theta) \prod_{t=2}^T p(y_t \mid y_{t-1}, \theta).$$

The factors are **conditional on the previous day**, not independent marginals.

Independent but not identical: two sensors

Two sensors measure the same temperature μ independently, but have different noise levels:

$$Y_i \mid \mu \sim \mathcal{N}(\mu, \sigma_i^2), \quad \sigma_1 \neq \sigma_2.$$

The likelihood is a product, but its factors use different σ_i .

i.i.d. is a useful special case—not a universal law.

Let R_1 mean “rain today” and R_2 mean “rain tomorrow”. Suppose

$$P(R_1) = 0.4, \quad P(R_2|R_1) = 0.8, \quad P(R_2|\neg R_1) = 0.2.$$

1. Use the conditional rule

$$\begin{aligned} P(R_1, R_2) &= P(R_1)P(R_2|R_1) \\ &= 0.4 \times 0.8 = 0.32. \end{aligned}$$

Both days are rainy with probability 0.32.

Let R_1 mean “rain today” and R_2 mean “rain tomorrow”. Suppose

$$P(R_1) = 0.4, \quad P(R_2|R_1) = 0.8, \quad P(R_2|\neg R_1) = 0.2.$$

1. Use the conditional rule

$$\begin{aligned} P(R_1, R_2) &= P(R_1)P(R_2|R_1) \\ &= 0.4 \times 0.8 = 0.32. \end{aligned}$$

Both days are rainy with probability 0.32.

2. Test independence

$$P(R_2) = 0.4 \times 0.8 + 0.6 \times 0.2 = 0.44.$$

If the days were independent,

$$P(R_1)P(R_2) = 0.4 \times 0.44 = 0.176,$$

not 0.32.

Let R_1 mean “rain today” and R_2 mean “rain tomorrow”. Suppose

$$P(R_1) = 0.4, \quad P(R_2|R_1) = 0.8, \quad P(R_2|\neg R_1) = 0.2.$$

1. Use the conditional rule

$$\begin{aligned} P(R_1, R_2) &= P(R_1)P(R_2|R_1) \\ &= 0.4 \times 0.8 = 0.32. \end{aligned}$$

Both days are rainy with probability 0.32.

2. Test independence

$$P(R_2) = 0.4 \times 0.8 + 0.6 \times 0.2 = 0.44.$$

If the days were independent,

$$P(R_1)P(R_2) = 0.4 \times 0.44 = 0.176,$$

not 0.32.

Dependence does not forbid multiplication: use a conditional probability.

Given smoke S , two alarms act independently, but their sensitivities differ:

$$P(A_1 = 1|S) = 0.9, \quad P(A_2 = 1|S) = 0.6.$$

Independence

Both alarms ring with probability

$$\begin{aligned} P(A_1 = 1, A_2 = 1|S) \\ = 0.9 \times 0.6 = 0.54. \end{aligned}$$

The product is valid.

Given smoke S , two alarms act independently, but their sensitivities differ:

$$P(A_1 = 1|S) = 0.9, \quad P(A_2 = 1|S) = 0.6.$$

Independence

Both alarms ring with probability

$$\begin{aligned} P(A_1 = 1, A_2 = 1|S) \\ = 0.9 \times 0.6 = 0.54. \end{aligned}$$

The product is valid.

Not identically distributed

Alarm 1 detects smoke 90% of the time;
Alarm 2 detects it 60% of the time.

$$0.9 \neq 0.6$$

We cannot reuse one common factor for both alarms.

Worked example: independent does not mean identical

Given smoke S , two alarms act independently, but their sensitivities differ:

$$P(A_1 = 1|S) = 0.9, \quad P(A_2 = 1|S) = 0.6.$$

Independence

Both alarms ring with probability

$$\begin{aligned} P(A_1 = 1, A_2 = 1|S) \\ = 0.9 \times 0.6 = 0.54. \end{aligned}$$

The product is valid.

Not identically distributed

Alarm 1 detects smoke 90% of the time;
Alarm 2 detects it 60% of the time.

$$0.9 \neq 0.6$$

We cannot reuse one common factor for both alarms.

Independent says “multiply”; identical says “reuse the same sampling rule”.

IID turns the joint likelihood into per-example factors

D DERIVATION

For i.i.d. supervised examples, treat the observed inputs as fixed. Conditional independence gives

$$p(\mathcal{D} \mid \theta) = \prod_{i=1}^n p_{\theta}(y_i \mid x_i).$$

For i.i.d. supervised examples, treat the observed inputs as fixed. Conditional independence gives

$$p(\mathcal{D} \mid \theta) = \prod_{i=1}^n p_{\theta}(y_i \mid x_i).$$

The **factorization** comes from independence; the shared factor form comes from the common model p_{θ} .

Next we compute this product, then introduce a transform that makes it easier to optimize.

IID turns the joint likelihood into per-example factors

D DERIVATION

For i.i.d. supervised examples, treat the observed inputs as fixed. Conditional independence gives

$$p(\mathcal{D} \mid \theta) = \prod_{i=1}^n p_{\theta}(y_i \mid x_i).$$

The **factorization** comes from independence; the shared factor form comes from the common model p_{θ} .

Next we compute this product, then introduce a transform that makes it easier to optimize.

This is the bridge from one-example likelihood to a dataset loss.

Work two tiny Bernoulli datasets

D DERIVATION

Compare a fair coin $\theta = 0.5$ with a head-biased coin $\theta = 0.8$.

Work two tiny Bernoulli datasets

Compare a fair coin $\theta = 0.5$ with a head-biased coin $\theta = 0.8$.

Observed data	$L(0.5)$	$L(0.8)$	Which explains it better?
H, T	$0.5(0.5) = 0.25$	$0.8(0.2) = 0.16$	fair coin
H, H	$0.5^2 = 0.25$	$0.8^2 = 0.64$	head-biased coin

Work two tiny Bernoulli datasets

Compare a fair coin $\theta = 0.5$ with a head-biased coin $\theta = 0.8$.

Observed data	$L(0.5)$	$L(0.8)$	Which explains it better?
H, T	$0.5(0.5) = 0.25$	$0.8(0.2) = 0.16$	fair coin
H, H	$0.5^2 = 0.25$	$0.8^2 = 0.64$	head-biased coin

Same candidate models, different fixed data. Likelihood ranks candidates by how well they explain the observations actually seen.

Work two tiny Bernoulli datasets

Compare a fair coin $\theta = 0.5$ with a head-biased coin $\theta = 0.8$.

Observed data	$L(0.5)$	$L(0.8)$	Which explains it better?
H, T	$0.5(0.5) = 0.25$	$0.8(0.2) = 0.16$	fair coin
H, H	$0.5^2 = 0.25$	$0.8^2 = 0.64$	head-biased coin

Same candidate models, different fixed data. Likelihood ranks candidates by how well they explain the observations actually seen.

Multiply one Bernoulli term per independent flip.

Two Gaussian measurements: multiply densities

D DERIVATION

Let $Y_i \mid \mu \sim \mathcal{N}(\mu, 1)$ independently, and observe $y_1 = 1, y_2 = 2$.

Two Gaussian measurements: multiply densities

D DERIVATION

Let $Y_i \mid \mu \sim \mathcal{N}(\mu, 1)$ independently, and observe $y_1 = 1, y_2 = 2$.

Candidate $\mu = 1.5$

$$p(y_1 \mid \mu) \approx 0.352$$

$$p(y_2 \mid \mu) \approx 0.352$$

$$L(1.5) \approx 0.352 \times 0.352 = 0.124$$

Candidate $\mu = 0$

$$p(y_1 \mid \mu) \approx 0.242$$

$$p(y_2 \mid \mu) \approx 0.054$$

$$L(0) \approx 0.242 \times 0.054 = 0.013$$

Two Gaussian measurements: multiply densities

Let $Y_i \mid \mu \sim \mathcal{N}(\mu, 1)$ independently, and observe $y_1 = 1, y_2 = 2$.

Candidate $\mu = 1.5$

$$p(y_1 \mid \mu) \approx 0.352$$

$$p(y_2 \mid \mu) \approx 0.352$$

$$L(1.5) \approx 0.352 \times 0.352 = 0.124$$

Candidate $\mu = 0$

$$p(y_1 \mid \mu) \approx 0.242$$

$$p(y_2 \mid \mu) \approx 0.054$$

$$L(0) \approx 0.242 \times 0.054 = 0.013$$

For continuous data we multiply **densities; $\mu = 1.5$ gives the larger likelihood.**

Maximum likelihood estimation

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \max_{\boldsymbol{\theta}} p(\mathcal{D} \mid \boldsymbol{\theta})$$

Maximum likelihood estimation

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} p(\mathcal{D} \mid \theta)$$

Maximum likelihood estimation (MLE). Choose the parameter values that assign the highest likelihood to the observed dataset.

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L01/00_likelihood_iid_worked_examples.ipynb

Explore `torch.distributions`, `log_prob`, sampling, i.i.d. factorization, coin-toss MLE, Gaussian regression, Bernoulli classification, and MAP with Gaussian, Laplace, and Student- t priors.

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L01/00_likelihood_iid_worked_examples.ipynb

Explore `torch.distributions`, `log_prob`, sampling, i.i.d. factorization, coin-toss MLE, Gaussian regression, Bernoulli classification, and MAP with Gaussian, Laplace, and Student- t priors.

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L01/03_robust_linear_regression.ipynb

Short companion: fit the same line with Gaussian/MSE, Laplace/MAE, and Student- t losses; then see which observations control each estimate.

Observed: $H, H, T, T, T, H, H, T, T, T$

Hence $n_H = 4$ and $n_T = 6$.

Observed: $H, H, T, T, T, H, H, T, T, T$

Hence $n_H = 4$ and $n_T = 6$.

Model: $p(H \mid \theta) = \theta, p(T \mid \theta) = 1 - \theta$.

Coin flips: a likelihood you can compute

Observed: $H, H, T, T, T, H, H, T, T, T$

Hence $n_H = 4$ and $n_T = 6$.


Model: $p(H \mid \theta) = \theta$, $p(T \mid \theta) = 1 - \theta$.

Because the flips are i.i.d.,

$$L(\theta) = \prod_i p(x_i \mid \theta)$$

$$= \theta^{n_H} (1 - \theta)^{n_T}$$

$$= \theta^4 (1 - \theta)^6$$



$$L(0.4) \approx 0.00119$$

$$L(0.7) \approx 0.000175$$

about $6.8 \times$ larger at $\theta = 0.4$

Coin flips: a likelihood you can compute

D DERIVATION

Observed: $H, H, T, T, T, H, H, T, T, T$

Hence $n_H = 4$ and $n_T = 6$.


Model: $p(H \mid \theta) = \theta$, $p(T \mid \theta) = 1 - \theta$.

Because the flips are i.i.d.,

$$L(\theta) = \prod_i p(x_i \mid \theta)$$

$$= \theta^{n_H} (1 - \theta)^{n_T}$$

$$= \theta^4 (1 - \theta)^6$$



$$L(0.4) \approx 0.00119$$

$$L(0.7) \approx 0.000175$$

about $6.8 \times$ larger at $\theta = 0.4$

The likelihood peaks at $\theta = 0.4$ — the observed head fraction $\frac{4}{10}$.

The **log-likelihood** is the logarithm of the likelihood — not a different model:

$$\ell(\theta) := \log L(\theta) = \log[\theta^4(1 - \theta)^6] = 4 \log \theta + 6 \log(1 - \theta)$$

The **log-likelihood** is the logarithm of the likelihood — not a different model:

$$\ell(\theta) := \log L(\theta) = \log[\theta^4(1 - \theta)^6] = 4 \log \theta + 6 \log(1 - \theta)$$

Candidate $\theta = 0.4$

$$L(0.4) \approx 0.00119$$

$$\ell(0.4) = 4 \log 0.4 + 6 \log 0.6 \approx -6.73$$

Candidate $\theta = 0.7$

$$L(0.7) \approx 0.000175$$

$$\ell(0.7) = 4 \log 0.7 + 6 \log 0.3 \approx -8.65$$

The **log-likelihood** is the logarithm of the likelihood — not a different model:

$$\ell(\theta) := \log L(\theta) = \log[\theta^4(1 - \theta)^6] = 4 \log \theta + 6 \log(1 - \theta)$$

Candidate $\theta = 0.4$

$$L(0.4) \approx 0.00119$$

$$\ell(0.4) = 4 \log 0.4 + 6 \log 0.6 \approx -6.73$$

Candidate $\theta = 0.7$

$$L(0.7) \approx 0.000175$$

$$\ell(0.7) = 4 \log 0.7 + 6 \log 0.3 \approx -8.65$$

Likelihoods below 1 have negative logs; among them, -6.73 is the larger value.

The **log-likelihood** is the logarithm of the likelihood — not a different model:

$$\ell(\theta) := \log L(\theta) = \log[\theta^4(1 - \theta)^6] = 4 \log \theta + 6 \log(1 - \theta)$$

Candidate $\theta = 0.4$

$$L(0.4) \approx 0.00119$$

$$\ell(0.4) = 4 \log 0.4 + 6 \log 0.6 \approx -6.73$$

Candidate $\theta = 0.7$

$$L(0.7) \approx 0.000175$$

$$\ell(0.7) = 4 \log 0.7 + 6 \log 0.3 \approx -8.65$$

Likelihoods below 1 have negative logs; among them, -6.73 is the larger value.

The observed sequence has higher log-likelihood at $\theta = 0.4$.

Why logs? The maximizing parameter does not change

D DERIVATION

The logarithm is **strictly increasing** on positive numbers:

$$a > b > 0 \implies \log a > \log b$$

Why logs? The maximizing parameter does not change

The logarithm is **strictly increasing** on positive numbers:

$$a > b > 0 \implies \log a > \log b$$

Therefore it preserves the ordering of all positive likelihood values:

$$\begin{aligned} L(0.4) &> L(0.7) \\ \text{apply the increasing function } \log \\ \log L(0.4) &> \log L(0.7) \end{aligned}$$

Why logs? The maximizing parameter does not change

The logarithm is **strictly increasing** on positive numbers:

$$a > b > 0 \implies \log a > \log b$$

Therefore it preserves the ordering of all positive likelihood values:

$$\begin{aligned} L(0.4) &> L(0.7) \\ \text{apply the increasing function } \log \\ \log L(0.4) &> \log L(0.7) \end{aligned}$$

$$\arg \max_{\theta} L(\theta) = \arg \max_{\theta} \log L(\theta)$$

Why logs? The maximizing parameter does not change

The logarithm is **strictly increasing** on positive numbers:

$$a > b > 0 \implies \log a > \log b$$

Therefore it preserves the ordering of all positive likelihood values:

$$\begin{aligned} L(0.4) &> L(0.7) \\ \text{apply the increasing function } \log \\ \log L(0.4) &> \log L(0.7) \end{aligned}$$

$$\arg \max_{\theta} L(\theta) = \arg \max_{\theta} \log L(\theta)$$

Both likelihood and log-likelihood rank $\theta = 0.4$ above $\theta = 0.7$.

Poisson example: logging preserves the peak


Suppose $Y \sim \text{Poisson}(\lambda)$ and the observed count is $y = 3$.

Poisson example: logging preserves the peak

Suppose $Y \sim \text{Poisson}(\lambda)$ and the observed count is $y = 3$.


LIKELIHOOD

$$L(\lambda) = e^{-\lambda} \frac{\lambda^3}{3!}$$



LOG-LIKELIHOOD

$$\log L(\lambda) = 3 \log \lambda - \lambda - \log 3!$$




Poisson example: logging preserves the peak

Suppose $Y \sim \text{Poisson}(\lambda)$ and the observed count is $y = 3$.


LIKELIHOOD

$$L(\lambda) = e^{-\lambda} \frac{\lambda^3}{3!}$$



LOG-LIKELIHOOD

$$\log L(\lambda) = 3 \log \lambda - \lambda - \log 3!$$



Both curves peak at $\hat{\lambda}_{\text{MLE}} = 3$: the log changes the vertical scale, not the maximizing parameter.

Why logs? Products become stable sums

For independent observations,

$$L(\boldsymbol{\theta}) = \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \implies \log L(\boldsymbol{\theta}) = \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

Why logs? Products become stable sums

For independent observations,

$$L(\boldsymbol{\theta}) = \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \implies \log L(\boldsymbol{\theta}) = \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

Worked numerical example. Suppose 200 observed outcomes each receive probability 0.01.

Direct product	$L = 0.01^{200} = 10^{-400} < 4.94 \times 10^{-324} \rightarrow \text{underflows to } 0$
Log space	$\log L = 200 \log(0.01) \approx -921.03 \rightarrow \text{finite and representable}$

Float64 range: $\pm 1.80 \times 10^{308}$ · smallest positive normal 2.23×10^{-308} · smallest subnormal 4.94×10^{-324} .

Why logs? Products become stable sums

For independent observations,

$$L(\boldsymbol{\theta}) = \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \implies \log L(\boldsymbol{\theta}) = \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

Worked numerical example. Suppose 200 observed outcomes each receive probability 0.01.

Direct product	$L = 0.01^{200} = 10^{-400} < 4.94 \times 10^{-324} \rightarrow \text{underflows to } 0$
Log space	$\log L = 200 \log(0.01) \approx -921.03 \rightarrow \text{finite and representable}$

Float64 range: $\pm 1.80 \times 10^{308}$ · smallest positive normal 2.23×10^{-308} · smallest subnormal 4.94×10^{-324} .

Logs make optimization easier and prevent numerical underflow.

Coin flips: solve for the stationary point

D DERIVATION

For $L(\theta) = \theta^{n_H} (1 - \theta)^{n_T}$, the log-likelihood is

$$\log L(\theta) = n_H \log \theta + n_T \log(1 - \theta)$$

Coin flips: solve for the stationary point

For $L(\theta) = \theta^{n_H} (1 - \theta)^{n_T}$, the log-likelihood is

$$\log L(\theta) = n_H \log \theta + n_T \log(1 - \theta)$$

Set the derivative to zero:

$$\frac{d}{d\theta} \log L(\theta) = \frac{n_H}{\theta} - \frac{n_T}{1 - \theta} = 0$$

Coin flips: solve for the stationary point

For $L(\theta) = \theta^{n_H} (1 - \theta)^{n_T}$, the log-likelihood is

$$\log L(\theta) = n_H \log \theta + n_T \log(1 - \theta)$$

Set the derivative to zero:

$$\frac{d}{d\theta} \log L(\theta) = \frac{n_H}{\theta} - \frac{n_T}{1 - \theta} = 0$$

$$\implies \theta^* = \frac{n_H}{n_H + n_T} = \frac{4}{10} = 0.4$$

For $L(\theta) = \theta^{n_H}(1 - \theta)^{n_T}$, the log-likelihood is

$$\log L(\theta) = n_H \log \theta + n_T \log(1 - \theta)$$

Set the derivative to zero:

$$\frac{d}{d\theta} \log L(\theta) = \frac{n_H}{\theta} - \frac{n_T}{1 - \theta} = 0$$

$$\Rightarrow \theta^* = \frac{n_H}{n_H + n_T} = \frac{4}{10} = 0.4$$

$\theta^* = 0.4$ is a **candidate** optimum. Next, check the curvature.

The second derivative confirms a maximum

Differentiate the log-likelihood again:

$$\frac{d^2}{d\theta^2} \log L(\theta) = -\frac{n_H}{\theta^2} - \frac{n_T}{(1-\theta)^2}$$

Differentiate the log-likelihood again:

$$\frac{d^2}{d\theta^2} \log L(\theta) = -\frac{n_H}{\theta^2} - \frac{n_T}{(1-\theta)^2}$$

For $n_H, n_T > 0$ and $0 < \theta < 1$, both terms are negative:

$$\frac{d^2}{d\theta^2} \log L(\theta) < 0$$

The second derivative confirms a maximum

Differentiate the log-likelihood again:

$$\frac{d^2}{d\theta^2} \log L(\theta) = -\frac{n_H}{\theta^2} - \frac{n_T}{(1-\theta)^2}$$

For $n_H, n_T > 0$ and $0 < \theta < 1$, both terms are negative:

$$\frac{d^2}{d\theta^2} \log L(\theta) < 0$$

At the stationary point $\theta^* = 0.4$,

$$-\frac{4}{0.16} - \frac{6}{0.36} = -25 - 16.67 \approx -41.67 < 0$$

The second derivative is negative everywhere, so the curve bends downward (is strictly concave). Its one stationary point is the unique global maximum: $\hat{\theta}_{\text{MLE}} = 0.4$.

Maximizing likelihood = minimizing NLL

D DERIVATION

$$\arg \max_{\theta} \prod_i p_{\theta}(y_i \mid \mathbf{x}_i) = \arg \max_{\theta} \sum_i \log p_{\theta}(y_i \mid \mathbf{x}_i)$$

Maximizing likelihood = minimizing NLL

D DERIVATION

$$\arg \max_{\boldsymbol{\theta}} \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) = \arg \max_{\boldsymbol{\theta}} \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

$$= \arg \min_{\boldsymbol{\theta}} \left(- \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \right)$$

Maximizing likelihood = minimizing NLL

$$\arg \max_{\boldsymbol{\theta}} \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) = \arg \max_{\boldsymbol{\theta}} \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

$$= \arg \min_{\boldsymbol{\theta}} \left(- \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \right)$$

$$\ell_{i(\boldsymbol{\theta})} = -\log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

Empirical risk = average training loss

D DERIVATION

In ML, **risk** means expected loss. Because the data-generating distribution is unknown, we estimate it from the observed training set.

Empirical risk = average training loss

D DERIVATION

In ML, **risk** means expected loss. Because the data-generating distribution is unknown, we estimate it from the observed training set.

Empirical risk = average loss on $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$.

Empirical risk = average training loss

In ML, **risk** means expected loss. Because the data-generating distribution is unknown, we estimate it from the observed training set.

Empirical risk = average loss on $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$.

Per example	$\ell_{i(\boldsymbol{\theta})} = -\log p_{\boldsymbol{\theta}}(y_i \mid \mathbf{x}_i)$
Training set	$\hat{R}(\boldsymbol{\theta}) = \frac{1}{n} \sum_i \ell_{i(\boldsymbol{\theta})} = -\frac{1}{n} \log p(\mathcal{D} \mid \boldsymbol{\theta})$

The final equality uses conditional independence: $p(\mathcal{D} \mid \boldsymbol{\theta}) = \prod_i p_{\boldsymbol{\theta}}(y_i \mid \mathbf{x}_i)$.

Empirical risk = average training loss

In ML, **risk** means expected loss. Because the data-generating distribution is unknown, we estimate it from the observed training set.

Empirical risk = average loss on $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$.

Per example	$\ell_{i(\theta)} = -\log p_{\theta}(y_i \mid x_i)$
Training set	$\hat{R}(\theta) = \frac{1}{n} \sum_i \ell_{i(\theta)} = -\frac{1}{n} \log p(\mathcal{D} \mid \theta)$

The final equality uses conditional independence: $p(\mathcal{D} \mid \theta) = \prod_i p_{\theta}(y_i \mid x_i)$.

Minimizing empirical NLL is maximum likelihood; $\frac{1}{n}$ only rescales the objective.

Recall the Uniform observation model:

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \begin{cases} 1/(b - a) & a \leq y_i \leq b \\ 0 & \text{otherwise} \end{cases}$$

Recall the Uniform observation model:

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \begin{cases} 1/(b - a) & a \leq y_i \leq b \\ 0 & \text{otherwise} \end{cases}$$

If an observed target lies outside its support, then

$$p_{\theta}(y_i \mid \mathbf{x}_i) = 0$$

Recall the Uniform observation model:

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \begin{cases} 1/(b - a) & a \leq y_i \leq b \\ 0 & \text{otherwise} \end{cases}$$

If an observed target lies outside its support, then

$$p_{\theta}(y_i \mid \mathbf{x}_i) = 0$$

Its contribution to the negative log-likelihood is

$$\ell_{i(\theta)} = -\log p_{\theta}(y_i \mid \mathbf{x}_i) = -\log 0 = +\infty$$

Recall the Uniform observation model:

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \begin{cases} 1/(b - a) & a \leq y_i \leq b \\ 0 & \text{otherwise} \end{cases}$$

If an observed target lies outside its support, then

$$p_{\theta}(y_i \mid \mathbf{x}_i) = 0$$

Its contribution to the negative log-likelihood is

$$\ell_{i(\theta)} = -\log p_{\theta}(y_i \mid \mathbf{x}_i) = -\log 0 = +\infty$$

Zero likelihood becomes an infinite training penalty.

For a $\text{Uniform}(a, b)$ observation model, what is the NLL if the target lies outside $[a, b]$?

A. 0

B. A fixed finite penalty

C. $-\log 0 = +\infty$

D. It depends only on σ

Correct: C — $-\log 0 = +\infty$

Why: The model assigns zero density outside its support, so the observation is impossible under the assumption.

Regression: why MSE?

Regression as signal plus noise

$$y_i = f_{\boldsymbol{\theta}}(\mathbf{x}_i) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$


Regression as signal plus noise

$$y_i = f_{\boldsymbol{\theta}}(\mathbf{x}_i) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$


Therefore the conditional output distribution is

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(f_{\boldsymbol{\theta}}(\mathbf{x}_i), \sigma^2)$$

What the assumption means visually



What the assumption means visually



Gaussian noise lives in the **output direction**, conditional on x — a vertical bell at every input.

$$p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - f_{\boldsymbol{\theta}}(\boldsymbol{x}_i))^2}{2\sigma^2}\right)$$

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - f_{\theta}(\mathbf{x}_i))^2}{2\sigma^2}\right)$$

Negative log of a single example:

$$-\log p_{\theta}(y_i \mid \mathbf{x}_i) = \frac{(y_i - f_{\theta}(\mathbf{x}_i))^2}{2\sigma^2} + C$$

MSE is Gaussian MLE

Sum over the dataset:

$$-\log p(\mathcal{D} \mid \boldsymbol{\theta}) = \frac{1}{2\sigma^2} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2 + C$$

MSE is Gaussian MLE

Sum over the dataset:

$$-\log p(\mathcal{D} \mid \boldsymbol{\theta}) = \frac{1}{2\sigma^2} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2 + C$$

With σ fixed, the minimizer is

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \min_{\boldsymbol{\theta}} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2$$

MSE is Gaussian MLE

Sum over the dataset:

$$-\log p(\mathcal{D} \mid \boldsymbol{\theta}) = \frac{1}{2\sigma^2} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2 + C$$

With σ fixed, the minimizer is

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \min_{\boldsymbol{\theta}} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2$$

MSE = Gaussian NLL, up to constants.

Linear regression makes f_θ explicit

D DERIVATION

For the one-dimensional picture,

$$\hat{y}_i = \theta_0 + \theta_1 x_i$$

For the one-dimensional picture,

$$\hat{y}_i = \theta_0 + \theta_1 x_i$$

For d features, prepend a constant 1 so the intercept is part of the same parameter vector:

$$\tilde{x}_i = (1, x_i), \quad \theta = (\theta_0, \theta_1, \dots, \theta_d)$$

Linear regression makes f_{θ} explicit

For the one-dimensional picture,

$$\hat{y}_i = \theta_0 + \theta_1 x_i$$

For d features, prepend a constant 1 so the intercept is part of the same parameter vector:

$$\tilde{x}_i = (1, x_i), \quad \theta = (\theta_0, \theta_1, \dots, \theta_d)$$

Then the same notation used throughout the lecture becomes

$$\hat{y}_i = f_{\theta}(x_i) = \theta^{\top} \tilde{x}_i$$

θ_0 is the intercept; $\theta_1, \dots, \theta_d$ are the slopes.

Substitute the linear mean into the observation model:

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i, \sigma^2)$$

Substitute the linear mean into the observation model:

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i, \sigma^2)$$

The vertical residual of point i is

$$r_i = y_i - \hat{y}_i = y_i - \boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i$$

Substitute the linear mean into the observation model:

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i, \sigma^2)$$

The vertical residual of point i is

$$r_i = y_i - \hat{y}_i = y_i - \boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i$$

With σ fixed, maximum likelihood therefore gives

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \min_{\boldsymbol{\theta}} \sum_i \underbrace{r_i^2}_{\text{squared vertical error}}$$

Ordinary least squares is Gaussian maximum likelihood for a linear mean.

The normal equations solve linear regression

D DERIVATION

Stack the augmented inputs into $\tilde{X} = [1 \quad X]$. Then

$$\hat{y} = \tilde{X}\theta, \quad J(\theta) = \|\mathbf{y} - \tilde{X}\theta\|_2^2$$

The normal equations solve linear regression

Stack the augmented inputs into $\tilde{\mathbf{X}} = [\mathbf{1} \quad \mathbf{X}]$. Then

$$\hat{y} = \tilde{\mathbf{X}}\boldsymbol{\theta}, \quad J(\boldsymbol{\theta}) = \|\mathbf{y} - \tilde{\mathbf{X}}\boldsymbol{\theta}\|_2^2$$

The gradient $\nabla_{\boldsymbol{\theta}} J$ collects one partial derivative per parameter. Set it to zero:

$$\nabla_{\boldsymbol{\theta}} J = 2\tilde{\mathbf{X}}^\top (\tilde{\mathbf{X}}\boldsymbol{\theta} - \mathbf{y}) = \mathbf{0}$$

The normal equations solve linear regression

Stack the augmented inputs into $\tilde{X} = [1 \quad X]$. Then

$$\hat{y} = \tilde{X}\theta, \quad J(\theta) = \|y - \tilde{X}\theta\|_2^2$$

The gradient $\nabla_{\theta}J$ collects one partial derivative per parameter. Set it to zero:

$$\nabla_{\theta}J = 2\tilde{X}^{\top}(\tilde{X}\theta - y) = 0$$

This gives the **normal equations**:

$$\tilde{X}^{\top}\tilde{X}\hat{\theta} = \tilde{X}^{\top}y$$

The normal equations solve linear regression

Stack the augmented inputs into $\tilde{\mathbf{X}} = [\mathbf{1} \quad \mathbf{X}]$. Then

$$\hat{y} = \tilde{\mathbf{X}}\boldsymbol{\theta}, \quad J(\boldsymbol{\theta}) = \|\mathbf{y} - \tilde{\mathbf{X}}\boldsymbol{\theta}\|_2^2$$

The gradient $\nabla_{\boldsymbol{\theta}} J$ collects one partial derivative per parameter. Set it to zero:

$$\nabla_{\boldsymbol{\theta}} J = 2\tilde{\mathbf{X}}^\top (\tilde{\mathbf{X}}\boldsymbol{\theta} - \mathbf{y}) = \mathbf{0}$$

This gives the **normal equations**:

$$\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \hat{\boldsymbol{\theta}} = \tilde{\mathbf{X}}^\top \mathbf{y}$$

If $\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}}$ is invertible,

$$\hat{\boldsymbol{\theta}} = (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top \mathbf{y}$$

Linear regression has a closed form; neural networks keep the MSE objective but use iterative optimization.

Why earlier ML courses could ignore σ^2

σ^2 is the observation-noise variance: the expected squared spread of targets around the mean prediction.

For n Gaussian observations, write $r_i = y_i - \mu_i$. The dataset NLL is

$$\mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \frac{1}{2\sigma^2} \sum_i r_i^2 + \frac{n}{2} \log \sigma^2 + C.$$

σ^2 is the observation-noise variance: the expected squared spread of targets around the mean prediction.

For n Gaussian observations, write $r_i = y_i - \mu_i$. The dataset NLL is

$$\mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \frac{1}{2\sigma^2} \sum_i r_i^2 + \frac{n}{2} \log \sigma^2 + C.$$

If one global σ^2 is fixed, then $\frac{n}{2} \log \sigma^2$ is constant and $\frac{1}{2\sigma^2}$ is only a positive scale factor:

$$\arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \arg \min_{\boldsymbol{\theta}} \sum_i r_i^2.$$

Why earlier ML courses could ignore σ^2

σ^2 is the observation-noise variance: the expected squared spread of targets around the mean prediction.

For n Gaussian observations, write $r_i = y_i - \mu_i$. The dataset NLL is

$$\mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \frac{1}{2\sigma^2} \sum_i r_i^2 + \frac{n}{2} \log \sigma^2 + C.$$

If one global σ^2 is fixed, then $\frac{n}{2} \log \sigma^2$ is constant and $\frac{1}{2\sigma^2}$ is only a positive scale factor:

$$\arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \arg \min_{\boldsymbol{\theta}} \sum_i r_i^2.$$

For point predictions, fixed-variance Gaussian MLE gives the same fitted model as MSE.

One global noise level.

Fit the mean model, then estimate the residual spread:

$$\hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{\mu}_i)^2.$$

Use this when a residual plot has roughly the same vertical width across inputs.

How is σ^2 chosen in practice?

One global noise level.

Fit the mean model, then estimate the residual spread:

$$\hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{\mu}_i)^2.$$

Use this when a residual plot has roughly the same vertical width across inputs.

Input-dependent noise.

Let the network predict both

$$(\mu_i, s_i) = f_{\theta}(\mathbf{x}_i),$$

$$s_i = \log \sigma_i^2, \quad \sigma_i^2 = e^{s_i} > 0.$$

Use this when residual spread changes with the input—for example, a funnel shape.

How is σ^2 chosen in practice?

One global noise level.

Fit the mean model, then estimate the residual spread:

$$\hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{\mu}_i)^2.$$

Use this when a residual plot has roughly the same vertical width across inputs.

Input-dependent noise.

Let the network predict both

$$(\mu_i, s_i) = f_{\theta}(\mathbf{x}_i),$$

$$s_i = \log \sigma_i^2, \quad \sigma_i^2 = e^{s_i} > 0.$$

Use this when residual spread changes with the input—for example, a funnel shape.

Fit both cases with Gaussian NLL; validate interval **coverage: the fraction of targets contained by intervals claiming a stated probability.**

HOMOSCEDASTIC



Network output: μ_i

Shared variance: σ^2

same uncertainty at every x_i


Homoscedastic vs input-dependent uncertainty

HOMOSCEDASTIC



Network output: μ_i
Shared variance: σ^2
same uncertainty at every x_i

HETEROSCEDASTIC



Network outputs: μ_i and s_i
 $s_i = \log \sigma_i^2$, $\sigma_i^2 = e^{s_i} > 0$
uncertainty changes with x_i


Homoscedastic vs input-dependent uncertainty

HOMOSCEDASTIC



Network output: μ_i
Shared variance: σ^2
same uncertainty at every x_i

HETEROSCEDASTIC




Network outputs: μ_i and s_i
 $s_i = \log \sigma_i^2$, $\sigma_i^2 = e^{s_i} > 0$
uncertainty changes with x_i

Predict **log-variance** so every real s_i maps to a positive variance e^{s_i} .

A few unusual observations can dominate a fitted line

Q QUESTION

SYNTHETIC · COMPUTED · exact 12-point teaching case; reproduced in the executed robust-notebook appendix




Ten typical observations lie near one linear trend.

— fit to the ten typical points — MSE fit after adding two outliers

A few unusual observations can dominate a fitted line

Q QUESTION

SYNTHETIC · COMPUTED · exact 12-point teaching case; reproduced in the executed robust-notebook appendix



— fit to the ten typical points — MSE fit after adding two outliers


Ten typical observations lie near one linear trend.

Two unusual observations have large vertical errors.

A few unusual observations can dominate a fitted line

Q QUESTION

SYNTHETIC · COMPUTED · exact 12-point teaching case; reproduced in the executed robust-notebook appendix



— fit to the ten typical points — MSE fit after adding two outliers

Ten typical observations lie near one linear trend.


Two unusual observations have large vertical errors.

The squared-error fit becomes much flatter. How can two points change the fitted line so strongly?

Robust regression asks for a noise model that allows occasional large residuals without letting them dominate the fit.

Tails are easier to compare on a log scale



For residual $r = y - \hat{y}$, compare unit-scale Gaussian, Laplace, and Student- t ; $\nu = 3$ is its degrees of freedom.



Tails are easier to compare on a log scale

For residual $r = y - \hat{y}$, compare unit-scale Gaussian, Laplace, and Student- t ; $\nu = 3$ is its degrees of freedom.

— Gaussian — Laplace — Student- t , $\nu = 3$



Smaller ν means heavier Student- t tails. Heavy tails make an occasional large residual less surprising.

A large residual is less surprising under heavy tails

D DERIVATION

Evaluate the three densities at the same residual, $r = 4$, using unit scale parameters.

Gaussian

density ≈ 0.000134

The model finds this residual extremely surprising.

Laplace

density ≈ 0.00916

More tail density makes the same residual less surprising.

Student- t , $\nu = 3$

density ≈ 0.00916

Heavy tails explicitly allow occasional extreme errors.

Laplace and Student- t assign about $68 \times$ the Gaussian density here.

The observation is not discarded; the assumed noise process simply says that such an error can occur.

One observation contributes one gradient signal

D DERIVATION

For observation i , the model predicts $\hat{y}_i$ and leaves residual $r_i = y_i - \hat{y}_i$.



Loss value ℓ_i

says how incompatible this observation is with the current prediction.

One observation contributes one gradient signal

D DERIVATION

For observation i , the model predicts $\hat{y}_i$ and leaves residual $r_i = y_i - \hat{y}_i$.



Loss value ℓ_i

says how incompatible this observation is with the current prediction.

Gradient magnitude $|\ell_{i'}(r_i)|$

says how strongly this observation scales the next parameter update.

One observation contributes one gradient signal

For observation i , the model predicts $\hat{y}_i$ and leaves residual $r_i = y_i - \hat{y}_i$.



Loss value ℓ_i

says how incompatible this observation is with the current prediction.

Gradient magnitude $|\ell_{i'}(r_i)|$

says how strongly this observation scales the next parameter update.

$$\nabla_{\theta} \ell_i = -\ell_{i'}(r_i) \nabla_{\theta} \hat{y}_i$$

The sign gives the correction direction; $|\ell_{i'}(r_i)|$ gives the size of this observation's gradient factor.

For one observation, $\hat{y}_i = f_{\theta}(x_i)$ and $r_i = y_i - \hat{y}_i$.



$$\ell(r) = -\log p_{R(r)} \quad \rightarrow \quad \psi(r) = \frac{d\ell}{dr} = -\frac{1}{p_{R(r)}} \frac{dp_{R(r)}}{dr}$$

For one observation, $\hat{y}_i = f_{\theta}(x_i)$ and $r_i = y_i - \hat{y}_i$.



$$\ell(r) = -\log p_{R(r)} \quad \rightarrow \quad \psi(r) = \frac{d\ell}{dr} = -\frac{1}{p_{R(r)}} \frac{dp_{R(r)}}{dr}$$

$$\nabla_{\theta} \ell_i = \psi(r_i) \nabla_{\theta} r_i = -\psi(r_i) \nabla_{\theta} f_{\theta}(x_i).$$

chain rule: residual-loss derivative $\times$ residual-to-model derivative

For one observation, $\hat{y}_i = f_{\theta}(x_i)$ and $r_i = y_i - \hat{y}_i$.



$$\ell(r) = -\log p_{R(r)} \quad \rightarrow \quad \psi(r) = \frac{d\ell}{dr} = -\frac{1}{p_{R(r)}} \frac{dp_{R(r)}}{dr}$$

$$\nabla_{\theta} \ell_i = \psi(r_i) \nabla_{\theta} r_i = -\psi(r_i) \nabla_{\theta} f_{\theta}(x_i).$$

chain rule: residual-loss derivative $\times$ residual-to-model derivative

$$\text{Linear example: } f_{\theta}(x) = \theta_0 + \theta_1 x \quad \rightarrow \quad \nabla_{\theta} \ell_i = -\psi(r_i)(1, x_i)^{\top}.$$

The density chooses $\psi(r)$; the model supplies $\nabla_{\theta} f$. Their product is the observation's parameter gradient.

Noise	Residual density	NLL $\ell(r)$	$\psi(r) = \ell'(r)$
	$p_{R(r)}$		
Gaussian	$\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{r^2}{2\sigma^2}}$	$\frac{r^2}{2\sigma^2} + C$	$\frac{r}{\sigma^2}$
Laplace	$\frac{1}{2s_L} e^{-\frac{ r }{s_L}}$	$\frac{ r }{s_L} + C$	$\frac{\text{sign}(r)}{s_L}$

Noise	Residual density	NLL $\ell(r)$	$\psi(r) = \ell'(r)$
	$p_R(r)$		
Gaussian	$\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{r^2}{2\sigma^2}}$	$\frac{r^2}{2\sigma^2} + C$	$\frac{r}{\sigma^2}$
Laplace	$\frac{1}{2s_L} e^{-\frac{ r }{s_L}}$	$\frac{ r }{s_L} + C$	$\frac{\text{sign}(r)}{s_L}$

$$\nabla_{\theta} \ell_i = -\psi(r_i) \nabla_{\theta} f_{\theta}(x_i)$$

At $r = 0$, $|r|$ has no single derivative; optimization uses any slope in $\left[-\frac{1}{s_L}, \frac{1}{s_L}\right]$, called a **subgradient**.
Away from zero, the gradient magnitude is $\frac{1}{s_L}$.

Gaussian influence grows with $|r|$; Laplace influence has constant magnitude.

Student- t residual: the full chain

D DERIVATION

For degrees of freedom ν and scale s :

ν controls tail heaviness; $s > 0$ sets residual scale; $c_{\nu,s} > 0$ normalizes the density to integrate to 1.

density $p_{R(r)} = c_{\nu,s} \left(1 + \frac{r^2}{\nu s^2}\right)^{-\frac{\nu+1}{2}}$

NLL $\ell(r) = \frac{\nu+1}{2} \log\left(1 + \frac{r^2}{\nu s^2}\right) + C$

r derivative $\psi(r) = \frac{d\ell}{dr} = (\nu + 1) \frac{r}{\nu s^2 + r^2}$

θ gradient $\nabla_{\theta} \ell_i = -(\nu + 1) \frac{r_i}{\nu s^2 + r_i^2} \nabla_{\theta} f_{\theta}(x_i)$

Student- t residual: the full chain

D DERIVATION

For degrees of freedom ν and scale s :


ν controls tail heaviness; $s > 0$ sets residual scale; $c_{\nu,s} > 0$ normalizes the density to integrate to 1.

$$\begin{aligned} \text{density} \quad p_{R(r)} &= c_{\nu,s} \left(1 + \frac{r^2}{\nu s^2}\right)^{-\frac{\nu+1}{2}} \\ \text{NLL} \quad \ell(r) &= \frac{\nu+1}{2} \log\left(1 + \frac{r^2}{\nu s^2}\right) + C \\ r \text{ derivative} \quad \psi(r) &= \frac{d\ell}{dr} = (\nu+1) \frac{r}{\nu s^2 + r^2} \\ \theta \text{ gradient} \quad \nabla_{\theta} \ell_i &= -(\nu+1) \frac{r_i}{\nu s^2 + r_i^2} \nabla_{\theta} f_{\theta}(x_i) \end{aligned}$$

Near zero, $\psi(r) \approx \left(\frac{\nu+1}{\nu s^2}\right)r$; for $|r|$ very large, $\psi(r) \approx \frac{\nu+1}{r} \rightarrow 0$.

Student- t treats small residuals approximately quadratically, but an extreme residual's gradient contribution eventually decreases.

Large residuals create different gradient magnitudes




— Gaussian — Laplace — Student- $t, \nu = 3$

At $r = 4$, the gradient factor is:

Gaussian

$$|\ell'(4)| = 4$$

Large residuals create different gradient magnitudes



— Gaussian — Laplace — Student- $t, \nu = 3$

At $r = 4$, the gradient factor is:


Gaussian

$$|\ell'(4)| = 4$$

Laplace

$$|\ell'(4)| = 1$$

Large residuals create different gradient magnitudes



At $r = 4$, the gradient factor is:

Gaussian

$$|\ell'(4)| = 4$$

Laplace

$$|\ell'(4)| = 1$$

Student- t

$$|\ell'(4)| = \frac{16}{19} \approx 0.84$$

Robust fitting keeps every observation, while limiting the gradient contribution of an extreme residual.

The noise model determines how errors are penalized

Noise model	Negative log-likelihood
— Gaussian	r^2
— Laplace	$ r $
— Uniform	0 inside; $+\infty$ outside
— Student- t	$\log\left(1 + \frac{r^2}{\nu}\right)$


$r = y - \hat{y}$ · swatches match plot

The noise model determines how errors are penalized

Noise model	Negative log-likelihood
— Gaussian	r^2
— Laplace	$ r $
— Uniform	0 inside; $+\infty$ outside
— Student- t	$\log\left(1 + \frac{r^2}{\nu}\right)$


$r = y - \hat{y}$ · swatches match plot

At $r = 0$, prediction equals target. Farther horizontally = larger error; higher vertically = larger NLL penalty.



The noise model determines how errors are penalized

Noise model	Negative log-likelihood
— Gaussian	r^2
— Laplace	$ r $
— Uniform	0 inside; $+\infty$ outside
— Student- t	$\log\left(1 + \frac{r^2}{\nu}\right)$




$r = y - \hat{y}$ · swatches match plot

At $r = 0$, prediction equals target. Farther horizontally = larger error; higher vertically = larger NLL penalty.

Gaussian penalizes large errors most; Student- t is more robust. The dashed Uniform boundaries mark an infinite penalty outside the allowed interval.

Fit the same data under three noise models




Learned $\hat{y} = \theta_0 + \theta_1 x$

model	fitted (θ_0, θ_1)
Gaussian	(-0.040, 0.397)
Laplace	(0.044, 0.906)
Student- t	(0.053, 0.882)

$$\sigma = 1 \cdot s_L = 1 \cdot (\nu, s) = (3, 1)$$

Fit the same data under three noise models



$$\text{Learned } \hat{y} = \theta_0 + \theta_1 x$$


model	fitted (θ_0, θ_1)
Gaussian	(-0.040, 0.397)
Laplace	(0.044, 0.906)
Student- t	(0.053, 0.882)

$$\sigma = 1 \cdot s_L = 1 \cdot (\nu, s) = (3, 1)$$

All twelve observations are retained. The models differ only in the residual likelihood optimized.

The Gaussian fit is pulled flat by the two red points; Laplace and Student- t stay close to the trend supported by the ten typical points.

Student- t hyperparameters control robustness




— $\nu = 1$ — $\nu = 3$ — $\nu = 30$ (all $s = 1$)

Learned coefficients

ν	s	fitted (θ_0, θ_1)
1	1	(0.064, 0.928)
3	0.5	(0.065, 0.934)
3	1	(0.053, 0.882)
3	1.5	(0.036, 0.807)
30	1	(-0.014, 0.581)

Smaller ν gives heavier tails. Larger s delays the transition from quadratic to tail-robust behaviour.

Student- t hyperparameters control robustness



— $\nu = 1$ — $\nu = 3$ — $\nu = 30$ (all $s = 1$)

Learned coefficients

ν	s	fitted (θ_0, θ_1)
1	1	(0.064, 0.928)
3	0.5	(0.065, 0.934)
3	1	(0.053, 0.882)
3	1.5	(0.036, 0.807)
30	1	(-0.014, 0.581)

Smaller ν gives heavier tails. Larger s delays the transition from quadratic to tail-robust behaviour.

For unregularized Gaussian and Laplace fits, changing the fixed scale only rescales the objective and leaves (θ_0, θ_1) unchanged; Student- t scale changes relative residual weights.

ν and s are modelling choices: cross-validation or domain knowledge should choose how quickly an observation is treated as a tail event.

Interactive: likelihood $\rightarrow$ loss

I INTERACTIVE

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/likelihood-to-loss

likelihood-to-loss.html — three synchronized panels: the **noise density** $p(r)$, the **loss** $-\log p(r)$, and **data with a fitted line**.

Controls: Gaussian / Laplace / Uniform / Student- t · noise scale · slope & intercept · **add an outlier** · per-point contributions.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/likelihood-to-loss

likelihood-to-loss.html — three synchronized panels: the **noise density** $p(r)$, the **loss** $-\log p(r)$, and **data with a fitted line**.

Controls: Gaussian / Laplace / Uniform / Student- t · noise scale · slope & intercept · **add an outlier** · per-point contributions.

Ask: which model reacts most to an outlier? why does Uniform give an infinite barrier? which loss is robust?

Classification: cross-entropy

Classification should output probabilities

For K classes:

$$p_{\theta}(Y = k \mid \mathbf{x}) \geq 0, \quad \sum_{k=1}^K p_{\theta}(Y = k \mid \mathbf{x}) = 1$$


A network first produces unconstrained class scores, called **logits**; sigmoid or softmax converts them to probabilities. We derive both below.

Classification should output probabilities

For K classes:

$$p_{\theta}(Y = k \mid \mathbf{x}) \geq 0, \quad \sum_{k=1}^K p_{\theta}(Y = k \mid \mathbf{x}) = 1$$

A network first produces unconstrained class scores, called **logits**; sigmoid or softmax converts them to probabilities. We derive both below.



f_{θ} may be linear or nonlinear; the probability constraints come from the output transformation.

Binary classification: Bernoulli model

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Bernoulli}(p_{\boldsymbol{\theta}}(\mathbf{x}))$$

Binary classification: Bernoulli model

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Bernoulli}(p_{\boldsymbol{\theta}}(\mathbf{x}))$$

Compact one-line form: Write $p = p_{\boldsymbol{\theta}}(\mathbf{x})$. Then

$$p_{\boldsymbol{\theta}}(y \mid \mathbf{x}) = p^y(1 - p)^{1-y}$$

Binary classification: Bernoulli model

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Bernoulli}(p_{\boldsymbol{\theta}}(\mathbf{x}))$$

Compact one-line form: Write $p = p_{\boldsymbol{\theta}}(\mathbf{x})$. Then

$$p_{\boldsymbol{\theta}}(y \mid \mathbf{x}) = p^y(1 - p)^{1-y}$$

$$y = 1 \Rightarrow p(y) = p$$

$$y = 0 \Rightarrow p(y) = 1 - p$$

$$-\log p_{\theta}(y \mid \boldsymbol{x}) = -\log(p^y(1-p)^{1-y})$$

Bernoulli NLL gives BCE

D DERIVATION

$$-\log p_{\theta}(y \mid \mathbf{x}) = -\log(p^y(1-p)^{1-y})$$

$$= -y \log p - (1-y) \log(1-p)$$

$$-\log p_{\theta}(y \mid \boldsymbol{x}) = -\log(p^y(1-p)^{1-y})$$

$$= -y \log p - (1-y) \log(1-p)$$

Binary cross-entropy = Bernoulli NLL.

$$\begin{aligned} -\log p_{\theta}(y \mid \mathbf{x}) &= -\log(p^y(1-p)^{1-y}) \\ &= -y \log p - (1-y) \log(1-p) \end{aligned}$$

Binary cross-entropy = Bernoulli NLL.

Check. True label $y = 1$: if $p = 0.8$ the loss is $-\log 0.8 \approx 0.22$; if $p = 0.2$ it is $-\log 0.2 \approx 1.61$.

Why BCE matches a binary label better than MSE

★ OPTIONAL

For one observed example, let the target be $y = 1$ and the model predict

$$p = P(Y = 1|\mathbf{x}) = 0.8.$$

Why BCE matches a binary label better than MSE

★ OPTIONAL

For one observed example, let the target be $y = 1$ and the model predict

$$p = P(Y = 1|\mathbf{x}) = 0.8.$$

Squared error

$$\ell_{\text{MSE}} = (y - p)^2 = (1 - 0.8)^2 = 0.04.$$

Its usual likelihood interpretation is Gaussian:

$$Y|\mathbf{x} \sim \mathcal{N}(0.8, \sigma^2).$$

Permits all real Y —not just labels.

Why BCE matches a binary label better than MSE

★ OPTIONAL

For one observed example, let the target be $y = 1$ and the model predict

$$p = P(Y = 1|\mathbf{x}) = 0.8.$$

Squared error

$$\ell_{\text{MSE}} = (y - p)^2 = (1 - 0.8)^2 = 0.04.$$

Its usual likelihood interpretation is Gaussian:

$$Y|\mathbf{x} \sim \mathcal{N}(0.8, \sigma^2).$$

Permits all real Y —not just labels.

Binary cross-entropy

$$\ell_{\text{BCE}} = -\log p = -\log 0.8 \approx 0.223.$$

Its likelihood is Bernoulli:

$$Y|\mathbf{x} \sim \text{Bernoulli}(0.8).$$

Permits only $Y = 0$ or $Y = 1$.

Why BCE matches a binary label better than MSE

★ OPTIONAL

For one observed example, let the target be $y = 1$ and the model predict

$$p = P(Y = 1|\mathbf{x}) = 0.8.$$

Squared error

$$\ell_{\text{MSE}} = (y - p)^2 = (1 - 0.8)^2 = 0.04.$$

Its usual likelihood interpretation is Gaussian:

$$Y|\mathbf{x} \sim \mathcal{N}(0.8, \sigma^2).$$

Permits all real Y —not just labels.

Binary cross-entropy

$$\ell_{\text{BCE}} = -\log p = -\log 0.8 \approx 0.223.$$

Its likelihood is Bernoulli:

$$Y|\mathbf{x} \sim \text{Bernoulli}(0.8).$$

Permits only $Y = 0$ or $Y = 1$.

MSE is possible; BCE is the negative log-likelihood matched to a binary observation.

Why use logits?

The network emits an unbounded score $z = f_{\theta}(x) \in \mathbb{R}$; the **sigmoid** maps it to a probability:

$$p = \sigma(z) = \frac{1}{1 + e^{-z}}$$

Why use logits?

The network emits an unbounded score $z = f_{\theta}(x) \in \mathbb{R}$; the **sigmoid** maps it to a probability:


$$p = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z \rightarrow -\infty : p \rightarrow 0$$

$$z = 0 : p = 0.5$$

$$z \rightarrow +\infty : p \rightarrow 1$$

In practice sigmoid + BCE are fused for numerical stability.



Logistic regression = Bernoulli MLE

D DERIVATION

Use an augmented input $\tilde{\mathbf{x}}_i$ so $\boldsymbol{\theta}$ contains both weights and the intercept:

$$p_i = \sigma(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i), \quad Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \text{Bernoulli}(p_i)$$

.

Logistic regression = Bernoulli MLE

D DERIVATION

Use an augmented input $\tilde{\mathbf{x}}_i$ so $\boldsymbol{\theta}$ contains both weights and the intercept:

$$p_i = \sigma(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i), \quad Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \text{Bernoulli}(p_i)$$

. Likelihood of the whole labelled dataset (i.i.d.):

$$L(\boldsymbol{\theta}) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

Logistic regression = Bernoulli MLE

Use an augmented input $\tilde{\mathbf{x}}_i$ so $\boldsymbol{\theta}$ contains both weights and the intercept:

$$p_i = \sigma(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i), \quad Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \text{Bernoulli}(p_i)$$

. Likelihood of the whole labelled dataset (i.i.d.):

$$L(\boldsymbol{\theta}) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

Its negative log — the objective we minimise:

$$J(\boldsymbol{\theta}) = - \sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)]$$

Logistic regression = Bernoulli MLE

D DERIVATION

Use an augmented input $\tilde{\mathbf{x}}_i$ so $\boldsymbol{\theta}$ contains both weights and the intercept:

$$p_i = \sigma(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i), \quad Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \text{Bernoulli}(p_i)$$

. Likelihood of the whole labelled dataset (i.i.d.):

$$L(\boldsymbol{\theta}) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

Its negative log — the objective we minimise:

$$J(\boldsymbol{\theta}) = - \sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)]$$

Logistic regression is Bernoulli maximum likelihood; its training objective is **binary cross-entropy**.

CATEGORICAL OBSERVATION

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Categorical}(p_1, \dots, p_K)$$

A **one-hot** target has one 1 at the true class and 0 everywhere else.

Example: class 2 of 3 $\rightarrow \mathbf{y} = (0, 1, 0)$

$$p_{\theta}(\mathbf{y} \mid \mathbf{x}) = \prod_{k=1}^K p_k^{y_k}$$

CATEGORICAL OBSERVATION

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Categorical}(p_1, \dots, p_K)$$

A **one-hot** target has one 1 at the true class and 0 everywhere else.


Example: class 2 of 3 $\rightarrow \mathbf{y} = (0, 1, 0)$

$$p_{\theta}(\mathbf{y} \mid \mathbf{x}) = \prod_{k=1}^K p_k^{y_k}$$

NETWORK PARAMETERIZATION

$$\mathbf{z} = f_{\theta}(\mathbf{x}), \quad p_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

softmax probabilities



CATEGORICAL OBSERVATION

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Categorical}(p_1, \dots, p_K)$$

A **one-hot** target has one 1 at the true class and 0 everywhere else.


Example: class 2 of 3 $\rightarrow \mathbf{y} = (0, 1, 0)$

$$p_{\theta}(\mathbf{y} \mid \mathbf{x}) = \prod_{k=1}^K p_k^{y_k}$$

NETWORK PARAMETERIZATION

$$\mathbf{z} = f_{\theta}(\mathbf{x}), \quad p_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

softmax probabilities



Softmax guarantees $p_k > 0$ and $\sum_k p_k = 1$, as a categorical model requires.

Categorical NLL gives cross-entropy

D DERIVATION

$$-\log p_{\theta}(\mathbf{y} \mid \mathbf{x}) = -\log \prod_k p_k^{y_k} = -\sum_k y_k \log p_k$$

Categorical NLL gives cross-entropy

D DERIVATION

$$-\log p_{\theta}(\mathbf{y} \mid \mathbf{x}) = -\log \prod_k p_k^{y_k} = -\sum_k y_k \log p_k$$

Because $\mathbf{y}$ is one-hot, only the true class survives:

$$\mathcal{L} = -\log p_{\text{true class}}$$

Confident mistakes are expensive

$$\ell(p_y) = -\log p_y$$

Confident mistakes are expensive

$$\ell(p_y) = -\log p_y$$

Read off the loss at three confidences for the **true** class:

Confident mistakes are expensive


$$\ell(p_y) = -\log p_y$$

Read off the loss at three confidences for the **true** class:

$$p_y = 0.9 \Rightarrow \ell \approx 0.105$$

$$p_y = 0.1 \Rightarrow \ell \approx 2.303$$

$$p_y = 0.001 \Rightarrow \ell \approx 6.908$$



Confident mistakes are expensive


$$\ell(p_y) = -\log p_y$$

Read off the loss at three confidences for the **true** class:

$$p_y = 0.9 \Rightarrow \ell \approx 0.105$$

$$p_y = 0.1 \Rightarrow \ell \approx 2.303$$

$$p_y = 0.001 \Rightarrow \ell \approx 6.908$$



The negative log makes a confident wrong prediction a large loss.

Why does cross-entropy penalize a confidently wrong prediction strongly?

A. It makes all class probabilities equal

B. The true-class probability is near zero, so $-\log p_y$ is large

C. It adds an L_2 penalty to the logits

D. The gradient is always zero

Correct: B — The true-class probability is near zero

Why: The negative log grows without bound as p_y approaches zero, matching the fact that the model declared the observed class nearly impossible.

Bayes' rule for learning

PROBABILITY LECTURE

$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

A – event or hypothesis to infer

B – observed event or information

Examples: disease given a test; urn given a sampled colour.

PROBABILITY LECTURE

$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

A – event or hypothesis to infer

B – observed event or information

Examples: disease given a test; urn given a sampled colour.

MACHINE / DEEP LEARNING

$$p(\boldsymbol{\theta} \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \boldsymbol{\theta})p(\boldsymbol{\theta})}{p(\mathcal{D})}$$

$p(\boldsymbol{\theta})$ – prior before seeing the dataset

$p(\mathcal{D} \mid \boldsymbol{\theta})$ – likelihood of the fixed dataset

$p(\mathcal{D})$ – evidence (normalizing constant)

$p(\boldsymbol{\theta} \mid \mathcal{D})$ – posterior after seeing the dataset

Question: which parameter settings remain plausible after seeing the data?

PROBABILITY LECTURE

$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

A – event or hypothesis to infer

B – observed event or information

Examples: disease given a test; urn given a sampled colour.

MACHINE / DEEP LEARNING

$$p(\boldsymbol{\theta} \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \boldsymbol{\theta})p(\boldsymbol{\theta})}{p(\mathcal{D})}$$

$p(\boldsymbol{\theta})$ – prior before seeing the dataset

$p(\mathcal{D} \mid \boldsymbol{\theta})$ – likelihood of the fixed dataset

$p(\mathcal{D})$ – evidence (normalizing constant)

$p(\boldsymbol{\theta} \mid \mathcal{D})$ – posterior after seeing the dataset

Question: which parameter settings remain plausible after seeing the data?

Replace A by $\boldsymbol{\theta}$ and B by $\mathcal{D}$: the algebra is unchanged; the interpretation changes.

Three heads do not prove $\theta = 1$

Q QUESTION

Let $\theta = P(H)$ and suppose the first three flips are $\mathcal{D} = (H, H, H)$.

Three heads do not prove $\theta = 1$


QUESTION

Let $\theta = P(H)$ and suppose the first three flips are $\mathcal{D} = (H, H, H)$.

H · H · H

$$L(\theta; \mathcal{D}) = p(\mathcal{D} \mid \theta) = \theta^3$$

$$L(\theta) = \theta^3$$




Three heads do not prove $\theta = 1$

QUESTION

Let $\theta = P(H)$ and suppose the first three flips are $\mathcal{D} = (H, H, H)$.

H • H • H

$$L(\theta) = \theta^3$$



$$L(\theta; \mathcal{D}) = p(\mathcal{D} \mid \theta) = \theta^3$$

$$\hat{\theta}_{\text{MLE}} = \arg \max_{0 \leq \theta \leq 1} \theta^3 = 1$$

Did we learn that the coin **always** lands heads—or did a merely head-biased coin produce a lucky run?


Three heads do not prove $\theta = 1$

QUESTION

Let $\theta = P(H)$ and suppose the first three flips are $\mathcal{D} = (H, H, H)$.

H • H • H

$$L(\theta) = \theta^3$$



$$L(\theta; \mathcal{D}) = p(\mathcal{D} \mid \theta) = \theta^3$$

$$\hat{\theta}_{\text{MLE}} = \arg \max_{0 \leq \theta \leq 1} \theta^3 = 1$$

Did we learn that the coin **always** lands heads—or did a merely head-biased coin produce a lucky run?

MLE identifies the best-fitting θ ; it does **not make $\theta = 1$ certain.**

A prior tempers the conclusion from three flips

Choose a smooth prior on $\theta \in [0, 1]$, highest near the fair-coin value 0.5.

Prior

$$p(\theta) = 6\theta(1 - \theta)$$

A prior tempers the conclusion from three flips

Choose a smooth prior on $\theta \in [0, 1]$, highest near the fair-coin value 0.5.

Prior

$$p(\theta) = 6\theta(1 - \theta)$$

Likelihood of H, H, H

$$p(\mathcal{D} \mid \theta) = \theta^3$$

A prior tempers the conclusion from three flips

Choose a smooth prior on $\theta \in [0, 1]$, highest near the fair-coin value 0.5.

Prior

$$p(\theta) = 6\theta(1 - \theta)$$

Likelihood of H, H, H


$$p(\mathcal{D} \mid \theta) = \theta^3$$

Posterior

$$p(\theta \mid \mathcal{D}) \propto \theta^3 p(\theta) \propto \theta^4(1 - \theta)$$

Likelihood peak (MLE): $\theta = 1$

Posterior peak: $\theta = \frac{4}{5} = 0.8$



A prior tempers the conclusion from three flips

Choose a smooth prior on $\theta \in [0, 1]$, highest near the fair-coin value 0.5.

Prior

$$p(\theta) = 6\theta(1 - \theta)$$

Likelihood of H, H, H


$$p(\mathcal{D} \mid \theta) = \theta^3$$

Posterior

$$p(\theta \mid \mathcal{D}) \propto \theta^3 p(\theta) \propto \theta^4(1 - \theta)$$

Likelihood peak (MLE): $\theta = 1$

Posterior peak: $\theta = \frac{4}{5} = 0.8$



With little data, the prior matters; as observations accumulate, the likelihood increasingly dominates.

Step 1: fixed data give fixed likelihoods

Suppose we consider only two possible coins.

Fair coin F

$$P(H|F) = 0.5$$

Heads-biased coin B

$$P(H|B) = 0.9$$

Step 1: fixed data give fixed likelihoods

D DERIVATION

Suppose we consider only two possible coins.

Fair coin F

$$P(H|F) = 0.5$$

Heads-biased coin B

$$P(H|B) = 0.9$$

Observed data: H · H · H

Step 1: fixed data give fixed likelihoods

D DERIVATION

Suppose we consider only two possible coins.

Fair coin F

$$P(H|F) = 0.5$$

Heads-biased coin B

$$P(H|B) = 0.9$$

Observed data: $H \cdot H \cdot H$

$$P(H, H, H|F) = 0.5^3 = 0.125$$

$$P(H, H, H|B) = 0.9^3 = 0.729$$

$$\text{likelihood ratio} = \frac{0.729}{0.125} \approx 5.83$$

HHH favours B over F by a factor of 5.83. These likelihoods do not depend on the prior.

Step 2: update a prior that favours the fair coin

Before seeing the flips, suppose our starting belief is $P(F) = 0.90$ and $P(B) = 0.10$.

Hypothesis	Prior	Likelihood	Prior $\times$ likelihood
Fair F	0.90	0.125	$0.90 \times$ $0.125 =$ 0.1125
Biased B	0.10	0.729	$0.10 \times$ $0.729 =$ 0.0729

Step 2: update a prior that favours the fair coin

D DERIVATION

Before seeing the flips, suppose our starting belief is $P(F) = 0.90$ and $P(B) = 0.10$.

Hypothesis	Prior	Likelihood	Prior $\times$ likelihood
Fair F	0.90	0.125	$0.90 \times$ $0.125 =$ 0.1125
Biased B	0.10	0.729	$0.10 \times$ $0.729 =$ 0.0729

$w_F = 0.1125$ and $w_B = 0.0729$ are **weights**, not probabilities yet.

Bayes first scores each hypothesis by prior $\times$ likelihood. The two scores need not add to 1.

Step 3: normalize the weights into probabilities

D DERIVATION

The normalizing constant is the sum of the two weights:

$$P(\mathcal{D}) = w_F + w_B = 0.1125 + 0.0729 = 0.1854.$$

$P(\mathcal{D})$ is the **evidence** (also called marginal likelihood); it makes posterior probabilities sum to 1.

Step 3: normalize the weights into probabilities

D DERIVATION

The normalizing constant is the sum of the two weights:

$$P(\mathcal{D}) = w_F + w_B = 0.1125 + 0.0729 = 0.1854.$$

$P(\mathcal{D})$ is the **evidence** (also called marginal likelihood); it makes posterior probabilities sum to 1.

Fair coin

$$P(F|\mathcal{D}) = \frac{0.1125}{0.1854} \approx 0.607$$

Heads-biased coin

$$P(B|\mathcal{D}) = \frac{0.0729}{0.1854} \approx 0.393$$

Step 3: normalize the weights into probabilities

D DERIVATION

The normalizing constant is the sum of the two weights:

$$P(\mathcal{D}) = w_F + w_B = 0.1125 + 0.0729 = 0.1854.$$

$P(\mathcal{D})$ is the **evidence** (also called marginal likelihood); it makes posterior probabilities sum to 1.

Fair coin

$$P(F|\mathcal{D}) = \frac{0.1125}{0.1854} \approx 0.607$$

Heads-biased coin

$$P(B|\mathcal{D}) = \frac{0.0729}{0.1854} \approx 0.393$$

The normalized values now add to 1: $0.607 + 0.393 = 1$. They are the **posterior** probabilities—our beliefs after observing HHH.

HHH raises belief in B from 0.10 to 0.393, but it does not overturn the strong fair-coin prior.

Step 4: start instead from a biased-coin prior

D DERIVATION

Now use the **same hypotheses and the same HHH**, but start from $P(F) = 0.10$ and $P(B) = 0.90$.

Hypothesis	Prior	Likelihood	Weight
Fair F	0.10	0.125	0.0125
Biased B	0.90	0.729	0.6561

Step 4: start instead from a biased-coin prior

D DERIVATION

Now use the **same hypotheses and the same HHH**, but start from $P(F) = 0.10$ and $P(B) = 0.90$.

Hypothesis	Prior	Likelihood	Weight
Fair F	0.10	0.125	0.0125
Biased B	0.90	0.729	0.6561

$$P(\mathcal{D}) = 0.0125 + 0.6561 = 0.6686$$

$$P(F|\mathcal{D}) \approx 0.019, \quad P(B|\mathcal{D}) \approx 0.981$$

The data again favours B by the same factor 5.83; this time the prior already favoured B .

What changed—and what stayed fixed?

Starting belief	Prior $P(B)$	Posterior $P(B H, H, H)$
Probably fair	0.10	0.393
Probably biased	0.90	0.981

What changed—and what stayed fixed?

Starting belief	Prior $P(B)$	Posterior $P(B H, H, H)$
Probably fair	0.10	0.393
Probably biased	0.90	0.981

Fixed: the observations and likelihoods. In both analyses, HHH multiplies the odds for B by 5.83.

Changed: the prior. Different starting beliefs therefore lead to different posterior probabilities.

Odds for B versus F mean $\frac{P(B)}{P(F)}$.

posterior odds = prior odds $\times$ likelihood ratio

**Bayes does not erase the prior; it updates it with the same observations.
More data makes the likelihood increasingly dominant.**

Maximum a posteriori (MAP) estimation

D DERIVATION

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta \mid \mathcal{D}) = \arg \max_{\theta} p(\mathcal{D} \mid \theta) p(\theta)$$

Maximum a posteriori (MAP) estimation

D DERIVATION

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta \mid \mathcal{D}) = \arg \max_{\theta} p(\mathcal{D} \mid \theta) p(\theta)$$

The evidence $p(\mathcal{D})$ is constant with respect to θ , so it does not change the maximizing parameter.

Maximum a posteriori (MAP) estimation

D DERIVATION

$$\hat{\boldsymbol{\theta}}_{\text{MAP}} = \arg \max_{\boldsymbol{\theta}} p(\boldsymbol{\theta} \mid \mathcal{D}) = \arg \max_{\boldsymbol{\theta}} p(\mathcal{D} \mid \boldsymbol{\theta}) p(\boldsymbol{\theta})$$

The evidence $p(\mathcal{D})$ is constant with respect to $\boldsymbol{\theta}$, so it does not change the maximizing parameter. Take negative logs:

$$\hat{\boldsymbol{\theta}}_{\text{MAP}} = \arg \min_{\boldsymbol{\theta}} \left(\underbrace{-\log p(\mathcal{D} \mid \boldsymbol{\theta})}_{\text{NLL}} + \underbrace{-\log p(\boldsymbol{\theta})}_{\text{negative log-prior}} \right)$$

Maximum a posteriori (MAP) estimation

D DERIVATION

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta \mid \mathcal{D}) = \arg \max_{\theta} p(\mathcal{D} \mid \theta) p(\theta)$$

The evidence $p(\mathcal{D})$ is constant with respect to θ , so it does not change the maximizing parameter. Take negative logs:

$$\hat{\theta}_{\text{MAP}} = \arg \min_{\theta} \left(\underbrace{-\log p(\mathcal{D} \mid \theta)}_{\text{NLL}} + \underbrace{-\log p(\theta)}_{\text{negative log-prior}} \right)$$

MAP objective: NLL + negative log-prior. In ML language: data loss + regularization.

Priors, MAP and regularization

Every learner has an inductive bias

V VISUAL

When several predictors fit the training data, the learner still needs a preference for **unseen inputs**. That preference is its **inductive bias**.

Every learner has an inductive bias

When several predictors fit the training data, the learner still needs a preference for **unseen inputs**. That preference is its **inductive bias**.

Model	Built-in preference → consequence
Linear regression	responses vary linearly with features → line or plane extrapolation
k-NN	nearby inputs tend to have similar outputs → local prediction
Decision tree	a few axis-aligned rules partition the space → piecewise-constant regions
CNN	useful local patterns repeat across positions → shared filters recognize translations

Every learner has an inductive bias

When several predictors fit the training data, the learner still needs a preference for **unseen inputs**. That preference is its **inductive bias**.

Model	Built-in preference → consequence
Linear regression	responses vary linearly with features → line or plane extrapolation
k-NN	nearby inputs tend to have similar outputs → local prediction
Decision tree	a few axis-aligned rules partition the space → piecewise-constant regions
CNN	useful local patterns repeat across positions → shared filters recognize translations

A prior is one explicit inductive bias; the model class, features, architecture, and optimization add others.


A neural-network prior ranks plausible parameter vectors

$$\theta \sim p(\theta)$$

For a neural network, θ contains all weights and biases. Before using the current dataset, a prior states which settings are more plausible.

It can prefer small or sparse weights (**sparse** = many weights at or near zero), smooth functions, and correlated parameters.

— narrow $\sigma = 0.55$ — broad $\sigma = 1.55$



narrow = stronger preference near 0


A neural-network prior ranks plausible parameter vectors

$$\theta \sim p(\theta)$$

For a neural network, θ contains all weights and biases. Before using the current dataset, a prior states which settings are more plausible.

It can prefer small or sparse weights (**sparse** = many weights at or near zero), smooth functions, and correlated parameters.

— narrow $\sigma = 0.55$ — broad $\sigma = 1.55$



narrow = stronger preference near 0


A prior makes part of the model's inductive bias explicit as a probability distribution.

Parameter priors have visible consequences

Example: a zero-centred prior makes parameter values near $\theta = \mathbf{0}$ more plausible before seeing data.

REGRESSION

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i, \sigma^2)$$




Parameter priors have visible consequences

Example: a zero-centred prior makes parameter values near $\theta = \mathbf{0}$ more plausible before seeing data.


REGRESSION

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(\theta^\top \tilde{\mathbf{x}}_i, \sigma^2)$$



BINARY CLASSIFICATION

$$p_{\theta}(Y = 1 \mid \mathbf{x}) = \sigma(\theta^\top \tilde{\mathbf{x}})$$




Parameter priors have visible consequences

Example: a zero-centred prior makes parameter values near $\theta = 0$ more plausible before seeing data.


REGRESSION

$$Y_i \mid X_i = x_i \sim \mathcal{N}(\theta^\top \tilde{x}_i, \sigma^2)$$



BINARY CLASSIFICATION

$$p_\theta(Y = 1 \mid x) = \sigma(\theta^\top \tilde{x})$$




This conclusion depends on the prior: a prior centred elsewhere can prefer probabilities away from 0.5.

Model complexity: underfitting to overfitting

Fit the same eight observations by least squares; only the polynomial degree changes.


dashed = underlying trend · orange = fitted polynomial · red dot = one noisy observation



Model complexity: underfitting to overfitting

Fit the same eight observations by least squares; only the polynomial degree changes.

dashed = underlying trend · orange = fitted polynomial · red dot = one noisy observation



underfit




captures the main trend

Model complexity: underfitting to overfitting

Fit the same eight observations by least squares; only the polynomial degree changes.


dashed = underlying trend · orange = fitted polynomial · red dot = one noisy observation



underfit



captures the main trend




overfit


Model complexity: underfitting to overfitting

Fit the same eight observations by least squares; only the polynomial degree changes.


dashed = underlying trend · orange = fitted polynomial · red dot = one noisy observation



underfit



captures the main trend



overfit

Higher degree lowers training error, but degree 7 chases the noisy observation and becomes unstable. Priors and regularization can prefer smoother solutions.

Start with four noisy points from a straight line


The data were generated from $f_{\text{true}(x)} = 2 + 3x$ with small measurement noise.

x	$f_{\text{true}(x)}$	observed y
-1	-1	-0.8
0	2	1.8
1	5	5.3
2	8	7.9

Start with four noisy points from a straight line

The data were generated from $f_{\text{true}(x)} = 2 + 3x$ with small measurement noise.

x	$f_{\text{true}(x)}$	observed y
-1	-1	-0.8
0	2	1.8
1	5	5.3
2	8	7.9




grey line = generating rule · dots = observed data

Start with four noisy points from a straight line

The data were generated from $f_{\text{true}(x)} = 2 + 3x$ with small measurement noise.

x	$f_{\text{true}(x)}$	observed y
-1	-1	-0.8
0	2	1.8
1	5	5.3
2	8	7.9



grey line = generating rule · dots = observed data

The learner sees only the four observations and must infer the intercept b and slope w .

The likelihood peaks near the generating line

Model each observation as

$$Y_i | x_i, b, w \sim \mathcal{N}(b + wx_i, 0.5^2).$$

The data score each candidate line

$$L(b, w) \propto \exp\left(-\frac{1}{2(0.5)^2} \sum_i (y_i - b - wx_i)^2\right).$$

candidate (b, w)	squared error
(2, 3)	0.180
(0, 1)	56.580
MLE (2.07, 2.96)	0.162

The likelihood peaks near the generating line

D DERIVATION


Model each observation as

$$Y_i|x_i, b, w \sim \mathcal{N}(b + wx_i, 0.5^2).$$

The data score each candidate line

$$L(b, w) \propto \exp\left(-\frac{1}{2(0.5)^2} \sum_i (y_i - b - wx_i)^2\right).$$

candidate (b, w)	squared error
(2, 3)	0.180
(0, 1)	56.580
MLE (2.07, 2.96)	0.162



The likelihood peaks near the generating line

D DERIVATION


Model each observation as

$$Y_i|x_i, b, w \sim \mathcal{N}(b + wx_i, 0.5^2).$$

The data score each candidate line

$$L(b, w) \propto \exp\left(-\frac{1}{2(0.5)^2} \sum_i (y_i - b - wx_i)^2\right).$$

candidate (b, w)	squared error
(2, 3)	0.180
(0, 1)	56.580
MLE (2.07, 2.96)	0.162



The observed data make intercepts near 2 and slopes near 3 most plausible.

Before data, the prior prefers parameters near zero

Choose a simple Gaussian prior over the two regression parameters:

$$\boldsymbol{\theta} = (b, w) \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}), \quad \text{here } \tau = 0.5.$$


τ is the prior standard deviation; $\mathbf{I}$ makes b and w independent with equal scale.

Before data, the prior prefers parameters near zero

Choose a simple Gaussian prior over the two regression parameters:

$$\boldsymbol{\theta} = (b, w) \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}), \quad \text{here } \tau = 0.5.$$

τ is the prior standard deviation; $\mathbf{I}$ makes b and w independent with equal scale.



What does it express?

$$-\log p(b, w) = \frac{1}{2\tau^2} (b^2 + w^2) + C = 2(b^2 + w^2) + C \quad \text{when } \tau = 0.5.$$

Before seeing this dataset:


- $b \approx 0$ and $w \approx 0$: small-intercept, flatter lines
- farther from $(0, 0)$: less prior density

Before data, the prior prefers parameters near zero

Choose a simple Gaussian prior over the two regression parameters:

$$\theta = (b, w) \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}), \quad \text{here } \tau = 0.5.$$

τ is the prior standard deviation; $\mathbf{I}$ makes b and w independent with equal scale.



What does it express?

$$-\log p(b, w) = \frac{1}{2\tau^2} (b^2 + w^2) + C = 2(b^2 + w^2) + C \quad \text{when } \tau = 0.5.$$

Before seeing this dataset:

- $b \approx 0$ and $w \approx 0$: small-intercept, flatter lines
- farther from $(0, 0)$: less prior density


The likelihood says “fit these data”; the prior says “prefer smaller coefficients”.

Bayes combines the regression likelihood and prior

$$p(b, w \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid b, w) p(b, w)}{p(\mathcal{D})}$$

Bayes combines the regression likelihood and prior

$$p(b, w \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid b, w) p(b, w)}{p(\mathcal{D})}$$



Contours: — prior — likelihood — posterior

● MLE (2.07, 2.96) ● posterior peak (1.79, 2.62) ● prior mean (0, 0)


The narrower prior visibly pulls the posterior peak away from the MLE and toward zero.

The prior changes the fitted line

— MLE, dashed — MAP, solid

Moderate prior: $\tau = 0.5$

MAP: $\hat{y} = 1.79 + 2.62x$




The prior changes the fitted line

— MLE, dashed — MAP, solid


Moderate prior: $\tau = 0.5$

$$\text{MAP: } \hat{y} = 1.79 + 2.62x$$



Much stronger prior: $\tau = 0.1$

$$\text{MAP: } \hat{y} = 0.44 + 0.68x$$




$$(b, w) : (2.07, 2.96)_{\text{MLE}} \rightarrow (1.79, 2.62)_{\text{moderate}} \rightarrow (0.44, 0.68)_{\text{strong}}$$

The prior changes the fitted line

— MLE, dashed — MAP, solid


Moderate prior: $\tau = 0.5$

$$\text{MAP: } \hat{y} = 1.79 + 2.62x$$



Much stronger prior: $\tau = 0.1$

$$\text{MAP: } \hat{y} = 0.44 + 0.68x$$



$$(b, w) : (2.07, 2.96)_{\text{MLE}} \rightarrow (1.79, 2.62)_{\text{moderate}} \rightarrow (0.44, 0.68)_{\text{strong}}$$

Reducing τ puts more prior mass near the origin and moves both MAP coefficients much closer to zero.

A Gaussian prior is L_2 regularization

Assume the parameters are centred at zero before seeing the data:

$$\boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}).$$

$$\underbrace{-\log p(\mathcal{D} \mid \boldsymbol{\theta})}_{\text{NLL}} + \underbrace{-\log p(\boldsymbol{\theta})}_{\text{negative log-prior}}$$

A Gaussian prior is L_2 regularization

Assume the parameters are centred at zero before seeing the data:

$$\boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}).$$

$$\underbrace{-\log p(\mathcal{D} \mid \boldsymbol{\theta})}_{\text{NLL}} + \underbrace{-\log p(\boldsymbol{\theta})}_{\text{negative log-prior}}$$

$$-\log p(\boldsymbol{\theta}) = \frac{1}{2\tau^2} \|\boldsymbol{\theta}\|_2^2 + C$$

$\|\boldsymbol{\theta}\|_2^2 = \sum_j \theta_j^2$ is the sum of squared coefficients.

A Gaussian prior is L_2 regularization

D DERIVATION

Assume the parameters are centred at zero before seeing the data:

$$\boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}).$$

$$\underbrace{-\log p(\mathcal{D} \mid \boldsymbol{\theta})}_{\text{NLL}} + \underbrace{-\log p(\boldsymbol{\theta})}_{\text{negative log-prior}}$$

$$-\log p(\boldsymbol{\theta}) = \frac{1}{2\tau^2} \|\boldsymbol{\theta}\|_2^2 + C$$

$\|\boldsymbol{\theta}\|_2^2 = \sum_j \theta_j^2$ is the sum of squared coefficients.

$$\Rightarrow \mathcal{L}_{\text{MAP}} = \text{NLL} + \underbrace{\frac{1}{2\tau^2} \|\boldsymbol{\theta}\|_2^2}_{\lambda} + C$$

$\lambda = \frac{1}{2\tau^2}$: **smaller** $\tau \rightarrow$ **larger** $\lambda \rightarrow$ **stronger shrinkage toward 0.**

Gaussian MAP gives ridge shrinkage, step by step

D DERIVATION

For one coefficient, let $z = \hat{\theta}_{\text{MLE}}$ be the likelihood centre and s its width:

$$p(\mathcal{D}|\theta) \propto \exp\left(-\frac{(\theta-z)^2}{2s^2}\right), \quad p(\theta) \propto \exp\left(-\frac{\theta^2}{2\tau^2}\right).$$

Gaussian MAP gives ridge shrinkage, step by step

D DERIVATION

For one coefficient, let $z = \hat{\theta}_{\text{MLE}}$ be the likelihood centre and s its width:

$$p(\mathcal{D}|\theta) \propto \exp\left(-\frac{(\theta-z)^2}{2s^2}\right), \quad p(\theta) \propto \exp\left(-\frac{\theta^2}{2\tau^2}\right).$$

$$J_2(\theta) = \underbrace{\frac{(\theta-z)^2}{2s^2}}_{\text{data NLL}} + \underbrace{\frac{\theta^2}{2\tau^2}}_{\text{Gaussian prior NLL}}.$$

Gaussian MAP gives ridge shrinkage, step by step

For one coefficient, let $z = \hat{\theta}_{\text{MLE}}$ be the likelihood centre and s its width:

$$p(\mathcal{D}|\theta) \propto \exp\left(-\frac{(\theta-z)^2}{2s^2}\right), \quad p(\theta) \propto \exp\left(-\frac{\theta^2}{2\tau^2}\right).$$

$$J_2(\theta) = \underbrace{\frac{(\theta-z)^2}{2s^2}}_{\text{data NLL}} + \underbrace{\frac{\theta^2}{2\tau^2}}_{\text{Gaussian prior NLL}}.$$

$$\frac{dJ_2}{d\theta} = \frac{\theta - z}{s^2} + \frac{\theta}{\tau^2} = 0$$

Gaussian MAP gives ridge shrinkage, step by step

For one coefficient, let $z = \hat{\theta}_{\text{MLE}}$ be the likelihood centre and s its width:

$$p(\mathcal{D}|\theta) \propto \exp\left(-\frac{(\theta-z)^2}{2s^2}\right), \quad p(\theta) \propto \exp\left(-\frac{\theta^2}{2\tau^2}\right).$$

$$J_2(\theta) = \underbrace{\frac{(\theta-z)^2}{2s^2}}_{\text{data NLL}} + \underbrace{\frac{\theta^2}{2\tau^2}}_{\text{Gaussian prior NLL}}.$$

$$\frac{dJ_2}{d\theta} = \frac{\theta - z}{s^2} + \frac{\theta}{\tau^2} = 0$$

$$\tau^2(\theta - z) + s^2\theta = 0 \quad \rightarrow \quad (\tau^2 + s^2)\theta = \tau^2 z$$

Gaussian MAP gives ridge shrinkage, step by step

D DERIVATION

For one coefficient, let $z = \hat{\theta}_{\text{MLE}}$ be the likelihood centre and s its width:

$$p(\mathcal{D}|\theta) \propto \exp\left(-\frac{(\theta-z)^2}{2s^2}\right), \quad p(\theta) \propto \exp\left(-\frac{\theta^2}{2\tau^2}\right).$$

$$J_2(\theta) = \underbrace{\frac{(\theta-z)^2}{2s^2}}_{\text{data NLL}} + \underbrace{\frac{\theta^2}{2\tau^2}}_{\text{Gaussian prior NLL}}.$$

$$\frac{dJ_2}{d\theta} = \frac{\theta - z}{s^2} + \frac{\theta}{\tau^2} = 0$$

$$\tau^2(\theta - z) + s^2\theta = 0 \quad \rightarrow \quad (\tau^2 + s^2)\theta = \tau^2 z$$

$$\hat{\theta}_{\text{MAP}} = \alpha z, \quad \alpha = \frac{\tau^2}{\tau^2 + s^2} \in (0, 1).$$

If $z = 0.6$ and $s = \tau$, then $\alpha = \frac{1}{2}$ and MAP = 0.3. Ridge reaches exactly zero only when $z = 0$.

Prior width controls how far MAP moves

Same likelihood: $L(\theta) \propto \mathcal{N}(\theta; \hat{\theta}_{\text{MLE}} = 2, s^2)$ with $s = 0.6$; prior: $p(\theta) = \mathcal{N}(0, \tau^2)$.

$$\hat{\theta}_{\text{MAP}} = \frac{\tau^2}{\tau^2 + s^2} \hat{\theta}_{\text{MLE}} \quad (\text{a shrinkage factor between 0 and 1})$$

— prior — likelihood — posterior

Prior width controls how far MAP moves


Same likelihood: $L(\theta) \propto \mathcal{N}(\theta; \hat{\theta}_{\text{MLE}} = 2, s^2)$ with $s = 0.6$; prior: $p(\theta) = \mathcal{N}(0, \tau^2)$.

$$\hat{\theta}_{\text{MAP}} = \frac{\tau^2}{\tau^2 + s^2} \hat{\theta}_{\text{MLE}} \quad (\text{a shrinkage factor between 0 and 1})$$

— prior — likelihood — posterior

Broad prior: $\tau = 2, \lambda = 0.125$

MLE = 2 $\rightarrow$ **MAP ≈ 1.83**



Prior width controls how far MAP moves


Same likelihood: $L(\theta) \propto \mathcal{N}(\theta; \hat{\theta}_{\text{MLE}} = 2, s^2)$ with $s = 0.6$; prior: $p(\theta) = \mathcal{N}(0, \tau^2)$.

$$\hat{\theta}_{\text{MAP}} = \frac{\tau^2}{\tau^2 + s^2} \hat{\theta}_{\text{MLE}} \quad (\text{a shrinkage factor between 0 and 1})$$

— prior — likelihood — posterior


Broad prior: $\tau = 2, \lambda = 0.125$

MLE = 2 → MAP ≈ 1.83



Narrow prior: $\tau = 0.5, \lambda = 2$

MLE = 2 → MAP ≈ 0.82



Prior width controls how far MAP moves


Same likelihood: $L(\theta) \propto \mathcal{N}(\theta; \hat{\theta}_{\text{MLE}} = 2, s^2)$ with $s = 0.6$; prior: $p(\theta) = \mathcal{N}(0, \tau^2)$.

$$\hat{\theta}_{\text{MAP}} = \frac{\tau^2}{\tau^2 + s^2} \hat{\theta}_{\text{MLE}} \quad (\text{a shrinkage factor between 0 and 1})$$

— prior — likelihood — posterior


Broad prior: $\tau = 2, \lambda = 0.125$

MLE = 2 $\rightarrow$ MAP ≈ 1.83



Narrow prior: $\tau = 0.5, \lambda = 2$

MLE = 2 $\rightarrow$ MAP ≈ 0.82



The narrower prior gives the larger λ and pulls MAP farther from the MLE toward zero.

A multivariate Gaussian prior has geometry


$$\boldsymbol{\theta} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \quad -\log p(\boldsymbol{\theta}) = \frac{1}{2}(\boldsymbol{\theta} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{\theta} - \boldsymbol{\mu}) + C$$

$\boldsymbol{\Sigma}$ is the covariance matrix; $\boldsymbol{\Sigma}^{-1}$ is the precision matrix, which determines how strongly each direction is penalized.

A multivariate Gaussian prior has geometry

$$\theta \sim \mathcal{N}(\mu, \Sigma) \quad -\log p(\theta) = \frac{1}{2}(\theta - \mu)^\top \Sigma^{-1}(\theta - \mu) + C$$

Σ is the covariance matrix; Σ^{-1} is the precision matrix, which determines how strongly each direction is penalized.




Axes: θ_1 and θ_2 are two coordinates of the parameter vector θ .

Colour = prior density $p(\theta)$: pale outer regions are less plausible; darker centres are more plausible.

A multivariate Gaussian prior has geometry

$$\theta \sim \mathcal{N}(\mu, \Sigma) \quad -\log p(\theta) = \frac{1}{2}(\theta - \mu)^\top \Sigma^{-1}(\theta - \mu) + C$$

Σ is the covariance matrix; Σ^{-1} is the precision matrix, which determines how strongly each direction is penalized.



Axes: θ_1 and θ_2 are two coordinates of the parameter vector θ .


Colour = prior density $p(\theta)$: pale outer regions are less plausible; darker centres are more plausible.

Covariance says **which directions in parameter space are plausible.**

Laplace is sharper at zero than a Normal prior

Compare zero-centred priors with the **same variance**, 1:


— Normal: $\sigma = 1$ — Laplace: $a = \frac{1}{\sqrt{2}}$



Laplace is sharper at zero than a Normal prior

Compare zero-centred priors with the **same variance**, 1:

— Normal: $\sigma = 1$ — Laplace: $a = \frac{1}{\sqrt{2}}$



Normal prior


- smooth around $\theta = 0$
- lighter tails

Laplace prior

- a sharp peak at $\theta = 0$
- heavier tails

Laplace is sharper at zero than a Normal prior

Compare zero-centred priors with the **same variance**, 1:



Normal prior

- smooth around $\theta = 0$
- lighter tails

Laplace prior

- a sharp peak at $\theta = 0$
- heavier tails

The Laplace peak creates a **corner** in the negative log-prior. Next we derive why that corner can set coefficients exactly to zero.

Laplace MAP gives an L_1 objective

D DERIVATION

Use the scalar likelihood centred at $z = \hat{\theta}_{\text{MLE}}$ and a zero-centred Laplace prior with scale a :

$$p(\mathcal{D}|\theta) \propto e^{-\frac{(\theta-z)^2}{2s^2}}, \quad p(\theta) = \frac{1}{2a} e^{-\frac{|\theta|}{a}}.$$

Laplace MAP gives an L_1 objective

D DERIVATION

Use the scalar likelihood centred at $z = \hat{\theta}_{\text{MLE}}$ and a zero-centred Laplace prior with scale a :

$$p(\mathcal{D}|\theta) \propto e^{-\frac{(\theta-z)^2}{2s^2}}, \quad p(\theta) = \frac{1}{2a} e^{-\frac{|\theta|}{a}}.$$

$$-\log p(\theta) = \frac{|\theta|}{a} + \log(2a).$$

Laplace MAP gives an L_1 objective

D DERIVATION

Use the scalar likelihood centred at $z = \hat{\theta}_{\text{MLE}}$ and a zero-centred Laplace prior with scale a :

$$p(\mathcal{D}|\theta) \propto e^{-\frac{(\theta-z)^2}{2s^2}}, \quad p(\theta) = \frac{1}{2a} e^{-\frac{|\theta|}{a}}.$$

$$-\log p(\theta) = \frac{|\theta|}{a} + \log(2a).$$

$$J_1(\theta) = \frac{(\theta - z)^2}{2s^2} + \frac{|\theta|}{a} + C.$$

Laplace MAP gives an L_1 objective

D DERIVATION

Use the scalar likelihood centred at $z = \hat{\theta}_{\text{MLE}}$ and a zero-centred Laplace prior with scale a :

$$p(\mathcal{D}|\theta) \propto e^{-\frac{(\theta-z)^2}{2s^2}}, \quad p(\theta) = \frac{1}{2a} e^{-\frac{|\theta|}{a}}.$$

$$-\log p(\theta) = \frac{|\theta|}{a} + \log(2a).$$

$$J_1(\theta) = \frac{(\theta - z)^2}{2s^2} + \frac{|\theta|}{a} + C.$$

For vectors, $\|\theta\|_1 = \sum_j |\theta_j|$. Multiply by s^2 ; the minimizer does not change:

$$\tilde{J}_1(\theta) = \frac{1}{2}(\theta - z)^2 + \lambda|\theta|, \quad \lambda = \frac{s^2}{a}.$$

The Laplace negative log-prior is proportional to $|\theta|$, so MAP becomes an L_1 -regularized optimization problem.

Soft thresholding follows from three cases

D DERIVATION

Differentiate $\tilde{J}_1(\theta) = \frac{1}{2}(\theta - z)^2 + \lambda|\theta|$ on either side of the corner.

Case 1: $\theta > 0$

$$\tilde{J}_1' = \theta - z + \lambda$$

$$0 = \theta - z + \lambda$$

$$\theta = z - \lambda, \text{ valid if } z > \lambda$$

Case 2: $\theta < 0$

$$\tilde{J}_1' = \theta - z - \lambda$$

$$0 = \theta - z - \lambda$$

$$\theta = z + \lambda, \text{ valid if } z < -\lambda$$

Case 3: $\theta = 0$

$$\partial|\theta|_0 = [-1, 1]$$

$$0 \in -z + \lambda[-1, 1]$$

$$\text{valid iff } |z| \leq \lambda$$

Soft thresholding follows from three cases

D DERIVATION

Differentiate $\tilde{J}_1(\theta) = \frac{1}{2}(\theta - z)^2 + \lambda|\theta|$ on either side of the corner.

Case 1: $\theta > 0$

$$\tilde{J}_{1'} = \theta - z + \lambda$$

$$0 = \theta - z + \lambda$$

$$\theta = z - \lambda, \text{ valid if } z > \lambda$$

Case 2: $\theta < 0$

$$\tilde{J}_{1'} = \theta - z - \lambda$$

$$0 = \theta - z - \lambda$$

$$\theta = z + \lambda, \text{ valid if } z < -\lambda$$

Case 3: $\theta = 0$

$$\partial|\theta|_0 = [-1, 1]$$

$$0 \in -z + \lambda[-1, 1]$$

$$\text{valid iff } |z| \leq \lambda$$

$$z > \lambda \\ \hat{\theta}_{\text{MAP}} = z - \lambda$$

$$|z| \leq \lambda \\ \hat{\theta}_{\text{MAP}} = 0$$

$$z < -\lambda \\ \hat{\theta}_{\text{MAP}} = z + \lambda$$

$$\hat{\theta}_{\text{MAP}} = \text{sign}(z) \max(|z| - \lambda, 0).$$

Values inside the threshold band $[-\lambda, \lambda]$ map exactly to zero; values outside are pulled toward zero by λ .


Ridge shrinks; L_1 can threshold to exactly zero

For one coefficient, let the data term be $D(\theta) = \frac{1}{2}(\theta - 0.6)^2$, centred at $z = 0.6$.

Normal prior $\rightarrow L_2$

$$J_2(\theta) = D(\theta) + \frac{1}{2}\theta^2$$

$$\hat{\theta}_{\text{MAP}} = \frac{z}{1+1} = 0.3$$



Ridge shrinks; L_1 can threshold to exactly zero


D DERIVATION

For one coefficient, let the data term be $D(\theta) = \frac{1}{2}(\theta - 0.6)^2$, centred at $z = 0.6$.

Normal prior $\rightarrow L_2$

$$J_2(\theta) = D(\theta) + \frac{1}{2}\theta^2$$


$$\hat{\theta}_{\text{MAP}} = \frac{z}{1+1} = 0.3$$



Laplace prior $\rightarrow L_1$

$$J_1(\theta) = D(\theta) + |\theta|$$

$$\hat{\theta}_{\text{MAP}} = \text{sign}(z) \max(|z| - 1, 0) = 0$$



Ridge shrinks; L_1 can threshold to exactly zero


D DERIVATION

For one coefficient, let the data term be $D(\theta) = \frac{1}{2}(\theta - 0.6)^2$, centred at $z = 0.6$.

Normal prior $\rightarrow L_2$

$$J_2(\theta) = D(\theta) + \frac{1}{2}\theta^2$$


$$\hat{\theta}_{\text{MAP}} = \frac{z}{1+1} = 0.3$$



Laplace prior $\rightarrow L_1$

$$J_1(\theta) = D(\theta) + |\theta|$$


$$\hat{\theta}_{\text{MAP}} = \text{sign}(z) \max(|z| - 1, 0) = 0$$



When the quadratic data term separates by coefficient—as with orthogonal features— L_1 applies this threshold coordinate by coordinate and can produce **sparsity**: many exact zeros.


Prior strength controls underfitting and overfitting

For the summed-NLL convention used above, $\lambda = \frac{1}{2\tau^2}$:



Prior strength controls underfitting and overfitting

For the summed-NLL convention used above, $\lambda = \frac{1}{2\tau^2}$:



DATA FIT DOMINATES

flexible model

sensitive to this sample · overfit

LIKELIHOOD + PRIOR

enough flexibility

better generalization


PRIOR DOMINATES

weights forced small

too rigid · underfit

Prior strength controls underfitting and overfitting

For the summed-NLL convention used above, $\lambda = \frac{1}{2\tau^2}$:



Regularization trades flexibility for stability; more is not always better.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/map-regularization

map-regularization.html — a two-parameter linear model so contours are visible. Shows **likelihood**, **prior**, and **posterior** contours with the **MLE**, **MAP**, and zero points.

Controls: amount of data · noise · prior variance · Gaussian vs Laplace · prior correlation.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/map-regularization

map-regularization.html — a two-parameter linear model so contours are visible. Shows **likelihood**, **prior**, and **posterior** contours with the **MLE**, **MAP**, and zero points.

Controls: amount of data · noise · prior variance · Gaussian vs Laplace · prior correlation.

Ask: more data? narrower prior? why does MAP → MLE as data grows? a correlated prior?

Full Bayes and deep learning

MAP is still one parameter vector

MLE $\hat{\theta}_{\text{MLE}}$ — one point

MAP $\hat{\theta}_{\text{MAP}}$ — one point

MAP is still one parameter vector

MLE $\hat{\theta}_{\text{MLE}}$ — one point

MAP $\hat{\theta}_{\text{MAP}}$ — one point


Full Bayes keeps the whole posterior $p(\theta \mid \mathcal{D})$ — a **distribution**, not a point.

Posterior predictive distribution

$$p(y^\star \mid \boldsymbol{x}^\star, \mathcal{D}) = \int p_\theta(y^\star \mid \boldsymbol{x}^\star) p(\boldsymbol{\theta} \mid \mathcal{D}) \mathrm{d}\boldsymbol{\theta}$$


Posterior predictive distribution

$$p(y^* \mid x^*, \mathcal{D}) = \int p_{\theta}(y^* \mid x^*) p(\theta \mid \mathcal{D}) d\theta$$



Posterior predictive distribution

$$p(y^* \mid x^*, \mathcal{D}) = \int p_{\theta}(y^* \mid x^*) p(\theta \mid \mathcal{D}) d\theta$$



Bayesian prediction **averages over plausible models** — and reports uncertainty.

Why DL usually uses point estimates

★ OPTIONAL

A modern network has millions to billions of parameters, so integrating $p(\boldsymbol{\theta} \mid \mathcal{D})$ is hard.

Why DL usually uses point estimates

★ OPTIONAL

A modern network has millions to billions of parameters, so integrating $p(\boldsymbol{\theta} \mid \mathcal{D})$ is hard. Common practical choices:

- MLE · MAP

Why DL usually uses point estimates

★ OPTIONAL

A modern network has millions to billions of parameters, so integrating $p(\boldsymbol{\theta} \mid \mathcal{D})$ is hard. Common practical choices:

- MLE · MAP
- deep **ensembles**

Why DL usually uses point estimates

★ OPTIONAL

A modern network has millions to billions of parameters, so integrating $p(\boldsymbol{\theta} \mid \mathcal{D})$ is hard. Common practical choices:

- MLE · MAP
- deep **ensembles**
- **dropout**-based approximations

Why DL usually uses point estimates

★ OPTIONAL

A modern network has millions to billions of parameters, so integrating $p(\boldsymbol{\theta} \mid \mathcal{D})$ is hard. Common practical choices:

- MLE · MAP
- deep **ensembles**
- **dropout**-based approximations
- **variational inference**: fit a tractable approximate posterior

The standard supervised DL objective

B is the current minibatch; $|B|$ is its number of examples.

$$\min_{\boldsymbol{\theta}} \underbrace{\frac{1}{|B|} \sum_{i \in B} -\log p_{\boldsymbol{\theta}}(y_i \mid \mathbf{x}_i)}_{\text{likelihood / data fit}} + \underbrace{\lambda \|\boldsymbol{\theta}\|_2^2}_{\text{prior / regularization}}$$

With an averaged NLL, λ absorbs the dataset or minibatch scaling convention.

The standard supervised DL objective

B is the current minibatch; $|B|$ is its number of examples.

$$\min_{\boldsymbol{\theta}} \underbrace{\frac{1}{|B|} \sum_{i \in B} -\log p_{\boldsymbol{\theta}}(y_i \mid \mathbf{x}_i)}_{\text{likelihood / data fit}} + \underbrace{\lambda \|\boldsymbol{\theta}\|_2^2}_{\text{prior / regularization}}$$

With an averaged NLL, λ absorbs the dataset or minibatch scaling convention.

probabilistic model + neural network + gradient optimization

The standard supervised DL objective

B is the current minibatch; $|B|$ is its number of examples.

$$\min_{\theta} \underbrace{\frac{1}{|B|} \sum_{i \in B} -\log p_{\theta}(y_i \mid \mathbf{x}_i)}_{\text{likelihood / data fit}} + \underbrace{\lambda \|\theta\|_2^2}_{\text{prior / regularization}}$$

With an averaged NLL, λ absorbs the dataset or minibatch scaling convention.

probabilistic model + neural network + gradient optimization

DL doesn't discard probabilistic ML — it **scales** it with expressive functions and gradient descent.

Gaussian likelihood

⇒ **MSE**

Categorical likelihood

⇒ **cross-entropy**

Gaussian prior

⇒ **L_2**


Laplace prior

⇒ **L_1**

MLE + prior

⇒ **MAP**

Every supervised objective begins with two assumptions



Loss $\leftrightarrow$ likelihood; regularizer $\leftrightarrow$ prior; prediction $\leftrightarrow$ plug-in or posterior predictive.

Next: from linear models to neural networks — **neurons, nonlinearities, and multilayer perceptrons.**