

Where Deep Learning Losses Come From

Likelihood, cross-entropy, and regularization from probability

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

Three familiar losses, one probabilistic origin

You have already met these losses:

$$\mathcal{L}_{\text{reg}} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2 \quad (\text{regression})$$

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_i [y_i \log p_i + (1 - y_i) \log(1 - p_i)] \quad (\text{binary classification})$$

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{data}} + \lambda \|\boldsymbol{\theta}\|_2^2 \quad (\text{regularized})$$

Objective. Derive each term from an explicit likelihood or prior.

Two probabilistic assumptions determine the two halves of the objective.

We will define the predictive distribution and likelihood; use i.i.d. to move from one example to a dataset;

then introduce log-likelihood, NLL, the prior $p(\theta)$, and MAP.



Probability primer

A point prediction summarizes a distribution


REGRESSION • $Y \in \mathbb{R}$

Point prediction: mean of $Y \mid \mathbf{X} = \mathbf{x}$

$$\hat{y}(\mathbf{x}) = \mathbb{E}_{\theta}[Y \mid \mathbf{X} = \mathbf{x}] = \mu_{\theta}(\mathbf{x})$$

Full predictive distribution

$$Y \mid \mathbf{X} = \mathbf{x} \\ \sim \mathcal{N}(\mu_{\theta}(\mathbf{x}), \sigma_{\theta}^2(\mathbf{x}))$$



BINARY CLASSIFICATION • $Y \in \{0, 1\}$


Point prediction: mode of $Y \mid \mathbf{X} = \mathbf{x}$

$$\hat{y}(\mathbf{x}) = \mathbf{1}[p_{\theta}(\mathbf{x}) \geq 0.5]$$

Full predictive distribution

$$Y \mid \mathbf{X} = \mathbf{x} \\ \sim \text{Bernoulli}(p_{\theta}(\mathbf{x}))$$

example: $p_{\theta}(\mathbf{x}) = 0.75$



Neural network: $\mathbf{x} \rightarrow$ distribution parameters. **Prediction:** $\hat{y} =$ mean or mode.

Discrete distributions assign probability to outcomes


The **support** $\mathcal{S}_Y$ is the set of values that Y can take.

Probability mass function (PMF)

$$p_{Y(y)} = P(Y = y), \quad y \in \mathcal{S}_Y$$

$$p_{Y(y)} \geq 0, \quad \sum_{y \in \mathcal{S}_Y} p_{Y(y)} = 1$$

Bernoulli(0.75) PMF



each bar is an outcome probability

Examples. Bernoulli: $\{0, 1\}$ · categorical: $\{1, \dots, K\}$ · Poisson: $\{0, 1, 2, \dots\}$

Exact values can have nonzero probability: here $P(Y = 1) = 0.75$.

Continuous distributions assign density across a support


The **support** may be bounded or unbounded:

$$\mathcal{S}_Y = [a, b]$$

for a Uniform variable; $\mathcal{S}_Y = \mathbb{R}$ for a Gaussian.

Probability density function (PDF)

$$p_{Y(y)} \geq 0, \quad \int_{\mathcal{S}_Y} p_{Y(y)} \, dy = 1$$



the total area under a PDF is 1


For continuous Y , probability comes from **area under the density.**

Density is not probability

For a continuous Y , probability comes from **area**:

$$P(a \leq Y \leq b) = \int_a^b p_{Y(y)} dy$$


$p_{Y(y)} \neq P(Y = y)$ — a density can **exceed 1**; only *integrals* are probabilities.



the shaded area is $P(a \leq Y \leq b)$


$$Y \sim \mathcal{U}(a, b) \quad p_{Y(y)} = \begin{cases} 1/(b-a) & a \leq y \leq b \\ 0 & \text{otherwise} \end{cases}$$

The support matters: outside $[a, b]$, the density is exactly 0.



$$Y \sim \mathcal{N}(\mu, \sigma^2) \quad p(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$

— $\mu = 0, \sigma = 1$ — $\mu = 0, \sigma = 1.8$ — $\mu = -1.5, \sigma = 1$



Only two knobs: **location** μ and **scale** σ .

From probability to likelihood

Read the model notation one symbol at a time

$$p_{\theta}(y \mid x)$$

X / x	X is the random input vector; x is one observed input
Y / y	Y is the random target; y is one possible or observed value
θ	the parameter vector — the model's learned weights and biases
$p_{\theta}(y \mid x)$	the probability (discrete Y) or density (continuous Y) assigned to y , given x

Read aloud: “model θ ’s probability or density for y , given x .”

Prediction fixes the model and varies the outcome

Example: $Y \in \{0, 1\}$. Fix the input x^* and one parameter setting θ_A .


$$p_{\theta_A}(Y = 0 \mid x^*) = 0.20$$

$$p_{\theta_A}(Y = 1 \mid x^*) = 0.80$$

$$0.20 + 0.80 = 1$$

We vary the possible value of Y while keeping θ_A and x^* fixed.

predictive distribution



$$\hat{y}(x^*) = \arg \max_y p_{\theta_A}(y \mid x^*) = 1$$

One observed label gives one likelihood score

Now observe $y^* = 1$ for the input x^* , so $\mathcal{D} = \{(x^*, 1)\}$.

$$L(\theta; \mathcal{D}) = p_{\theta}(y^* = 1 \mid x^*)$$


Candidate A: $L(\theta_A; \mathcal{D}) = 0.80$

Candidate B: $L(\theta_B; \mathcal{D}) = 0.30$

$$L(\theta_A; \mathcal{D}) > L(\theta_B; \mathcal{D})$$

θ_A makes the observed label $y^* = 1$ less surprising.

likelihood of the fixed observation



Likelihood is a **score as a function of θ** ; the scores need not sum to 1 over parameter settings.

One continuous observation gives a likelihood too


Suppose $Y \mid \mu \sim \mathcal{N}(\mu, 1)$ and we observe the fixed value $y^* = 1$.

$$L(\mu; y^* = 1) = p(y^* = 1 \mid \mu) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(1 - \mu)^2}{2}\right).$$

Candidate $\mu = 1$: $L(1) \approx 0.399$

Candidate $\mu = 0$: $L(0) \approx 0.242$

$\mu = 1$ centres the Gaussian at $y^* = 1$
and gives the largest **density**.



Discrete or continuous: fix what was observed, then ask which parameter values explain it better.

A dataset needs a joint likelihood

For two observations, the likelihood starts as a **joint** probability or density:

$$L(\boldsymbol{\theta}; \mathcal{D}) = p_{\boldsymbol{\theta}}(y_1, y_2 \mid \boldsymbol{x}_1, \boldsymbol{x}_2).$$

We may write a product only after making an assumption about how the observations relate.

Question: when is $p(y_1, y_2 \mid \boldsymbol{\theta}) = p(y_1 \mid \boldsymbol{\theta})p(y_2 \mid \boldsymbol{\theta})$ valid?

The product is a consequence of **conditional independence**; it is not automatic.

Independent means “no extra information”

After fixing θ , observing one outcome does not change the distribution of another:

$$p(y_2 \mid y_1, \theta) = p(y_2 \mid \theta).$$

The chain rule then becomes

$$p(y_1, y_2 \mid \theta) = p(y_1 \mid \theta)p(y_2 \mid y_1, \theta) = p(y_1 \mid \theta)p(y_2 \mid \theta).$$

Coin-flip example. If $P(H \mid \theta) = \theta$ and the flips are independent, then $P(H, H \mid \theta) = \theta^2$.

Independence is what licenses multiplication.

Identically distributed means “the same sampling rule”

D DERIVATION

Coin flips

$Y_i \mid \theta \sim \text{Bernoulli}(\theta)$ for every i .
Each flip uses the same bias θ .

Supervised learning

The pairs $(\mathbf{X}_i, Y_i)$ come from the same population, and one shared θ defines $p_\theta(y_i \mid \mathbf{x}_i)$.

“Independent” and “identically distributed” are two separate claims; one can hold without the other.

i.i.d. = independent observations + one common data-generating rule.

Not independent: daily weather

Tomorrow's weather depends on today's. A first-order Markov model writes

$$p(y_{1:T} \mid \theta) = p(y_1 \mid \theta) \prod_{t=2}^T p(y_t \mid y_{t-1}, \theta).$$

The factors are **conditional on the previous day**, not independent marginals.

Independent but not identical: two sensors

Two sensors measure the same temperature μ independently, but have different noise levels:

$$Y_i \mid \mu \sim \mathcal{N}(\mu, \sigma_i^2), \quad \sigma_1 \neq \sigma_2.$$

The likelihood is a product, but its factors use different σ_i .

i.i.d. is a useful special case—not a universal law.

Let R_1 mean “rain today” and R_2 mean “rain tomorrow”. Suppose

$$P(R_1) = 0.4, \quad P(R_2|R_1) = 0.8, \quad P(R_2|\neg R_1) = 0.2.$$

1. Use the conditional rule

$$\begin{aligned} P(R_1, R_2) &= P(R_1)P(R_2|R_1) \\ &= 0.4 \times 0.8 = 0.32. \end{aligned}$$

Both days are rainy with probability 0.32.

2. Test independence

$$P(R_2) = 0.4 \times 0.8 + 0.6 \times 0.2 = 0.44.$$

If the days were independent,

$$P(R_1)P(R_2) = 0.4 \times 0.44 = 0.176,$$

not 0.32.

Dependence does not forbid multiplication: use a conditional probability.

Worked example: independent does not mean identical

Given smoke S , two alarms act independently, but their sensitivities differ:

$$P(A_1 = 1|S) = 0.9, \quad P(A_2 = 1|S) = 0.6.$$

Independence

Both alarms ring with probability

$$\begin{aligned} P(A_1 = 1, A_2 = 1|S) \\ = 0.9 \times 0.6 = 0.54. \end{aligned}$$

The product is valid.

Not identically distributed

Alarm 1 detects smoke 90% of the time;
Alarm 2 detects it 60% of the time.

$$0.9 \neq 0.6$$

We cannot reuse one common factor for both alarms.

Independent says “multiply”; identical says “reuse the same sampling rule”.

IID turns the joint likelihood into per-example factors

D DERIVATION

For i.i.d. supervised examples, treat the observed inputs as fixed. Conditional independence gives

$$p(\mathcal{D} \mid \theta) = \prod_{i=1}^n p_{\theta}(y_i \mid \mathbf{x}_i).$$

Taking logs turns the product into the sum used by every training loop:

$$\log p(\mathcal{D} \mid \theta) = \sum_{i=1}^n \log p_{\theta}(y_i \mid \mathbf{x}_i).$$

The **factorization** comes from independence; the shared factor form comes from the common model p_{θ} .

This is the bridge from one-example likelihood to a dataset loss.

Work two tiny Bernoulli datasets

Compare a fair coin $\theta = 0.5$ with a head-biased coin $\theta = 0.8$.

Observed data	$L(0.5)$	$L(0.8)$	Which explains it better?
H, T	$0.5(0.5) = 0.25$	$0.8(0.2) = 0.16$	fair coin
H, H	$0.5^2 = 0.25$	$0.8^2 = 0.64$	head-biased coin

Same candidate models, different fixed data. Likelihood ranks candidates by how well they explain the observations actually seen.

Multiply one Bernoulli term per independent flip.

Two Gaussian measurements: multiply densities

Let $Y_i \mid \mu \sim \mathcal{N}(\mu, 1)$ independently, and observe $y_1 = 1, y_2 = 2$.

Candidate $\mu = 1.5$

$$p(y_1 \mid \mu) \approx 0.352$$

$$p(y_2 \mid \mu) \approx 0.352$$

$$L(1.5) \approx 0.352 \times 0.352 = 0.124$$

Candidate $\mu = 0$

$$p(y_1 \mid \mu) \approx 0.242$$

$$p(y_2 \mid \mu) \approx 0.054$$

$$L(0) \approx 0.242 \times 0.054 = 0.013$$

$$\log L(1.5) \approx -2.09, \quad \log L(0) \approx -4.34.$$

For continuous data we multiply **densities; $\mu = 1.5$ gives the larger likelihood.**

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L01/00_likelihood_iid_worked_examples.ipynb

Explore `torch.distributions`, `log_prob`, sampling, i.i.d. factorization, coin-toss MLE, Gaussian regression, Bernoulli classification, and MAP with Gaussian, Laplace, and Student- t priors.

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L01/03_robust_linear_regression.ipynb

Short companion: fit the same line with Gaussian/MSE, Laplace/MAE, and Student- t losses; then see which observations control each estimate.

In class: predict each result before running the cell; then connect the code to the matching derivation in the slides.

Coin flips: a likelihood you can compute

Observed: $H, H, T, T, T, H, H, T, T, T$

Hence $n_H = 4$ and $n_T = 6$.


Model: $p(H \mid \theta) = \theta$, $p(T \mid \theta) = 1 - \theta$.

Because the flips are i.i.d.,

$$L(\theta) = \prod_i p(x_i \mid \theta)$$

$$= \theta^{n_H} (1 - \theta)^{n_T}$$

$$= \theta^4 (1 - \theta)^6$$



$$L(0.4) \approx 0.00119$$

$$L(0.7) \approx 0.000175$$

about $6.8 \times$ larger at $\theta = 0.4$

The likelihood peaks at $\theta = 0.4$ — the observed head fraction $\frac{4}{10}$.

The **log-likelihood** is the logarithm of the likelihood — not a different model:

$$\ell(\theta) := \log L(\theta) = \log[\theta^4(1 - \theta)^6] = 4 \log \theta + 6 \log(1 - \theta)$$

Candidate $\theta = 0.4$

$$L(0.4) \approx 0.00119$$

$$\ell(0.4) = 4 \log 0.4 + 6 \log 0.6 \approx -6.73$$

Candidate $\theta = 0.7$

$$L(0.7) \approx 0.000175$$

$$\ell(0.7) = 4 \log 0.7 + 6 \log 0.3 \approx -8.65$$

Likelihoods below 1 have negative logs; among them, -6.73 is the larger value.

The observed sequence has higher log-likelihood at $\theta = 0.4$.

Why logs? The maximizing parameter does not change

D DERIVATION

The logarithm is **strictly increasing** on positive numbers:

$$a > b > 0 \implies \log a > \log b$$

Therefore it preserves the ordering of all positive likelihood values:

$$\begin{aligned} L(0.4) &> L(0.7) \\ \text{apply the increasing function } \log \\ \log L(0.4) &> \log L(0.7) \end{aligned}$$

$$\arg \max_{\theta} L(\theta) = \arg \max_{\theta} \log L(\theta)$$


Both likelihood and log-likelihood rank $\theta = 0.4$ above $\theta = 0.7$.

Poisson example: logging preserves the peak

Suppose $Y \sim \text{Poisson}(\lambda)$ and the observed count is $y = 3$.


LIKELIHOOD

$$L(\lambda) = e^{-\lambda} \frac{\lambda^3}{3!}$$



LOG-LIKELIHOOD

$$\log L(\lambda) = 3 \log \lambda - \lambda - \log 3!$$



Both curves peak at $\hat{\lambda}_{\text{MLE}} = 3$: the log changes the vertical scale, not the maximizing parameter.

Why logs? Products become stable sums

For independent observations,

$$L(\boldsymbol{\theta}) = \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \implies \log L(\boldsymbol{\theta}) = \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

Worked numerical example. Suppose 200 observed outcomes each receive probability 0.01.

Direct product	$L = 0.01^{200} = 10^{-400} < 4.94 \times 10^{-324} \rightarrow \text{underflows to } 0$
Log space	$\log L = 200 \log(0.01) \approx -921.03 \rightarrow \text{finite and representable}$

Float64 range: $\pm 1.80 \times 10^{308}$ · smallest positive normal 2.23×10^{-308} · smallest subnormal 4.94×10^{-324} .

Logs make optimization easier and prevent numerical underflow.

For $L(\theta) = \theta^{n_H}(1 - \theta)^{n_T}$, the log-likelihood is

$$\log L(\theta) = n_H \log \theta + n_T \log(1 - \theta)$$

Set the derivative to zero:

$$\frac{d}{d\theta} \log L(\theta) = \frac{n_H}{\theta} - \frac{n_T}{1 - \theta} = 0$$

$$\implies \theta^* = \frac{n_H}{n_H + n_T} = \frac{4}{10} = 0.4$$

$\theta^* = 0.4$ is a **candidate** optimum. Next, check the curvature.

The second derivative confirms a maximum

Differentiate the log-likelihood again:

$$\frac{d^2}{d\theta^2} \log L(\theta) = -\frac{n_H}{\theta^2} - \frac{n_T}{(1-\theta)^2}$$

For $n_H, n_T > 0$ and $0 < \theta < 1$, both terms are negative:

$$\frac{d^2}{d\theta^2} \log L(\theta) < 0$$

At the stationary point $\theta^* = 0.4$,

$$-\frac{4}{(0.4)^2} - \frac{6}{(0.6)^2} = -25 - 16.67 \approx -41.67 < 0$$

Strict concavity + stationary point $\rightarrow$ unique global maximum: $\hat{\theta}_{\text{MLE}} = \theta^* = \frac{4}{10} = 0.4$.

Maximum likelihood estimation

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} p(\mathcal{D} \mid \theta)$$

Interpretation. Choose the parameter values that assign the **highest likelihood** to the observed dataset.

Maximizing likelihood = minimizing NLL

$$\arg \max_{\boldsymbol{\theta}} \prod_i p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) = \arg \max_{\boldsymbol{\theta}} \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

$$= \arg \min_{\boldsymbol{\theta}} \left(- \sum_i \log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i) \right)$$

$$\ell_{i(\boldsymbol{\theta})} = -\log p_{\boldsymbol{\theta}}(y_i \mid \boldsymbol{x}_i)$$

Empirical risk = average training loss

In ML, **risk** means expected loss. Because the data-generating distribution is unknown, we estimate it from the observed training set.

Empirical risk = average loss on $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$.

Per example	$\ell_{i(\theta)} = -\log p_{\theta}(y_i \mid x_i)$
Training set	$\hat{R}(\theta) = \frac{1}{n} \sum_i \ell_{i(\theta)} = -\frac{1}{n} \log p(\mathcal{D} \mid \theta)$

The final equality uses conditional independence: $p(\mathcal{D} \mid \theta) = \prod_i p_{\theta}(y_i \mid x_i)$.

Minimizing empirical NLL is maximum likelihood; $\frac{1}{n}$ only rescales the objective.

Recall the Uniform observation model:

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \begin{cases} 1/(b - a) & a \leq y_i \leq b \\ 0 & \text{otherwise} \end{cases}$$

If an observed target lies outside its support, then

$$p_{\theta}(y_i \mid \mathbf{x}_i) = 0$$

Its contribution to the negative log-likelihood is

$$\ell_{i(\theta)} = -\log p_{\theta}(y_i \mid \mathbf{x}_i) = -\log 0 = +\infty$$

Zero likelihood becomes an infinite training penalty.

For a $\text{Uniform}(a, b)$ observation model, what is the NLL if the target lies outside $[a, b]$?

A. 0

B. A fixed finite penalty

C. $-\log 0 = +\infty$

D. It depends only on σ

Correct: C — $-\log 0 = +\infty$

Why: The model assigns zero density outside its support, so the observation is impossible under the assumption.

Regression: why MSE?


Regression as signal plus noise

$$y_i = f_{\boldsymbol{\theta}}(\mathbf{x}_i) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$

Therefore the conditional output distribution is

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(f_{\boldsymbol{\theta}}(\mathbf{x}_i), \sigma^2)$$

What the assumption means visually



Gaussian noise lives in the **output direction**, conditional on x — a vertical bell at every input.

$$p_{\theta}(y_i \mid \mathbf{x}_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - f_{\theta}(\mathbf{x}_i))^2}{2\sigma^2}\right)$$

Negative log of a single example:

$$-\log p_{\theta}(y_i \mid \mathbf{x}_i) = \frac{(y_i - f_{\theta}(\mathbf{x}_i))^2}{2\sigma^2} + C$$

MSE is Gaussian MLE

Sum over the dataset:

$$-\log p(\mathcal{D} \mid \boldsymbol{\theta}) = \frac{1}{2\sigma^2} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2 + C$$

With σ fixed, the minimizer is

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \min_{\boldsymbol{\theta}} \sum_i (y_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^2$$

MSE = Gaussian NLL, up to constants.

Linear regression makes f_{θ} explicit

For the one-dimensional picture,

$$\hat{y}_i = \theta_0 + \theta_1 x_i$$

For d features, prepend a constant 1 so the intercept is part of the same parameter vector:

$$\tilde{x}_i = (1, x_i), \quad \theta = (\theta_0, \theta_1, \dots, \theta_d)$$

Then the same notation used throughout the lecture becomes

$$\hat{y}_i = f_{\theta}(x_i) = \theta^{\top} \tilde{x}_i$$

θ_0 is the intercept; $\theta_1, \dots, \theta_d$ are the slopes.

Substitute the linear mean into the observation model:

$$Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \mathcal{N}(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i, \sigma^2)$$

The vertical residual of point i is

$$r_i = y_i - \hat{y}_i = y_i - \boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i$$

With σ fixed, maximum likelihood therefore gives

$$\hat{\boldsymbol{\theta}}_{\text{MLE}} = \arg \min_{\boldsymbol{\theta}} \sum_i \underbrace{r_i^2}_{\text{squared vertical error}}$$

Ordinary least squares is Gaussian maximum likelihood for a linear mean.

The normal equations solve linear regression

Stack the augmented inputs into $\tilde{X} = [1 \quad X]$. Then

$$\hat{y} = \tilde{X}\theta, \quad J(\theta) = \|y - \tilde{X}\theta\|_2^2$$

Differentiate and set the gradient to zero:

$$\nabla_{\theta} J = 2\tilde{X}^{\top}(\tilde{X}\theta - y) = 0$$

This gives the **normal equations**:

$$\tilde{X}^{\top} \tilde{X} \hat{\theta} = \tilde{X}^{\top} y$$

If $\tilde{X}^{\top} \tilde{X}$ is invertible,

$$\hat{\theta} = (\tilde{X}^{\top} \tilde{X})^{-1} \tilde{X}^{\top} y$$

Linear regression has a closed form; neural networks keep the MSE objective but use iterative optimization.

Why earlier ML courses could ignore σ^2

σ^2 is the observation-noise variance: the expected squared spread of targets around the mean prediction.

For n Gaussian observations, write $r_i = y_i - \mu_i$. The dataset NLL is

$$\mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \frac{1}{2\sigma^2} \sum_i r_i^2 + \frac{n}{2} \log \sigma^2 + C.$$

If one global σ^2 is fixed, then $\frac{n}{2} \log \sigma^2$ is constant and $\frac{1}{2\sigma^2}$ is only a positive scale factor:

$$\arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}; \sigma^2) = \arg \min_{\boldsymbol{\theta}} \sum_i r_i^2.$$

For point predictions, fixed-variance Gaussian MLE gives the same fitted model as MSE.

How is σ^2 chosen in practice?

One global noise level.

Fit the mean model, then estimate the residual spread:

$$\hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{\mu}_i)^2.$$

Use this when a residual plot has roughly the same vertical width across inputs.

Input-dependent noise.

Let the network predict both

$$(\mu_i, s_i) = f_{\theta}(\mathbf{x}_i),$$

$$s_i = \log \sigma_i^2, \quad \sigma_i^2 = e^{s_i} > 0.$$

Use this when residual spread changes with the input—for example, a funnel shape.

Fit both cases with Gaussian NLL; validate whether uncertainty intervals attain their promised coverage.


Homoscedastic vs input-dependent uncertainty

HOMOSCEDASTIC



Network output: μ_i
Shared variance: σ^2
same uncertainty at every x_i

HETEROSCEDASTIC




Network outputs: μ_i and s_i
 $s_i = \log \sigma_i^2$, $\sigma_i^2 = e^{s_i} > 0$
uncertainty changes with x_i

Predict **log-variance** so every real s_i maps to a positive variance e^{s_i} .

A few unusual observations can dominate a fitted line

Q QUESTION



— fit to the ten typical points — MSE fit after adding two outliers

Ten typical observations lie near one linear trend.

Two unusual observations have large vertical errors.



The squared-error fit becomes much flatter. How can two points change the fitted line so strongly?

Robust regression asks for a noise model that allows occasional large residuals without letting them dominate the fit.

Tails are easier to compare on a log scale

For a fitted mean $\hat{y}$, the residual is $r = y - \hat{y}$. Compare three unit-scale models for that residual.

— Gaussian — Laplace — Student- t , $\nu = 3$



Heavy tails assign more density—and therefore less surprise—to an occasional large residual.

A large residual is less surprising under heavy tails

D DERIVATION

Evaluate the three densities at the same residual, $r = 4$, using unit scale parameters.

Gaussian

density ≈ 0.000134

The model finds this residual extremely surprising.

Laplace

density ≈ 0.00916

More tail density makes the same residual less surprising.

Student- t , $\nu = 3$

density ≈ 0.00916

Heavy tails explicitly allow occasional extreme errors.

Laplace and Student- t assign about $68 \times$ the Gaussian density here.

The observation is not discarded; the assumed noise process simply says that such an error can occur.

One observation contributes one gradient signal

For observation i , the model predicts $\hat{y}_i$ and leaves residual $r_i = y_i - \hat{y}_i$.



Loss value ℓ_i

says how incompatible this observation is with the current prediction.

Gradient magnitude $|\ell_{i'}(r_i)|$

says how strongly this observation scales the next parameter update.

$$\nabla_{\theta} \ell_i = -\ell_{i'}(r_i) \nabla_{\theta} \hat{y}_i$$

The sign gives the correction direction; $|\ell_{i'}(r_i)|$ gives the size of this observation's gradient factor.

For one observation, $\hat{y}_i = f_{\theta}(x_i)$ and $r_i = y_i - \hat{y}_i$.



$$\psi(r) = -\frac{d}{dr} \log p_{R(r)}$$

$$\nabla_{\theta} \ell_i = \psi(r_i) \nabla_{\theta} r_i = -\psi(r_i) \nabla_{\theta} f_{\theta}(x_i).$$

residual derivative $\times$ residual-to-model derivative

For $f_{\theta}(x) = b + wx$ with $\theta = (b, w)$:

$$\nabla_{\theta} \ell_i = -\psi(r_i)(1, x_i)^{\top}.$$

The density chooses $\psi(r)$; the model supplies $\nabla_{\theta} f$. Their product is the observation's parameter gradient.

Noise	Residual density	NLL $\ell(r)$	$\psi(r) = \ell'(r)$
	$p_R(r)$		
Gaussian	$\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{r^2}{2\sigma^2}}$	$\frac{r^2}{2\sigma^2} + C$	$\frac{r}{\sigma^2}$
Laplace	$\frac{1}{2b} e^{-\frac{ r }{b}}$	$\frac{ r }{b} + C$	$\frac{\text{sign}(r)}{b}$

$$\nabla_{\theta} \ell_i = -\psi(r_i) \nabla_{\theta} f_{\theta}(x_i)$$

At $r = 0$, Laplace uses the subgradient interval $\left[-\frac{1}{b}, \frac{1}{b}\right]$. Away from zero, its gradient magnitude is capped at $\frac{1}{b}$.

Gaussian influence grows with $|r|$; Laplace influence has constant magnitude.

Student- t residual: the full chain

D DERIVATION


For degrees of freedom ν and scale s :

$$\begin{aligned} \text{density} \quad p_{R(r)} &= c_{\nu,s} \left(1 + \frac{r^2}{\nu s^2}\right)^{-\frac{\nu+1}{2}} \\ \text{NLL} \quad \ell(r) &= \frac{\nu+1}{2} \log\left(1 + \frac{r^2}{\nu s^2}\right) + C \\ r \text{ derivative} \quad \psi(r) &= \frac{d\ell}{dr} = (\nu+1) \frac{r}{\nu s^2 + r^2} \\ \theta \text{ gradient} \quad \nabla_{\theta} \ell_i &= -(\nu+1) \frac{r_i}{\nu s^2 + r_i^2} \nabla_{\theta} f_{\theta}(x_i) \end{aligned}$$

Near zero, $\psi(r) \approx \left(\frac{\nu+1}{\nu s^2}\right)r$; for $|r|$ very large, $\psi(r) \approx \frac{\nu+1}{r} \rightarrow 0$.

Student- t treats small residuals approximately quadratically, but an extreme residual's gradient contribution eventually decreases.

Large residuals create different gradient magnitudes



At $r = 4$, the gradient factor is:

Gaussian

$$|\ell'(4)| = 4$$

Laplace

$$|\ell'(4)| = 1$$


Student- t

$$|\ell'(4)| = \frac{16}{19} \approx 0.84$$

Robust fitting keeps every observation, while limiting the gradient contribution of an extreme residual.

The noise model determines how errors are penalized

Noise model	Negative log-likelihood
— Gaussian	r^2
— Laplace	$ r $
— Uniform	0 inside; $+\infty$ outside
— Student- t	$\log\left(1 + \frac{r^2}{\nu}\right)$




$r = y - \hat{y}$ · swatches match plot

At $r = 0$, prediction equals target. Farther horizontally = larger error; higher vertically = larger NLL penalty.

Gaussian penalizes large errors most; Student- t is more robust. The dashed Uniform boundaries mark an infinite penalty outside the allowed interval.

Fit the same data under three noise models



Learned $\hat{y} = b + wx$


model	fitted (b, w)
Gaussian	$(-0.040, 0.397)$
Laplace	$(0.044, 0.906)$
Student- t	$(0.053, 0.882)$

$$\sigma = 1 \cdot b = 1 \cdot (\nu, s) = (3, 1)$$

All twelve observations are retained. The models differ only in the residual likelihood optimized.

The Gaussian fit is pulled flat by the two red points; Laplace and Student- t stay close to the trend supported by the ten typical points.

Student- t hyperparameters control robustness



— $\nu = 1$ — $\nu = 3$ — $\nu = 30$ (all $s = 1$)

Learned coefficients

ν	s	fitted (b, w)
1	1	(0.064, 0.928)
3	0.5	(0.065, 0.934)
3	1	(0.053, 0.882)
3	1.5	(0.036, 0.807)
30	1	(-0.014, 0.581)

Smaller ν gives heavier tails. Larger s delays the transition from quadratic to tail-robust behaviour.

For unregularized Gaussian and Laplace fits, changing the fixed scale only rescales the objective and leaves (b, w) unchanged; Student- t scale changes relative residual weights.

ν and s are modelling choices: cross-validation or domain knowledge should choose how quickly an observation is treated as a tail event.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/likelihood-to-loss

likelihood-to-loss.html — three synchronized panels: the **noise density** $p(r)$, the **loss** $-\log p(r)$, and **data with a fitted line**.

Controls: Gaussian / Laplace / Uniform / Student- t · noise scale · slope & intercept · **add an outlier** · per-point contributions.


Ask: which model reacts most to an outlier? why does Uniform give an infinite barrier? which loss is robust?

Classification: cross-entropy

Classification should output probabilities

For K classes:

$$p_{\theta}(Y = k \mid \mathbf{x}) \geq 0, \quad \sum_{k=1}^K p_{\theta}(Y = k \mid \mathbf{x}) = 1$$



f_{θ} may be **linear** (softmax regression) or **nonlinear** (a neural network).
It produces **logits** $z_k \in \mathbb{R}$; softmax converts them to **probabilities** $p_k \in [0, 1]$.

Binary classification: Bernoulli model

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Bernoulli}(p_{\boldsymbol{\theta}}(\mathbf{x}))$$

Compact one-line form:

$$p_{\boldsymbol{\theta}}(y \mid \mathbf{x}) = p^y(1 - p)^{1-y}$$

$$y = 1 \Rightarrow p(y) = p$$

$$y = 0 \Rightarrow p(y) = 1 - p$$

$$\begin{aligned} -\log p_{\theta}(y \mid \mathbf{x}) &= -\log(p^y(1-p)^{1-y}) \\ &= -y \log p - (1-y) \log(1-p) \end{aligned}$$

Binary cross-entropy = Bernoulli NLL.

Check. True label $y = 1$: if $p = 0.8$ the loss is $-\log 0.8 \approx 0.22$; if $p = 0.2$ it is $-\log 0.2 \approx 1.61$.

Why use logits?

The network emits an unbounded score $z = f_{\theta}(x) \in \mathbb{R}$; the **sigmoid** maps it to a probability:


$$p = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z \rightarrow -\infty : p \rightarrow 0$$

$$z = 0 : p = 0.5$$

$$z \rightarrow +\infty : p \rightarrow 1$$

In practice sigmoid + BCE are fused for numerical stability.



Logistic regression = Bernoulli MLE

D DERIVATION

Use an augmented input $\tilde{\mathbf{x}}_i$ so $\boldsymbol{\theta}$ contains both weights and the intercept:

$$p_i = \sigma(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i), \quad Y_i \mid \mathbf{X}_i = \mathbf{x}_i \sim \text{Bernoulli}(p_i)$$

. Likelihood of the whole labelled dataset (i.i.d.):

$$L(\boldsymbol{\theta}) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

Its negative log — the objective we minimise:

$$J(\boldsymbol{\theta}) = - \sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)]$$

Logistic regression is Bernoulli maximum likelihood; its training objective is **binary cross-entropy**.

CATEGORICAL OBSERVATION

$$Y \mid \mathbf{X} = \mathbf{x} \sim \text{Categorical}(p_1, \dots, p_K)$$

A **one-hot** target has one 1 at the true class and 0 everywhere else.


Example: class 2 of 3 $\rightarrow \mathbf{y} = (0, 1, 0)$

$$p_{\theta}(\mathbf{y} \mid \mathbf{x}) = \prod_{k=1}^K p_k^{y_k}$$

NETWORK PARAMETERIZATION

$$\mathbf{z} = f_{\theta}(\mathbf{x}), \quad p_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

softmax probabilities



Softmax guarantees $p_k > 0$ and $\sum_k p_k = 1$, as a categorical model requires.

Categorical NLL gives cross-entropy

D DERIVATION

$$-\log p_{\theta}(\mathbf{y} \mid \mathbf{x}) = -\log \prod_k p_k^{y_k} = -\sum_k y_k \log p_k$$

Because $\mathbf{y}$ is one-hot, only the true class survives:

$$\mathcal{L} = -\log p_{\text{true class}}$$

Confident mistakes are expensive


$$\ell(p_y) = -\log p_y$$

Read off the loss at three confidences for the **true** class:

$$p_y = 0.9 \Rightarrow \ell \approx 0.105$$

$$p_y = 0.1 \Rightarrow \ell \approx 2.303$$

$$p_y = 0.001 \Rightarrow \ell \approx 6.908$$



The negative log makes a confident wrong prediction a large loss.

Why does cross-entropy penalize a confidently wrong prediction strongly?

A. It makes all class probabilities equal

B. The true-class probability is near zero, so $-\log p_y$ is large

C. It adds an L_2 penalty to the logits

D. The gradient is always zero

Correct: B — The true-class probability is near zero

Why: The negative log grows without bound as p_y approaches zero, matching the fact that the model declared the observed class nearly impossible.

A binary label takes only two values

★ OPTIONAL

For one observed example, let the target be $y = 1$ and the model predict

$$p = P(Y = 1|\mathbf{x}) = 0.8.$$

Squared error

$$\ell_{\text{MSE}} = (y - p)^2 = (1 - 0.8)^2 = 0.04.$$

Its usual likelihood interpretation is Gaussian:

$$Y|\mathbf{x} \sim \mathcal{N}(0.8, \sigma^2).$$

Permits all real Y —not just labels.

Binary cross-entropy

$$\ell_{\text{CE}} = -\log p = -\log 0.8 \approx 0.223.$$

Its likelihood is Bernoulli:

$$Y|\mathbf{x} \sim \text{Bernoulli}(0.8).$$

Permits only $Y = 0$ or $Y = 1$.

MSE is possible; cross-entropy is the negative log-likelihood matched to a binary observation.

Bayes' rule for learning

Bayes' rule: from events to model parameters

PROBABILITY LECTURE

$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}$$

A – event or hypothesis to infer

B – observed evidence

Examples: disease given a test; urn given a sampled colour.

MACHINE / DEEP LEARNING

$$p(\boldsymbol{\theta} \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \boldsymbol{\theta})p(\boldsymbol{\theta})}{p(\mathcal{D})}$$

$\boldsymbol{\theta}$ – model parameters to infer

$\mathcal{D}$ – observed training dataset

Question: which parameter settings remain plausible after seeing the data?

Replace A by $\boldsymbol{\theta}$ and B by $\mathcal{D}$: the algebra is unchanged; the interpretation changes.


Three heads do not prove $\theta = 1$

QUESTION

Let $\theta = P(H)$ and suppose the first three flips are $\mathcal{D} = (H, H, H)$.

H • H • H

$$L(\theta) = \theta^3$$



$$L(\theta; \mathcal{D}) = p(\mathcal{D} \mid \theta) = \theta^3$$

$$\hat{\theta}_{\text{MLE}} = \arg \max_{0 \leq \theta \leq 1} \theta^3 = 1$$

Did we learn that the coin **always** lands heads—or did a merely head-biased coin produce a lucky run?

MLE identifies the best-fitting θ ; it does **not make $\theta = 1$ certain.**

A prior tempers the conclusion from three flips

Choose a smooth prior on $\theta \in [0, 1]$, highest near the fair-coin value 0.5.

Prior

$$p(\theta) = 6\theta(1 - \theta)$$

Likelihood of H, H, H


$$p(\mathcal{D} \mid \theta) = \theta^3$$

Posterior

$$p(\theta \mid \mathcal{D}) \propto \theta^3 p(\theta) \propto \theta^4(1 - \theta)$$

Likelihood peak (MLE): $\theta = 1$

Posterior peak: $\theta = \frac{4}{5} = 0.8$



With little data, the prior matters; as observations accumulate, the likelihood increasingly dominates.

Priors, MAP and regularization

Every learner has an inductive bias

When several predictors fit the training data, the learner still needs a preference for **unseen inputs**. That preference is its **inductive bias**.

Model	Built-in preference → consequence
Linear regression	responses vary linearly with features → line or plane extrapolation
k-NN	nearby inputs tend to have similar outputs → local prediction
Decision tree	a few axis-aligned rules partition the space → piecewise-constant regions
CNN	useful local patterns repeat across positions → shared filters recognize translations

A prior is one explicit inductive bias; the model class, features, architecture, and optimization add others.

Step 1: the evidence gives the same likelihoods

D DERIVATION

Suppose we consider only two possible coins.

Fair coin F

$$P(H|F) = 0.5$$

Heads-biased coin B

$$P(H|B) = 0.9$$

Observed data: $H \cdot H \cdot H$

$$P(H, H, H|F) = 0.5^3 = 0.125$$

$$P(H, H, H|B) = 0.9^3 = 0.729$$

$$\text{likelihood ratio} = \frac{0.729}{0.125} \approx 5.83$$

HHH favours B over F by a factor of 5.83. These likelihoods do not depend on the prior.

Step 2: update a prior that favours the fair coin

D DERIVATION

Before seeing the flips, suppose our starting belief is $P(F) = 0.90$ and $P(B) = 0.10$.

Hypothesis	Prior	Likelihood	Prior $\times$ likelihood
Fair F	0.90	0.125	$0.90 \times$ $0.125 =$ 0.1125
Biased B	0.10	0.729	$0.10 \times$ $0.729 =$ 0.0729

$w_F = 0.1125$ and $w_B = 0.0729$ are **weights**, not probabilities yet.

Bayes first scores each hypothesis by prior $\times$ likelihood. The two scores need not add to 1.

Step 3: normalize the weights into probabilities

D DERIVATION

The normalizing constant is the sum of the two weights:

$$P(\mathcal{D}) = w_F + w_B = 0.1125 + 0.0729 = 0.1854.$$

Fair coin

$$P(F|\mathcal{D}) = \frac{0.1125}{0.1854} \approx 0.607$$

Heads-biased coin

$$P(B|\mathcal{D}) = \frac{0.0729}{0.1854} \approx 0.393$$

The normalized values now add to 1: $0.607 + 0.393 = 1$. They are the **posterior** probabilities—our beliefs after observing HHH.

HHH raises belief in B from 0.10 to 0.393, but it does not overturn the strong fair-coin prior.

Step 4: start instead from a biased-coin prior

D DERIVATION

Now use the **same hypotheses and the same HHH**, but start from $P(F) = 0.10$ and $P(B) = 0.90$.

Hypothesis	Prior	Likelihood	Weight
Fair F	0.10	0.125	0.0125
Biased B	0.90	0.729	0.6561

$$P(\mathcal{D}) = 0.0125 + 0.6561 = 0.6686$$

$$P(F|\mathcal{D}) \approx 0.019, \quad P(B|\mathcal{D}) \approx 0.981$$

The data again favours B by the same factor 5.83; this time the prior already favoured B .

What changed—and what stayed fixed?

Starting belief	Prior $P(B)$	Posterior $P(B H, H, H)$
Probably fair	0.10	0.393
Probably biased	0.90	0.981

Fixed: the observations and likelihoods. In both analyses, HHH multiplies the odds for B by 5.83.

Changed: the prior. Different starting beliefs therefore lead to different posterior probabilities.

$$\text{posterior odds} = \text{prior odds} \times \text{likelihood ratio}$$

Bayes does not erase the prior; it updates it with exactly the same evidence. More data makes the likelihood increasingly dominant.


A neural-network prior ranks plausible parameter vectors

$$\theta \sim p(\theta)$$

For a neural network, θ contains all weights and biases. Before using the current dataset, a prior states which settings are more plausible.

It can prefer small or sparse weights, smooth functions, and correlated parameters.

— narrow $\sigma = 0.55$ — broad $\sigma = 1.55$



narrow = stronger preference near 0


A prior makes part of the model's inductive bias explicit as a probability distribution.

Parameter priors have visible consequences

Example: a zero-centred prior makes parameter values near $\theta = 0$ more plausible before seeing data.


REGRESSION

$$Y_i \mid X_i = x_i \sim \mathcal{N}(\theta^\top \tilde{x}_i, \sigma^2)$$



BINARY CLASSIFICATION

$$p_{\theta}(Y = 1 \mid x) = \sigma(\theta^\top \tilde{x})$$




This conclusion depends on the prior: a prior centred elsewhere can prefer probabilities away from 0.5.

Model complexity: underfitting to overfitting

Fit the same eight observations by least squares; only the polynomial degree changes.


dashed = underlying trend · orange = fitted polynomial · red dot = one noisy observation



underfit



captures the main trend




overfit

Higher degree lowers training error, but degree 7 chases the noisy observation and becomes unstable. Priors and regularization can prefer smoother solutions.

Start with four noisy points from a straight line

The data were generated from $f_{\text{true}(x)} = 2 + 3x$ with small measurement noise.

x	$f_{\text{true}(x)}$	observed y
-1	-1	-0.8
0	2	1.8
1	5	5.3
2	8	7.9



grey line = generating rule · dots = observed data

The learner sees only the four observations and must infer the intercept b and slope w .

The likelihood peaks near the generating line

D DERIVATION


Model each observation as

$$Y_i|x_i, b, w \sim \mathcal{N}(b + wx_i, 0.5^2).$$

The data score each candidate line

$$L(b, w) \propto \exp\left(-\frac{1}{2(0.5)^2} \sum_i (y_i - b - wx_i)^2\right).$$

candidate (b, w)	squared error
(2, 3)	0.180
(0, 1)	56.580
MLE (2.07, 2.96)	0.162




The observed data make intercepts near 2 and slopes near 3 most plausible.

Before data, the prior prefers parameters near zero

Choose a simple Gaussian prior over the two regression parameters:

$$\theta = (b, w) \sim \mathcal{N}(\mathbf{0}, 0.5^2 \mathbf{I}).$$



What does it express?

$$-\log p(b, w) = \frac{1}{2(0.5)^2} (b^2 + w^2) + C = 2(b^2 + w^2) + C.$$


Before seeing this dataset:

- $b \approx 0$: intercept near zero
- $w \approx 0$: a flatter line
- farther from $(0, 0)$: less prior density

The likelihood says “fit these data”; the prior says “prefer smaller coefficients”.

Bayes combines the regression likelihood and prior

$$p(b, w \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid b, w) p(b, w)}{p(\mathcal{D})}$$



Contours: — prior — likelihood — posterior

● MLE (2.07, 2.96) ● posterior peak (1.79, 2.62) ● prior mean (0, 0)


The narrower prior visibly pulls the posterior peak away from the MLE and toward zero.

The prior changes the fitted line

— MLE, dashed — MAP, solid


Moderate prior: $\tau = 0.5$

$$\text{MAP: } \hat{y} = 1.79 + 2.62x$$



Much stronger prior: $\tau = 0.1$

$$\text{MAP: } \hat{y} = 0.44 + 0.68x$$



$$(b, w) : (2.07, 2.96)_{\text{MLE}} \rightarrow (1.79, 2.62)_{\text{moderate}} \rightarrow (0.44, 0.68)_{\text{strong}}$$

Reducing τ puts more prior mass near the origin and moves both MAP coefficients much closer to zero.

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta \mid \mathcal{D}) = \arg \max_{\theta} p(\mathcal{D} \mid \theta) p(\theta)$$

Take negative logs:

$$\hat{\theta}_{\text{MAP}} = \arg \min_{\theta} \left(\underbrace{-\log p(\mathcal{D} \mid \theta)}_{\text{NLL}} + \underbrace{-\log p(\theta)}_{\text{negative log-prior}} \right)$$

MAP objective: NLL + negative log-prior

In ML language: data loss + regularization.

A Gaussian prior is L_2 regularization

D DERIVATION

Assume the parameters are centred at zero before seeing the data:

$$\boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}).$$

$$\underbrace{-\log p(\mathcal{D} \mid \boldsymbol{\theta})}_{\text{NLL}} + \underbrace{-\log p(\boldsymbol{\theta})}_{\text{negative log-prior}}$$

$$-\log p(\boldsymbol{\theta}) = \frac{1}{2\tau^2} \|\boldsymbol{\theta}\|_2^2 + C$$

$$\Rightarrow \mathcal{L}_{\text{MAP}} = \text{NLL} + \underbrace{\frac{1}{2\tau^2}}_{\lambda} \|\boldsymbol{\theta}\|_2^2 + C$$

$\lambda = \frac{1}{2\tau^2}$: **smaller** $\tau \rightarrow$ **larger** $\lambda \rightarrow$ **stronger shrinkage toward 0.**

Gaussian MAP gives ridge shrinkage, step by step

D DERIVATION

For one coefficient, summarize the data likelihood by $z = \hat{\theta}_{\text{MLE}}$:

$$p(\mathcal{D}|\theta) \propto \exp\left(-\frac{(\theta-z)^2}{2s^2}\right), \quad p(\theta) \propto \exp\left(-\frac{\theta^2}{2\tau^2}\right).$$

$$J_2(\theta) = \underbrace{\frac{(\theta-z)^2}{2s^2}}_{\text{data NLL}} + \underbrace{\frac{\theta^2}{2\tau^2}}_{\text{Gaussian prior NLL}}.$$

$$\frac{dJ_2}{d\theta} = \frac{\theta - z}{s^2} + \frac{\theta}{\tau^2} = 0$$

$$\tau^2(\theta - z) + s^2\theta = 0 \quad \rightarrow \quad (\tau^2 + s^2)\theta = \tau^2 z$$

$$\hat{\theta}_{\text{MAP}} = \alpha z, \quad \alpha = \frac{\tau^2}{\tau^2 + s^2} \in (0, 1).$$

If $z = 0.6$ and $s = \tau$, then $\alpha = \frac{1}{2}$ and MAP = 0.3. Ridge reaches exactly zero only when $z = 0$.

Prior width controls how far MAP moves


Same likelihood: $L(\theta) \propto \mathcal{N}(\theta; \hat{\theta}_{\text{MLE}} = 2, s^2)$ with $s = 0.6$; prior: $p(\theta) = \mathcal{N}(0, \tau^2)$.

$$\hat{\theta}_{\text{MAP}} = \frac{\tau^2}{\tau^2 + s^2} \hat{\theta}_{\text{MLE}} \quad (\text{a shrinkage factor between 0 and 1})$$

— prior — likelihood — posterior


Broad prior: $\tau = 2, \lambda = 0.125$

MLE = 2 → MAP ≈ 1.83



Narrow prior: $\tau = 0.5, \lambda = 2$


MLE = 2 → MAP ≈ 0.82



The narrower prior gives the larger λ and pulls MAP farther from the MLE toward zero.

A multivariate Gaussian prior has geometry

$$\theta \sim \mathcal{N}(\mu, \Sigma) \quad -\log p(\theta) = \frac{1}{2}(\theta - \mu)^\top \Sigma^{-1}(\theta - \mu) + C$$




Axes: θ_1 and θ_2 are two coordinates of the parameter vector θ .

Colour = prior density $p(\theta)$: pale outer regions are less plausible; darker centres are more plausible.

Covariance says **which directions in parameter space are plausible.**

Laplace is sharper at zero than a Normal prior

Compare zero-centred priors with the **same variance**, 1:



Normal prior

- smooth around $\theta = 0$
- lighter tails

Laplace prior

- a sharp peak at $\theta = 0$
- heavier tails

The Laplace peak creates a **corner** in the negative log-prior—the source of L_1 sparsity.

Laplace MAP gives an L_1 objective

D DERIVATION

Use the same scalar likelihood, centred at $z = \hat{\theta}_{\text{MLE}}$, and a zero-centred Laplace prior with scale a :

$$p(\mathcal{D}|\theta) \propto e^{-\frac{(\theta-z)^2}{2s^2}}, \quad p(\theta) = \frac{1}{2a} e^{-\frac{|\theta|}{a}}.$$

$$-\log p(\theta) = \frac{|\theta|}{a} + \log(2a).$$

$$J_1(\theta) = \frac{(\theta - z)^2}{2s^2} + \frac{|\theta|}{a} + C.$$

Multiply by the positive constant s^2 ; the minimizer does not change:

$$\tilde{J}_1(\theta) = \frac{1}{2}(\theta - z)^2 + \lambda|\theta|, \quad \lambda = \frac{s^2}{a}.$$

The Laplace negative log-prior is proportional to $|\theta|$, so MAP becomes an L_1 -regularized optimization problem.

Soft thresholding follows from three cases

D DERIVATION

Differentiate $\tilde{J}_1(\theta) = \frac{1}{2}(\theta - z)^2 + \lambda|\theta|$ on either side of the corner.

Case 1: $\theta > 0$

$$\tilde{J}_{1'} = \theta - z + \lambda$$

$$0 = \theta - z + \lambda$$

$$\theta = z - \lambda, \text{ valid if } z > \lambda$$

Case 2: $\theta < 0$

$$\tilde{J}_{1'} = \theta - z - \lambda$$

$$0 = \theta - z - \lambda$$

$$\theta = z + \lambda, \text{ valid if } z < -\lambda$$

Case 3: $\theta = 0$

$$\partial|\theta|_0 = [-1, 1]$$

$$0 \in -z + \lambda[-1, 1]$$

$$\text{valid iff } |z| \leq \lambda$$

$$z > \lambda \\ \hat{\theta}_{\text{MAP}} = z - \lambda$$

$$|z| \leq \lambda \\ \hat{\theta}_{\text{MAP}} = 0$$

$$z < -\lambda \\ \hat{\theta}_{\text{MAP}} = z + \lambda$$

$$\hat{\theta}_{\text{MAP}} = \text{sign}(z) \max(|z| - \lambda, 0).$$

Values inside the threshold band $[-\lambda, \lambda]$ map exactly to zero; values outside are pulled toward zero by λ .

Ridge shrinks; L_1 can threshold to exactly zero


D DERIVATION

For one coefficient, let the data term be $D(\theta) = \frac{1}{2}(\theta - 0.6)^2$, centred at $z = 0.6$.

Normal prior $\rightarrow L_2$

$$J_2(\theta) = D(\theta) + \frac{1}{2}\theta^2$$


$$\hat{\theta}_{\text{MAP}} = \frac{z}{1+1} = 0.3$$



Laplace prior $\rightarrow L_1$

$$J_1(\theta) = D(\theta) + |\theta|$$


$$\hat{\theta}_{\text{MAP}} = \text{sign}(z) \max(|z| - 1, 0) = 0$$



Applied coefficient by coefficient, L_1 can make many MAP estimates **exactly zero: sparsity.**

Prior strength controls underfitting and overfitting

For the summed-NLL convention used above, $\lambda = \frac{1}{2\tau^2}$:



Regularization trades variance for bias; more is not always better.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/map-regularization

map-regularization.html — a two-parameter linear model so contours are visible. Shows **likelihood**, **prior**, and **posterior** contours with the **MLE**, **MAP**, and zero points.

Controls: amount of data · noise · prior variance · Gaussian vs Laplace · prior correlation.

Ask: more data? narrower prior? why does MAP → MLE as data grows? a correlated prior?

Full Bayes and deep learning

MAP is still one parameter vector


MLE $\hat{\theta}_{\text{MLE}}$ — one point

MAP $\hat{\theta}_{\text{MAP}}$ — one point

Full Bayes keeps the whole posterior $p(\theta \mid \mathcal{D})$ — a **distribution**, not a point.

Posterior predictive distribution

$$p(y^* \mid x^*, \mathcal{D}) = \int p_{\theta}(y^* \mid x^*) p(\theta \mid \mathcal{D}) d\theta$$



Bayesian prediction **averages over plausible models** — and reports uncertainty.

Why DL usually uses point estimates

★ OPTIONAL

A modern network has millions to billions of parameters, so integrating $p(\boldsymbol{\theta} \mid \mathcal{D})$ is hard. Common practical choices:

- MLE · MAP
- deep **ensembles**
- **dropout**-based approximations
- **variational** inference

The standard supervised DL objective

$$\min_{\boldsymbol{\theta}} \underbrace{\frac{1}{|B|} \sum_{i \in B} -\log p_{\boldsymbol{\theta}}(y_i \mid \mathbf{x}_i)}_{\text{likelihood / data fit}} + \underbrace{\lambda \|\boldsymbol{\theta}\|_2^2}_{\text{prior / regularization}}$$

With an averaged NLL, λ absorbs the dataset or minibatch scaling convention.


probabilistic model + neural network + gradient optimization

DL doesn't discard probabilistic ML — it **scales** it with expressive functions and gradient descent.

Gaussian likelihood
Categorical likelihood
Gaussian prior
Laplace prior
MLE + prior

⇒ **MSE**
⇒ **cross-entropy**
⇒ **L_2**
⇒ **L_1**
⇒ **MAP**

Every supervised objective begins with two assumptions



Loss $\leftrightarrow$ likelihood; regularizer $\leftrightarrow$ prior; prediction $\leftrightarrow$ plug-in or posterior predictive.

Next: from linear models to neural networks — **neurons, nonlinearities, and multilayer perceptrons.**