

Why These Losses?

A probabilistic refresher for Deep Learning

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

Three familiar equations

You have already met these losses:

$$\mathcal{L}_{\text{reg}} = \frac{1}{n} \sum_i (y_i - \hat{y}_i)^2 \quad \mathcal{L}_{\text{class}} = -\frac{1}{n} \sum_i \log p_{\theta}(y_i \mid x_i)$$

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{data}} + \lambda \|\theta\|_2^2$$

Why exactly these three expressions? Today: they are not arbitrary.

Which is it?

- Historical convention?
- Convenient gradients?
- Empirically successful?
- **Derived from probability?**
- All of the above?

Reveal. Primarily **consequences of probabilistic assumptions** — that is the thread for the whole lecture.

Assumption about outputs

$$p_{\theta}(y \mid x)$$

likelihood $\downarrow$ negative log-likelihood

training loss

Assumption about parameters

$$p(\theta)$$

Bayes' rule $\downarrow$ MAP

regularization

Minimal probability revision

A model should describe uncertainty

$$\hat{y} = f_{\theta}(x)$$

a single number

$$Y \mid X \{=\} x \sim p_{\theta}(\cdot \mid x)$$

a full distribution


- regression network → a **Gaussian mean**
- classifier → **categorical** probabilities
- language model → categorical distribution over **tokens**

A neural network parameterizes a probability distribution.

For a continuous Y , probability comes from **area**:

$$P(a \leq Y \leq b) = \int_a^b p(y) \, dy$$


$p(y) \neq P(Y\{=\}y)$ — a density can **exceed 1**; only *integrals* are probabilities.




$$Y \sim \mathcal{U}(a, b) \quad p(y) = \begin{cases} 1/(b-a) & a \leq y \leq b \\ 0 & \text{otherwise} \end{cases}$$

Q. NLL *outside* $[a, b]$? **A.** $-\log 0 = +\infty$.

A likelihood can impose a **hard constraint**.



$$Y \sim \mathcal{N}(\mu, \sigma^2) \quad p(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$



Only two knobs: **location** μ and **scale** σ .


Why the Gaussian produces a square

D


Take the negative log:

$$-\log p(y) = \underbrace{\frac{(y - \mu)^2}{2\sigma^2}}_{\text{depends on } \mu} + \underbrace{\frac{1}{2} \log(2\pi\sigma^2)}_{\text{constant in } \mu}$$

The only μ -dependent term is the **squared error** $(y - \mu)^2$.



$$\theta \sim \mathcal{N}(\mu, \Sigma) \quad -\log p(\theta) = \frac{1}{2}(\theta - \mu)^\top \Sigma^{-1}(\theta - \mu) + C$$



Covariance says **which directions in parameter space are plausible.**

Independence turns products into sums

Conditionally independent observations:

$$p(\mathcal{D} \mid \theta) = \prod_{i=1}^n p_{\theta}(y_i \mid x_i)$$

Take logs — the product becomes a **sum**:

$$\log p(\mathcal{D} \mid \theta) = \sum_{i=1}^n \log p_{\theta}(y_i \mid x_i)$$

Why a loss is an **average over examples**, and why minibatches work: $\frac{1}{|B|} \sum_{i \in B} \ell_i$.

Maximum likelihood → training losses

The one model $p_{\theta}(y \mid x)$, read two ways:

Fix θ , vary y

→ a **distribution** over outcomes.

Fix observed y , vary θ

→ a **likelihood** over parameters.

The likelihood is *not* “the probability that θ is correct.”

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} p(\mathcal{D} \mid \theta)$$

Interpretation. Choose the parameters under which the observed dataset would be **least surprising**.

$$\begin{aligned}\arg \max_{\theta} \prod_i p_{\theta}(y_i \mid x_i) &= \arg \max_{\theta} \sum_i \log p_{\theta}(y_i \mid x_i) \\ &= \arg \min_{\theta} \left(- \sum_i \log p_{\theta}(y_i \mid x_i) \right)\end{aligned}$$

$$\ell_i(\theta) = -\log p_{\theta}(y_i \mid x_i)$$

ML notation — empirical risk:

$$\hat{R}(\theta) = \frac{1}{n} \sum_i \ell(y_i, f_\theta(x_i))$$

Probabilistic reading of the per-example loss:

$$\ell(y_i, f_\theta(x_i)) = -\log p_\theta(y_i \mid x_i)$$


Choosing a loss = choosing a probabilistic observation model.

**Regression — where does MSE
come from?**

$$y_i = f_{\theta}(x_i) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$

Therefore the conditional output distribution is

$$Y_i \mid x_i, \theta \sim \mathcal{N}(f_{\theta}(x_i), \sigma^2)$$



Gaussian noise lives in the **output direction**, conditional on x — a vertical bell at every input.

$$p_{\theta}(y_i \mid x_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - f_{\theta}(x_i))^2}{2\sigma^2}\right)$$

Negative log of a single example:

$$-\log p_{\theta}(y_i \mid x_i) = \frac{(y_i - f_{\theta}(x_i))^2}{2\sigma^2} + C$$

Sum over the dataset:

$$-\log p(\mathcal{D} \mid \theta) = \frac{1}{2\sigma^2} \sum_i (y_i - f_\theta(x_i))^2 + C$$

With σ fixed, the minimizer is

$$\hat{\theta}_{\text{MLE}} = \arg \min_{\theta} \sum_i (y_i - f_\theta(x_i))^2$$

MSE = Gaussian NLL, up to constants.


$$\ell_i = \frac{(y_i - \mu_i)^2}{2\sigma^2} + \frac{1}{2} \log \sigma^2 + C$$

- small σ : errors penalized **strongly**
- large σ : errors deemed **less surprising**
- a predicted σ_i needs the $\frac{1}{2} \log \sigma_i^2$ term — it stops the model claiming infinite uncertainty

Heteroscedastic regression: let the network output both $(\mu_i, \log \sigma_i^2) = f_\theta(x_i)$.

Noise model	NLL
Gaussian	r^2
Laplace	$ r $
Uniform	hard interval
Student- t	robust

$$r = y - \hat{y}$$



INTERACTIVE

likelihood-to-loss.html — three synchronized panels: the **noise density** $p(r)$, the **loss** $-\log p(r)$, and **data with a fitted line**.

Controls: Gaussian / Laplace / Uniform / Student- t · noise scale · slope & intercept
· **add an outlier** · per-point contributions.


Ask: which model reacts most to an outlier? why does Uniform give an infinite barrier? which loss is robust?

Classification and cross-entropy

Classification should output probabilities

For K classes:

$$p_{\theta}(y\{=\}k \mid x) \geq 0, \quad \sum_{k=1}^K p_{\theta}(y\{=\}k \mid x) = 1$$



logits $z_k \in \mathbb{R} \rightarrow$ **probabilities** $p_k \in [0, 1]$.

$$Y \mid x \sim \text{Bernoulli}(p_\theta(x))$$

Compact one-line form:

$$p(y \mid x, \theta) = p^y (1 - p)^{1-y}$$

$$y \{=\} 1 \Rightarrow p(y) \{=\} p$$

$$y \{=\} 0 \Rightarrow p(y) \{=\} 1 - p$$

$$\begin{aligned} -\log p(y \mid x, \theta) &= -\log(p^y(1-p)^{1-y}) \\ &= -y \log p - (1-y) \log(1-p) \end{aligned}$$

Binary cross-entropy = Bernoulli NLL.

Why use logits?

V

The network emits an unbounded score $z = f_{\theta}(x) \in \mathbb{R}$; the **sigmoid** maps it to a probability:


$$p = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z \rightarrow -\infty : p \rightarrow 0$$

$$z = 0 : p = 0.5$$

$$z \rightarrow +\infty : p \rightarrow 1$$

In practice sigmoid + BCE are fused for numerical stability.



$$Y \mid x \sim \text{Categorical}(p_1, \dots, p_K)$$


For a one-hot target y :

$$p(y \mid x, \theta) = \prod_{k=1}^K p_k^{y_k}$$

From logits to softmax

$$p_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

Two guaranteed properties: $p_k > 0$ and $\sum_k p_k = 1$.



Shift invariance: $\text{softmax}(z) = \text{softmax}(z + c\mathbf{1})$.

$$-\log p(y \mid x, \theta) = -\log \prod_k p_k^{y_k} = -\sum_k y_k \log p_k$$

Because y is one-hot, only the true class survives:


$$\mathcal{L} = -\log p_{\text{true class}}$$

$$\ell(p_y) = -\log p_y$$

$$p_y = 0.9 \Rightarrow \ell \approx 0.105$$

$$p_y = 0.1 \Rightarrow \ell \approx 2.303$$


$$p_y = 0.001 \Rightarrow \ell \approx 6.908$$



Q. Why punish a *confidently wrong* prediction so hard?

For softmax followed by cross-entropy:

$$\frac{\partial \mathcal{L}}{\partial z_k} = p_k - y_k, \quad \nabla_z \mathcal{L} = p - y$$



Predicted p , target y , and their difference — a preview of backprop.

INTERACTIVE

[softmax-cross-entropy.html](#) — controls z_1, z_2, z_3 , the true class, and temperature T ; presets for **correct & confident** / **correct & unsure** / **wrong & unsure** / **wrong & confident**. Shows $z \rightarrow \text{softmax}(z/T) \rightarrow -\log p_y \rightarrow p - y$.

Ask students to predict the loss before it is revealed.

MSE is not *impossible* — it corresponds to the **wrong observation model** for class indicators.

MSE assumes $Y_k = p_k + \varepsilon_k, \varepsilon_k \sim \mathcal{N}(0, \sigma^2)$

Reality is $Y \sim \text{Categorical}(p)$

Cross-entropy also gives **stronger gradients** on confident errors (where MSE's gradient vanishes).

Bayes' rule, directly in classification


$$p(y\{=\}k \mid x) = \frac{p(x \mid y\{=\}k) p(y\{=\}k)}{p(x)}$$

Prediction ignores the constant denominator:

$$\hat{y} = \arg \max_k p(x \mid y\{=\}k) p(y\{=\}k)$$

- $p(y\{=\}k)$ — class **prior**
- $p(x \mid y\{=\}k)$ — class-conditional **likelihood**
- $p(y\{=\}k \mid x)$ — **posterior** class probability

$$X \mid Y\{=\}0 \sim \mathcal{N}(\mu_0, \sigma_0^2), \quad X \mid Y\{=\}1 \sim \mathcal{N}(\mu_1, \sigma_1^2)$$




Two class-conditional densities, a test point x^* , and its likelihood under each class.


Priors move the decision boundary

V

Balanced $p(Y \{=\} 1) = 0.5$



Rare $p(Y \{=\} 1) = 0.1$



$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$

INTERACTIVE


bayes-classifier.html — controls the means μ_0, μ_1 , the variances, the **prior** $p(Y\{=\}1)$, and a test point x^* . Shows $p(x \mid Y\{=\}0)$, $p(x \mid Y\{=\}1)$, the posterior $p(Y\{=\}1 \mid x)$, and the decision boundary together.

Try: balanced $\rightarrow$ drop the positive prior to 0.1 $\rightarrow$ why does the boundary move? then widen one variance.

Priors, MAP and regularization

The limitation of MLE

MLE asks only: *how well does θ explain the data?*



Many parameters fit a small dataset. **Which should we prefer?**

Put a prior over parameters

$$\theta \sim p(\theta)$$


Before seeing this dataset, which parameter values are plausible?

- small weights · sparse weights · smooth functions · correlated parameters

A prior is an **inductive bias**.

Bayes' rule over parameters

$$p(\theta \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \theta) p(\theta)}{p(\mathcal{D})}$$



$p(\mathcal{D}) = \int p(\mathcal{D} \mid \theta) p(\theta) d\theta$ — the evidence (we stop here).

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta \mid \mathcal{D}) = \arg \max_{\theta} p(\mathcal{D} \mid \theta) p(\theta)$$

Take negative logs:

$$\hat{\theta}_{\text{MAP}} = \arg \min_{\theta} (-\log p(\mathcal{D} \mid \theta) - \log p(\theta))$$

data loss + regularization

Assume $\theta \sim \mathcal{N}(0, \tau^2 I)$. Then

$$-\log p(\theta) = \frac{1}{2\tau^2} \theta^\top \theta + C$$

$$\Rightarrow \mathcal{L}_{\text{MAP}} = -\log p(\mathcal{D} \mid \theta) + \lambda \|\theta\|_2^2, \quad \lambda \propto \frac{1}{\tau^2}$$

narrow prior $\Rightarrow$ **stronger** regularization; **broad** prior $\Rightarrow$ **weaker**.

For $\theta \sim \mathcal{N}(0, \Sigma)$ the penalty is $\frac{1}{2}\theta^\top \Sigma^{-1}\theta$.

- isotropic $\Sigma \rightarrow$ ordinary L_2
- unequal variances $\rightarrow$ penalize parameters **differently**
- correlations $\rightarrow$ penalize particular **directions**

$$p(\theta_j) = \frac{1}{2b} \exp\left(-|\theta_j| \frac{1}{b}\right) \Rightarrow -\log p(\theta) = \frac{1}{b} \sum_j |\theta_j| + C$$

Laplace prior $\Rightarrow L_1$

Heavier tails at zero $\rightarrow$ **sparsity** (the lasso geometry, if time permits).

INTERACTIVE

[map-regularization.html](#) — a two-parameter linear model so contours are visible. Shows **likelihood**, **prior**, and **posterior** contours with the **MLE**, **MAP**, and zero points.

Controls: amount of data · noise · prior variance · Gaussian vs Laplace · prior correlation.

Ask: more data? narrower prior? why does MAP → MLE as data grows? a correlated prior?

Full Bayes and the deep-learning objective

MAP is still one parameter vector


MLE $\hat{\theta}_{\text{MLE}}$ — one point

MAP $\hat{\theta}_{\text{MAP}}$ — one point

Full Bayes keeps the whole posterior $p(\theta \mid \mathcal{D})$ — a **distribution**, not a point.

Posterior predictive distribution

$$p(y^* \mid x^*, \mathcal{D}) = \int p(y^* \mid x^*, \theta) p(\theta \mid \mathcal{D}) d\theta$$



Bayesian prediction **averages over plausible models** — and reports uncertainty.

A modern network has millions–billions of parameters, so integrating $p(\theta \mid \mathcal{D})$ is hard. Common practical choices:

- MLE · MAP
- deep **ensembles**
- **dropout**-based approximations
- **variational** inference

The standard supervised DL objective

$$\min_{\theta} \underbrace{\frac{1}{|B|} \sum_{i \in B} -\log p_{\theta}(y_i \mid x_i)}_{\text{likelihood / data fit}} + \underbrace{\lambda \|\theta\|_2^2}_{\text{prior / regularization}}$$

probabilistic model + neural network + gradient optimization

DL doesn't discard probabilistic ML — it **scales** it with expressive functions and gradient descent.

Complete each:

Gaussian likelihood $\Rightarrow$?

$\Rightarrow$ **MSE**

Categorical likelihood $\Rightarrow$?

$\Rightarrow$ **cross-entropy**

Gaussian prior $\Rightarrow$?

$\Rightarrow$ L_2

Laplace prior $\Rightarrow$?

$\Rightarrow$ L_1

MLE + prior $\Rightarrow$?

$\Rightarrow$ **MAP**

Final mental model

See a **loss** → ask “*what likelihood produced it?*”

See a **regularizer** → ask “*what prior produced it?*”

See a **prediction** → ask “*point estimate or posterior?*”

Next: computation graphs, gradients, and backpropagation.