

Derivatives, Gradients, Jacobians & Hessians

The calculus behind optimization and backpropagation

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

Why a calculus lecture?

Every step of training is a gradient step:

Why a calculus lecture?

Every step of training is a gradient step:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t)$$

Why a calculus lecture?

Every step of training is a gradient step:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t)$$

So we must be fluent with

Why a calculus lecture?

Every step of training is a gradient step:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t)$$

So we must be fluent with

$$\frac{d}{dx}, \quad \frac{\partial}{\partial x_i}, \quad \nabla f, \quad J, \quad H.$$

Why a calculus lecture?

Every step of training is a gradient step:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t)$$

So we must be fluent with

$$\frac{d}{dx}, \quad \frac{\partial}{\partial x_i}, \quad \nabla f, \quad J, \quad H.$$

These five objects are the entire vocabulary of optimization and backprop.

Five objects, one theme — each is a **derivative** of a function of a particular shape:

Five objects, one theme — each is a **derivative** of a function of a particular shape:


Object	Function type	Tells us
derivative $f'(x)$	$\mathbb{R} \rightarrow \mathbb{R}$	slope
partial $\frac{\partial f}{\partial x_i}$	$\mathbb{R}^d \rightarrow \mathbb{R}$	sensitivity to one coordinate
gradient ∇f	$\mathbb{R}^d \rightarrow \mathbb{R}$	steepest-ascent direction
Jacobian J	$\mathbb{R}^n \rightarrow \mathbb{R}^m$	local linear map
Hessian H	$\mathbb{R}^d \rightarrow \mathbb{R}$	curvature

The scalar derivative

$$f'(x) = \frac{df}{dx} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon) - f(x)}{\varepsilon}$$


Derivative as slope

$$f'(x) = \frac{df}{dx} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon) - f(x)}{\varepsilon}$$



Derivative as slope

$$f'(x) = \frac{df}{dx} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon) - f(x)}{\varepsilon}$$



the secant slope $\rightarrow$ the tangent slope as $\varepsilon \rightarrow 0$

The derivative is a local linear model

Near a point x :

The derivative is a local linear model

Near a point x :

$$f(x + \Delta x) \approx f(x) + f'(x) \Delta x$$

The derivative is a local linear model

Near a point x :

$$f(x + \Delta x) \approx f(x) + f'(x) \Delta x$$

derivative = best local linear approximation

The derivative is a local linear model

Near a point x :

$$f(x + \Delta x) \approx f(x) + f'(x) \Delta x$$

derivative = best local linear approximation

Example: $f(x) = x^2$ at $x = 3$, so $f(3) = 9$ and $f'(3) = 6$:

The derivative is a local linear model

Near a point x :

$$f(x + \Delta x) \approx f(x) + f'(x) \Delta x$$

derivative = best local linear approximation

Example: $f(x) = x^2$ at $x = 3$, so $f(3) = 9$ and $f'(3) = 6$:

$$f(3 + \Delta x) \approx 9 + 6 \Delta x.$$

Test the line $f(3 + \Delta x) \approx 9 + 6 \Delta x$ at a **small** step $\Delta x = 0.1$:

Test the line $f(3 + \Delta x) \approx 9 + 6 \Delta x$ at a **small** step $\Delta x = 0.1$: Linear prediction: $9 + 6 \cdot 0.1 = 9.6$.

Test the line $f(3 + \Delta x) \approx 9 + 6 \Delta x$ at a **small** step $\Delta x = 0.1$: Linear prediction: $9 + 6 \cdot 0.1 = 9.6$. True value: $f(3.1) = 3.1^2 = 9.61$.

Test the line $f(3 + \Delta x) \approx 9 + 6 \Delta x$ at a **small** step $\Delta x = 0.1$: Linear prediction: $9 + 6 \cdot 0.1 = 9.6$. True value: $f(3.1) = 3.1^2 = 9.61$. Error: $9.61 - 9.6 = 0.01$ — tiny.

Test the line $f(3 + \Delta x) \approx 9 + 6 \Delta x$ at a **small** step $\Delta x = 0.1$: Linear prediction: $9 + 6 \cdot 0.1 = 9.6$. True value: $f(3.1) = 3.1^2 = 9.61$. Error: $9.61 - 9.6 = 0.01$ — tiny.

Now a **large** step $\Delta x = 1$: the line gives $9 + 6 = 15$, but $f(4) = 16$ — error 1. The linear model degrades as Δx grows.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/derivative-tangent

drag the point and Δx ; watch the tangent, the secant, and the approximation error. Functions: x^2 , $\sin x$, e^x , $\log(1 + e^x)$.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/derivative-tangent

drag the point and Δx ; watch the tangent, the secant, and the approximation error. Functions: x^2 , $\sin x$, e^x , $\log(1 + e^x)$.

Explore how the approximation error grows as the step leaves the local tangent regime.

When is a first-order linear approximation most reliable?

A. For a large step on a sharply curved function

B. For a small step where curvature is modest

C. Only at a global minimum

D. Whenever the gradient is zero


Correct: B — For a small step where curvature is modest

Why: The derivative captures local slope; omitted higher-order terms become important over large steps or sharp bends.


Partial derivatives and surfaces

A function of two inputs is a **surface**: $f(x, y) = x^2 + y^2$.

A function of two inputs is a **surface**: $f(x, y) = x^2 + y^2$.



A function of two inputs is a **surface**: $f(x, y) = x^2 + y^2$.



input $(x, y) \rightarrow$ height $z = f(x, y)$

Partial derivatives

$\frac{\partial f}{\partial x}$ — change x , **hold y fixed**. $\frac{\partial f}{\partial y}$ — change y , hold x fixed.

Partial derivatives

$\frac{\partial f}{\partial x}$ — change x , **hold y fixed**. $\frac{\partial f}{\partial y}$ — change y , hold x fixed. For $f(x, y) = x^2 + y^2$:

Partial derivatives

$\frac{\partial f}{\partial x}$ — change x , **hold y fixed**. $\frac{\partial f}{\partial y}$ — change y , hold x fixed. For $f(x, y) = x^2 + y^2$:

$$\frac{\partial f}{\partial x} = 2x, \quad \frac{\partial f}{\partial y} = 2y$$

Partial derivatives

$\frac{\partial f}{\partial x}$ — change x , **hold y fixed**. $\frac{\partial f}{\partial y}$ — change y , hold x fixed. For $f(x, y) = x^2 + y^2$:

$$\frac{\partial f}{\partial x} = 2x, \quad \frac{\partial f}{\partial y} = 2y$$

$$\text{At } (x, y) = (3, 1): \quad \frac{\partial f}{\partial x} = 6, \quad \frac{\partial f}{\partial y} = 2.$$


A partial is the slope of a slice

Fix one input and you are left with a 1-D curve — the partial is **its** slope.

A partial is the slope of a slice

Fix one input and you are left with a 1-D curve — the partial is **its** slope.


slices = partial derivatives



A partial is the slope of a slice

Fix one input and you are left with a 1-D curve — the partial is **its** slope.

slices = partial derivatives



For $f = x^2 + 3y^2$ at $(1, 2)$: $\frac{\partial f}{\partial x} = 2$, $\frac{\partial f}{\partial y} = 12$ — the function climbs **faster** along y .

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/surface-partials

move a point on a 3-D surface and read off the two slice slopes. Functions: $x^2 + y^2$, $x^2 + 3y^2$, $\sin x \cos y$.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/surface-partials

move a point on a 3-D surface and read off the two slice slopes. Functions: $x^2 + y^2$, $x^2 + 3y^2$, $\sin x \cos y$.

Use the two slice slopes to compare the axes at a chosen point.

The gradient

Gradient: stack the partials

$$\nabla f(x) = \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_d} \end{pmatrix}$$

Gradient: stack the partials

$$\nabla f(x) = \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_d} \end{pmatrix}$$

For $f(x, y) = x^2 + y^2$: $\nabla f = \begin{pmatrix} 2x \\ 2y \end{pmatrix}$.

Gradient: stack the partials

$$\nabla f(x) = \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_d} \end{pmatrix}$$

For $f(x, y) = x^2 + y^2$: $\nabla f = \begin{pmatrix} 2x \\ 2y \end{pmatrix}$. At $(3, 1)$: $\nabla f = \begin{pmatrix} 6 \\ 2 \end{pmatrix}$ — pointing **outward**, uphill from the bowl's centre.

Same first-order Taylor idea, now with a **dot product**:

Same first-order Taylor idea, now with a **dot product**:

$$f(x + \Delta x) \approx f(x) + \nabla f(x)^\top \Delta x$$

Same first-order Taylor idea, now with a **dot product**:

$$f(x + \Delta x) \approx f(x) + \nabla f(x)^\top \Delta x$$

the vector version of $f(x + \Delta x) \approx f(x) + f'(x)\Delta x$

Directional derivative

For a **unit** direction u :

Directional derivative

For a **unit** direction u :

$$D_u f(x) = \nabla f(x)^\top u$$

Directional derivative

For a **unit** direction u :

$$D_u f(x) = \nabla f(x)^\top u$$

The rate of change along u is just the **projection** of the gradient onto u .

Directional derivative

For a **unit** direction u :

$$D_u f(x) = \nabla f(x)^\top u$$

The rate of change along u is just the **projection** of the gradient onto u .

Example: $f = x^2 + y^2$ at $(1, 2)$, so $\nabla f = \begin{pmatrix} 2 \\ 4 \end{pmatrix}$. Take $u = (3, 4)/5 = (0.6, 0.8)$:

Directional derivative

For a **unit** direction u :

$$D_u f(x) = \nabla f(x)^\top u$$

The rate of change along u is just the **projection** of the gradient onto u .

Example: $f = x^2 + y^2$ at $(1, 2)$, so $\nabla f = \begin{pmatrix} 2 \\ 4 \end{pmatrix}$. Take $u = (3, 4)/5 = (0.6, 0.8)$:

$$D_u f = 2 \cdot 0.6 + 4 \cdot 0.8 = 1.2 + 3.2 = 4.4$$

Directional derivative

For a **unit** direction u :

$$D_u f(x) = \nabla f(x)^\top u$$

The rate of change along u is just the **projection** of the gradient onto u .

Example: $f = x^2 + y^2$ at $(1, 2)$, so $\nabla f = \begin{pmatrix} 2 \\ 4 \end{pmatrix}$. Take $u = (3, 4)/5 = (0.6, 0.8)$:

$$D_u f = 2 \cdot 0.6 + 4 \cdot 0.8 = 1.2 + 3.2 = 4.4$$

compare $\|\nabla f\| = \sqrt{20} \approx 4.47$ – the **gradient** direction is a touch steeper.

The gradient points uphill

By Cauchy–Schwarz, over all unit u :

The gradient points uphill

By Cauchy–Schwarz, over all unit u :

$$\nabla f(x)^\top u \leq \|\nabla f(x)\|$$

The gradient points uphill

By Cauchy–Schwarz, over all unit u :

$$\nabla f(x)^\top u \leq \|\nabla f(x)\|$$

with equality at $u = \nabla f(x) / \|\nabla f(x)\|$.

The gradient points uphill

By Cauchy–Schwarz, over all unit u :

$$\nabla f(x)^\top u \leq \|\nabla f(x)\|$$

with equality at $u = \nabla f(x) / \|\nabla f(x)\|$.


∇f = steepest ascent • $-\nabla f$ = steepest descent

The gradient is perpendicular to contours

A contour is the set $f(x, y) = c$. Along it, f does not change, so for any tangent v :
 $\nabla f^\top v = 0$.


The gradient is perpendicular to contours

A contour is the set $f(x, y) = c$. Along it, f does not change, so for any tangent v : $\nabla f^\top v = 0$.



The gradient is perpendicular to contours

A contour is the set $f(x, y) = c$. Along it, f does not change, so for any tangent v : $\nabla f^\top v = 0$.



so ∇f is **orthogonal** to the contour — it points straight across the level lines.

Gradient descent on the contour map

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

Gradient descent on the contour map

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

Example: at $x_t = (3, 1)$ with $\nabla f = (6, 2)$ and $\eta = 0.1$:

Gradient descent on the contour map

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

Example: at $x_t = (3, 1)$ with $\nabla f = (6, 2)$ and $\eta = 0.1$:

$$x_{t+1} = (3, 1) - 0.1(6, 2) = (2.4, 0.8)$$

Gradient descent on the contour map

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

Example: at $x_t = (3, 1)$ with $\nabla f = (6, 2)$ and $\eta = 0.1$:

$$x_{t+1} = (3, 1) - 0.1(6, 2) = (2.4, 0.8)$$

Each step moves **perpendicular** to the current contour, downhill. The learning rate η sets the step length.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/optimizer-race

Watch gradient arrows on a contour map and a descent trajectory as you change η and the start point. Surfaces: $x^2 + y^2$, $x^2 + 10y^2$, $\sin x + \cos y$.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/optimizer-race

Watch gradient arrows on a contour map and a descent trajectory as you change η and the start point. Surfaces: $x^2 + y^2$, $x^2 + 10y^2$, $\sin x + \cos y$.

Move the start point to the bottom of the surface and inspect the update.

At a differentiable local minimum, what does a plain gradient-descent step do?

A. Moves uphill

B. Leaves the point unchanged because $\nabla f = 0$

C. Doubles the learning rate

D. Must move along a contour

Correct: B — Leaves the point unchanged

Why: At a differentiable local minimum the gradient is zero, so $x_{\{t+1\}} = x_t - \eta \nabla f(x_t) = x_t$.

The Hessian and curvature

Second derivative in 1-D

$$f''(x) = \frac{d^2 f}{dx^2} = \text{how fast the slope changes}$$

Second derivative in 1-D

$$f''(x) = \frac{d^2 f}{dx^2} = \text{how fast the slope changes}$$

$$f(x) = x^2 \Rightarrow f'' = 2 \text{ (bowl, curves \textcolor{brown}{up})};$$

Second derivative in 1-D


$$f''(x) = \frac{d^2 f}{dx^2} = \text{how fast the slope changes}$$

$$f(x) = x^2 \Rightarrow f'' = 2 \text{ (bowl, curves \textbf{up})}; \quad f(x) = -x^2 \Rightarrow f'' = -2 \text{ (hill, curves \textbf{down})}.$$


The **sign** of f'' names the shape:

The **sign** of f'' names the shape:


$f'' > 0$ convex



$f'' < 0$ concave






$f'' = 0$ flat



Convex, concave, flat

The **sign** of f'' names the shape:



$f'' > 0$ convex (bowl) · $f'' < 0$ concave (hill) · $f'' = 0$ flat (inflection)

The Hessian: matrix of second partials

$$H = \nabla^2 f, \quad H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$$

The Hessian: matrix of second partials

$$H = \nabla^2 f, \quad H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$$

For two variables:

The Hessian: matrix of second partials

$$H = \nabla^2 f, \quad H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$$

For two variables:

$$H = \begin{pmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial y \partial x} & \frac{\partial^2 f}{\partial y^2} \end{pmatrix}$$

$$f(x, y) = x^2 + 3y^2 \implies \nabla f = \begin{pmatrix} 2x \\ 6y \end{pmatrix}$$

$$f(x, y) = x^2 + 3y^2 \implies \nabla f = \begin{pmatrix} 2x \\ 6y \end{pmatrix}$$

Differentiate each partial again, one entry at a time:

$$f(x, y) = x^2 + 3y^2 \implies \nabla f = \begin{pmatrix} 2x \\ 6y \end{pmatrix}$$

Differentiate each partial again, one entry at a time:

$$\partial_{xx} = \frac{\partial}{\partial x}(2x) = 2$$

$$f(x, y) = x^2 + 3y^2 \implies \nabla f = \begin{pmatrix} 2x \\ 6y \end{pmatrix}$$

Differentiate each partial again, one entry at a time:

$$\partial_{xx} = \frac{\partial}{\partial x}(2x) = 2$$

$$\partial_{yy} = \frac{\partial}{\partial y}(6y) = 6, \quad \partial_{xy} = \frac{\partial}{\partial y}(2x) = 0$$

$$f(x, y) = x^2 + 3y^2 \implies \nabla f = \begin{pmatrix} 2x \\ 6y \end{pmatrix}$$

Differentiate each partial again, one entry at a time:

$$\partial_{xx} = \frac{\partial}{\partial x}(2x) = 2$$

$$\partial_{yy} = \frac{\partial}{\partial y}(6y) = 6, \quad \partial_{xy} = \frac{\partial}{\partial y}(2x) = 0$$

$$H = \begin{pmatrix} 2 & 0 \\ 0 & 6 \end{pmatrix}$$

$$f(x, y) = x^2 + 3y^2 \implies \nabla f = \begin{pmatrix} 2x \\ 6y \end{pmatrix}$$

Differentiate each partial again, one entry at a time:

$$\partial_{xx} = \frac{\partial}{\partial x}(2x) = 2$$

$$\partial_{yy} = \frac{\partial}{\partial y}(6y) = 6, \quad \partial_{xy} = \frac{\partial}{\partial y}(2x) = 0$$

$$H = \begin{pmatrix} 2 & 0 \\ 0 & 6 \end{pmatrix}$$

Curvature is **larger** in the y direction — the bowl is steeper that way.

The full second-order Taylor expansion:

The full second-order Taylor expansion:

$$f(x + \Delta x) \approx f(x) + \underbrace{\nabla f(x)^\top \Delta x}_{\text{tilt (gradient)}} + \underbrace{\frac{1}{2} \Delta x^\top H(x) \Delta x}_{\text{curvature (Hessian)}}$$

The full second-order Taylor expansion:

$$f(x + \Delta x) \approx f(x) + \underbrace{\nabla f(x)^\top \Delta x}_{\text{tilt (gradient)}} + \underbrace{\frac{1}{2} \Delta x^\top H(x) \Delta x}_{\text{curvature (Hessian)}}$$

gradient tilts · Hessian bends

For a quadratic $f(x) = \frac{1}{2}x^\top Ax$: $\nabla f = Ax$, $H = A$.

For a quadratic $f(x) = \frac{1}{2}x^\top Ax$: $\nabla f = Ax$, $H = A$.

- **eigenvectors** of $H \rightarrow$ the curvature **directions**;

For a quadratic $f(x) = \frac{1}{2}x^\top Ax$: $\nabla f = Ax$, $H = A$.

- **eigenvectors** of $H \rightarrow$ the curvature **directions**;
- **eigenvalues** of $H \rightarrow$ the curvature **magnitudes**.

Classifying critical points

At a point where $\nabla f = 0$, the Hessian eigenvalues decide the shape:

Classifying critical points

At a point where $\nabla f = 0$, the Hessian eigenvalues decide the shape:

Hessian eigenvalues	Local shape
all positive	local minimum
all negative	local maximum
mixed signs	saddle point
some zero	inconclusive / flat

Classifying critical points

At a point where $\nabla f = 0$, the Hessian eigenvalues decide the shape:

Hessian eigenvalues	Local shape
all positive	local minimum
all negative	local maximum
mixed signs	saddle point
some zero	inconclusive / flat

the eigenvalue **signs** alone classify the point.

A saddle point

$$f(x, y) = x^2 - y^2, \quad \nabla f = \begin{pmatrix} 2x \\ -2y \end{pmatrix}, \quad H = \begin{pmatrix} 2 & 0 \\ 0 & -2 \end{pmatrix}$$

$$f(x, y) = x^2 - y^2, \quad \nabla f = \begin{pmatrix} 2x \\ -2y \end{pmatrix}, \quad H = \begin{pmatrix} 2 & 0 \\ 0 & -2 \end{pmatrix}$$


Eigenvalues $+2$ and -2 have **mixed signs** $\rightarrow$ a saddle.

A saddle point

$$f(x, y) = x^2 - y^2, \quad \nabla f = \begin{pmatrix} 2x \\ -2y \end{pmatrix}, \quad H = \begin{pmatrix} 2 & 0 \\ 0 & -2 \end{pmatrix}$$

Eigenvalues $+2$ and -2 have **mixed signs** $\rightarrow$ a saddle.

$f = x^2 - y^2$ (saddle)




A saddle point

$$f(x, y) = x^2 - y^2, \quad \nabla f = \begin{pmatrix} 2x \\ -2y \end{pmatrix}, \quad H = \begin{pmatrix} 2 & 0 \\ 0 & -2 \end{pmatrix}$$

Eigenvalues $+2$ and -2 have **mixed signs** $\rightarrow$ a saddle.

$$f = x^2 - y^2 \text{ (saddle)}$$





$\nabla f = 0$ at the origin, but it is **not** a minimum

Why curvature matters for optimization

The gradient picks a **direction**, but curvature decides a good **step size**.


Why curvature matters for optimization

The gradient picks a **direction**, but curvature decides a good **step size**.



Why curvature matters for optimization

The gradient picks a **direction**, but curvature decides a good **step size**.



Elongated contours (a large **condition number**) make first-order gradient descent **zigzag** — much curvature one way, little the other.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/optimizer-race

Reshape a quadratic $ax^2 + by^2 + cxy$, see its Hessian eigenvectors, and watch the descent path zigzag.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/optimizer-race

Reshape a quadratic $ax^2 + by^2 + cxy$, see its Hessian eigenvectors, and watch the descent path zigzag.

Observe how one shared step size alternately overshoots the steep direction and creeps along the flat one.

Follow-along: exact quadratic GD Colab — predict both trajectories before running.

The Jacobian

When the gradient is not enough

The gradient is for **scalar** outputs, $f : \mathbb{R}^d \rightarrow \mathbb{R}$.

When the gradient is not enough

The gradient is for **scalar** outputs, $f : \mathbb{R}^d \rightarrow \mathbb{R}$. But network layers map **vectors to vectors**, $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ — e.g. $z = Wx + b$.

When the gradient is not enough

The gradient is for **scalar** outputs, $f : \mathbb{R}^d \rightarrow \mathbb{R}$. But network layers map **vectors to vectors**, $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ — e.g. $z = Wx + b$. We need the **Jacobian**.

Jacobian: matrix of all partials

For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$:

Jacobian: matrix of all partials

For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$:

$$J = \frac{\partial f}{\partial x} \in \mathbb{R}^{m \times n}, \quad J_{ij} = \frac{\partial f_i}{\partial x_j}$$

Jacobian: matrix of all partials

For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$:

$$J = \frac{\partial f}{\partial x} \in \mathbb{R}^{m \times n}, \quad J_{ij} = \frac{\partial f_i}{\partial x_j}$$

$$f(x + \Delta x) \approx f(x) + J(x) \Delta x \text{ — a local linear map}$$

Two Jacobians you already know

D DERIVATION

Two building blocks whose Jacobians you can write down instantly:

Two Jacobians you already know

Two building blocks whose Jacobians you can write down instantly:

Linear layer $f(x) = Wx + b$:

$$J = W$$

$$\text{since } \frac{\partial z_i}{\partial x_j} = W_{ij}.$$

Elementwise $h = \varphi(z)$:

$$J = \text{diag}(\varphi'(z_1), \dots, \varphi'(z_m))$$

(for ReLU, entries are 0 or 1).

Worked example: a $2 \rightarrow 2$ Jacobian

D DERIVATION

$f(x_1, x_2) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_1 x_2) \end{pmatrix}$ — the map from the PyTorch demo later.

Worked example: a $2 \rightarrow 2$ Jacobian

D DERIVATION

$f(x_1, x_2) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_1 x_2) \end{pmatrix}$ — the map from the PyTorch demo later. Row 1, partials of $f_1 = x_1^2 + x_2$: $\frac{\partial f_1}{\partial x_1} = 2x_1$, $\frac{\partial f_1}{\partial x_2} = 1$.

Worked example: a $2 \rightarrow 2$ Jacobian

D DERIVATION

$f(x_1, x_2) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_1 x_2) \end{pmatrix}$ — the map from the PyTorch demo later. Row 1, partials of $f_1 = x_1^2 + x_2$: $\frac{\partial f_1}{\partial x_1} = 2x_1$, $\frac{\partial f_1}{\partial x_2} = 1$. Row 2, partials of $f_2 = \sin(x_1 x_2)$: $\frac{\partial f_2}{\partial x_1} = x_2 \cos(x_1 x_2)$, $\frac{\partial f_2}{\partial x_2} = x_1 \cos(x_1 x_2)$.

Worked example: a $2 \rightarrow 2$ Jacobian

$f(x_1, x_2) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_1 x_2) \end{pmatrix}$ — the map from the PyTorch demo later. Row 1, partials of $f_1 = x_1^2 + x_2$: $\frac{\partial f_1}{\partial x_1} = 2x_1$, $\frac{\partial f_1}{\partial x_2} = 1$. Row 2, partials of $f_2 = \sin(x_1 x_2)$: $\frac{\partial f_2}{\partial x_1} = x_2 \cos(x_1 x_2)$, $\frac{\partial f_2}{\partial x_2} = x_1 \cos(x_1 x_2)$.

$$J = \begin{pmatrix} 2x_1 & 1 \\ x_2 \cos(x_1 x_2) & x_1 \cos(x_1 x_2) \end{pmatrix}$$

Worked example: a $2 \rightarrow 2$ Jacobian

$f(x_1, x_2) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_1 x_2) \end{pmatrix}$ — the map from the PyTorch demo later. Row 1, partials of $f_1 = x_1^2 + x_2$: $\frac{\partial f_1}{\partial x_1} = 2x_1$, $\frac{\partial f_1}{\partial x_2} = 1$. Row 2, partials of $f_2 = \sin(x_1 x_2)$: $\frac{\partial f_2}{\partial x_1} = x_2 \cos(x_1 x_2)$, $\frac{\partial f_2}{\partial x_2} = x_1 \cos(x_1 x_2)$.

$$J = \begin{pmatrix} 2x_1 & 1 \\ x_2 \cos(x_1 x_2) & x_1 \cos(x_1 x_2) \end{pmatrix}$$

$$\text{At } (x_1, x_2) = (1, 2): J = \begin{pmatrix} 2 & 1 \\ 2 \cos 2 & \cos 2 \end{pmatrix}.$$

Worked example: a $2 \rightarrow 2$ Jacobian

D DERIVATION


$f(x_1, x_2) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_1 x_2) \end{pmatrix}$ — the map from the PyTorch demo later. Row 1, partials of $f_1 = x_1^2 + x_2$: $\frac{\partial f_1}{\partial x_1} = 2x_1$, $\frac{\partial f_1}{\partial x_2} = 1$. Row 2, partials of $f_2 = \sin(x_1 x_2)$: $\frac{\partial f_2}{\partial x_1} = x_2 \cos(x_1 x_2)$, $\frac{\partial f_2}{\partial x_2} = x_1 \cos(x_1 x_2)$.


$$J = \begin{pmatrix} 2x_1 & 1 \\ x_2 \cos(x_1 x_2) & x_1 \cos(x_1 x_2) \end{pmatrix}$$

$$\text{At } (x_1, x_2) = (1, 2): J = \begin{pmatrix} 2 & 1 \\ 2 \cos 2 & \cos 2 \end{pmatrix}.$$


Each entry is one partial — a 2×2 local linear map, built without ever storing a big matrix.

The chain rule, with Jacobians






$$\frac{\partial v}{\partial x} = \frac{\partial v}{\partial u} \frac{\partial u}{\partial x} \quad \Rightarrow \quad J_{v,x} = J_{v,u} J_{u,x}$$



$$\frac{\partial v}{\partial x} = \frac{\partial v}{\partial u} \frac{\partial u}{\partial x} \implies J_{v,x} = J_{v,u} J_{u,x}$$


Shapes: with $x \in \mathbb{R}^2$, $u \in \mathbb{R}^3$, $v \in \mathbb{R}^2$, then $J_{u,x} \in \mathbb{R}^{3 \times 2}$ and $J_{v,u} \in \mathbb{R}^{2 \times 3}$:



$$\frac{\partial v}{\partial x} = \frac{\partial v}{\partial u} \frac{\partial u}{\partial x} \implies J_{v,x} = J_{v,u} J_{u,x}$$

Shapes: with $x \in \mathbb{R}^2$, $u \in \mathbb{R}^3$, $v \in \mathbb{R}^2$, then $J_{u,x} \in \mathbb{R}^{3 \times 2}$ and $J_{v,u} \in \mathbb{R}^{2 \times 3}$:

$$J_{v,x} = J_{v,u} J_{u,x} \in \mathbb{R}^{2 \times 2}$$



$$\frac{\partial v}{\partial x} = \frac{\partial v}{\partial u} \frac{\partial u}{\partial x} \implies J_{v,x} = J_{v,u} J_{u,x}$$

Shapes: with $x \in \mathbb{R}^2$, $u \in \mathbb{R}^3$, $v \in \mathbb{R}^2$, then $J_{u,x} \in \mathbb{R}^{3 \times 2}$ and $J_{v,u} \in \mathbb{R}^{2 \times 3}$:

$$J_{v,x} = J_{v,u} J_{u,x} \in \mathbb{R}^{2 \times 2}$$

The inner dimension (3) cancels — exactly like ordinary matrix multiplication.

Why we avoid full Jacobians

★ OPTIONAL

A single layer with $x, h \in \mathbb{R}^{4096}$ has a Jacobian of $4096 \times 4096 \approx 1.7 \times 10^7$ entries.

Why we avoid full Jacobians

★ OPTIONAL

A single layer with $x, h \in \mathbb{R}^{4096}$ has a Jacobian of $4096 \times 4096 \approx 1.7 \times 10^7$ entries. That is one dense matrix **per layer, per example** — far too big to form and multiply.

Why we avoid full Jacobians

★ OPTIONAL

A single layer with $x, h \in \mathbb{R}^{4096}$ has a Jacobian of $4096 \times 4096 \approx 1.7 \times 10^7$ entries. That is one dense matrix **per layer, per example** — far too big to form and multiply.

Frameworks never materialize J . They compute **vector–Jacobian products** instead.

If $y = f(x)$ and $\mathcal{L} = \mathcal{L}(y)$:

If $y = f(x)$ and $\mathcal{L} = \mathcal{L}(y)$:

$$\underbrace{\bar{x}}_{\frac{\partial \mathcal{L}}{\partial x}} = J^\top \underbrace{\bar{y}}_{\frac{\partial \mathcal{L}}{\partial y}}$$

If $y = f(x)$ and $\mathcal{L} = \mathcal{L}(y)$:

$$\underbrace{\bar{x}}_{\frac{\partial \mathcal{L}}{\partial x}} = J^\top \underbrace{\bar{y}}_{\frac{\partial \mathcal{L}}{\partial y}}$$

one VJP per layer — no matrix ever formed

Backprop is a chain of VJPs


For $x \rightarrow h_1 \rightarrow h_2 \rightarrow \mathcal{L}$:

$$\nabla_x \mathcal{L} = J_{h_1, x}^\top J_{h_2, h_1}^\top \nabla_{h_2} \mathcal{L}$$

Backprop is a chain of VJPs

For $x \rightarrow h_1 \rightarrow h_2 \rightarrow \mathcal{L}$:


$$\nabla_x \mathcal{L} = J_{h_1, x}^\top J_{h_2, h_1}^\top \nabla_{h_2} \mathcal{L}$$



Backprop is a chain of VJPs

For $x \rightarrow h_1 \rightarrow h_2 \rightarrow \mathcal{L}$:

$$\nabla_x \mathcal{L} = J_{h_1, x}^\top J_{h_2, h_1}^\top \nabla_{h_2} \mathcal{L}$$



Evaluate **right to left**: $\bar{h}_2 \rightarrow \bar{h}_1 \rightarrow \bar{x}$ — vectors only.

INTERACTIVE


↗ nipunbatra.github.io/interactive-articles/jacobian-local-map

a 2-D map warps a grid; the local Jacobian sends a small circle to an ellipse.

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/jacobian-local-map

a 2-D map warps a grid; the local Jacobian sends a small circle to an ellipse.



Where each object shows up in deep learning

Derivative → activations

Scalar activation derivatives drive backprop through nonlinearities:

Derivative → activations

Scalar activation derivatives drive backprop through nonlinearities:

$$\frac{d}{dz}\sigma(z) = \sigma(z)(1 - \sigma(z))$$

Derivative → activations

Scalar activation derivatives drive backprop through nonlinearities:

$$\frac{d}{dz}\sigma(z) = \sigma(z)(1 - \sigma(z))$$

$$\text{ReLU}'(z) = \mathbb{1}[z > 0]$$

Gradient → learning & attribution

- parameter updates: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$;

Gradient → learning & attribution

- parameter updates: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$;
- saliency maps: $\nabla_x f_y(x)$;

Gradient → learning & attribution

- parameter updates: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$;
- saliency maps: $\nabla_x f_y(x)$;
- adversarial examples: $x_{\text{adv}} = x + \varepsilon \text{sign}(\nabla_x \mathcal{L})$.

Appears in backprop, sensitivity analysis, normalizing flows ($\log|\det J|$), neural ODEs, implicit layers.

Appears in backprop, sensitivity analysis, normalizing flows ($\log|\det J|$), neural ODEs, implicit layers.

But almost always accessed through a product Jv or $J^\top v$ — never the full matrix.

Newton's method $x_{t+1} = x_t - H^{-1} \nabla f$, Laplace approximation, sharpness / flatness, influence functions.

Newton's method $x_{t+1} = x_t - H^{-1} \nabla f$, Laplace approximation, sharpness / flatness, influence functions.

But the full H is usually too large — we use Hessian–vector products Hv instead.

Why DL stays first-order

For P parameters, $\nabla_{\theta}\mathcal{L} \in \mathbb{R}^P$ but $H \in \mathbb{R}^{P \times P}$.

Why DL stays first-order

For P parameters, $\nabla_{\theta}\mathcal{L} \in \mathbb{R}^P$ but $H \in \mathbb{R}^{P \times P}$.

$P = 10^8 \implies H$ **has** 10^{16} **entries**

Why DL stays first-order

For P parameters, $\nabla_{\theta} \mathcal{L} \in \mathbb{R}^P$ but $H \in \mathbb{R}^{P \times P}$.

$P = 10^8 \implies H$ **has** 10^{16} **entries**

So practice uses SGD, momentum, Adam, and diagonal / Hessian-vector approximations.

The autodiff connection

Autodiff is none of the naive options

not symbolic differentiation

not finite differences

Autodiff is none of the naive options

not symbolic differentiation

not finite differences

Automatic differentiation applies the chain rule to the **executed program** — exact, and cheap.

Scalar loss $\rightarrow$ reverse mode

Training maps $\theta \in \mathbb{R}^P \rightarrow \mathcal{L} \in \mathbb{R}$ (many inputs, one output).

Scalar loss $\rightarrow$ reverse mode

Training maps $\theta \in \mathbb{R}^P \rightarrow \mathcal{L} \in \mathbb{R}$ (many inputs, one output).

reverse-mode autodiff computes $\nabla_{\theta} \mathcal{L}$ in one forward-pass cost

Non-scalar output needs a seed vector

If $y = f(x) \in \mathbb{R}^m$, `backward()` must know **which** scalar to differentiate — you supply a vector v and it returns (same vector as the $J^\top \bar{y}$ above, shaped like x)

Non-scalar output needs a seed vector

If $y = f(x) \in \mathbb{R}^m$, `backward()` must know **which** scalar to differentiate — you supply a vector v and it returns (same vector as the $J^\top \bar{y}$ above, shaped like x)

$$J^\top v.$$

Non-scalar output needs a seed vector

If $y = f(x) \in \mathbb{R}^m$, `backward()` must know **which** scalar to differentiate — you supply a vector v and it returns (same vector as the $J^\top \bar{y}$ above, shaped like x)

$J^\top v$.

This is why PyTorch errors on `.backward()` of a non-scalar tensor without a `gradient=` argument.

Compute the very Jacobian we found by hand — then a VJP:

Compute the very Jacobian we found by hand — then a VJP:

```
import torch
x = torch.tensor([1.0, 2.0], requires_grad=True)
def f(x):
    return torch.stack([x[0]**2 + x[1], torch.sin(x[0]*x[1])])

J = torch.autograd.functional.jacobian(f, x)    # full 2x2 Jacobian
v = torch.tensor([1.0, 3.0])
f(x).backward(v)                               # x.grad == J^T v  (a VJP)
```

Compute the very Jacobian we found by hand — then a VJP:

```
import torch
x = torch.tensor([1.0, 2.0], requires_grad=True)
def f(x):
    return torch.stack([x[0]**2 + x[1], torch.sin(x[0]*x[1])])

J = torch.autograd.functional.jacobian(f, x)  # full 2x2 Jacobian
v = torch.tensor([1.0, 3.0])
f(x).backward(v)                             # x.grad == J^T v  (a VJP)
```

jacobian forms the whole matrix; backward(v) returns only $v^\top J$.

```
def g(x):  
    return x[0]**2 + 3*x[1]**2 + x[0]*x[1]  
  
H = torch.autograd.functional.hessian(g, x)    # -> [[2, 1], [1, 6]]
```

```
def g(x):  
    return x[0]**2 + 3*x[1]**2 + x[0]*x[1]  
  
H = torch.autograd.functional.hessian(g, x)    # -> [[2, 1], [1, 6]]
```

$H = \begin{pmatrix} 2 & 1 \\ 1 & 6 \end{pmatrix}$ — matches the second partials by hand.

Follow-along: exact Jacobian / VJP / Hessian Colab — calculate first, then verify.

Summary

Six objects, six shapes, six jobs:

Six objects, six shapes, six jobs:

Object	Shape	Main DL use
$f'(x)$	scalar	activation derivative
∇f	d	parameter update
J	$m \times n$	vector-map sensitivity
$J^\top v$	n	backpropagation
H	$d \times d$	curvature
Hv	d	second-order methods

Final mental model — one idea

derivative = slope

gradient = direction of steepest ascent

Jacobian = local linear map

Hessian = curvature

backprop = an efficient chain of vector-Jacobian products

We can now read every derivative object in deep learning.

Next: **computation graphs and backpropagation** — the chain rule, run backward through the program.