

Next-Token Prediction

Tokens, embeddings, and one complete neural language-model calculation

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

L11 classified spatial sites; L12 classifies the next time site

PUBLIC L11 · SEGMENTATION

one categorical distribution at
every spatial location
output support $H \times W$



PUBLIC L12 · LANGUAGE

one categorical distribution at the
next ordered location
output support V

PUBLIC L11 · SEGMENTATION

one categorical distribution at
every spatial location
output support $H \times W$



PUBLIC L12 · LANGUAGE


one categorical distribution at the
next ordered location
output support V

The classifier and cross-entropy stay. Order, causality, and feeding predictions back are new.

Commit before the model reveals anything

QUESTION

HELD CASE · OPENING COMMITMENT Vocabulary [., a, e, i, n, v]; visible context .n; predict the next token.



Commit before the model reveals anything

QUESTION

HELD CASE · OPENING COMMITMENT Vocabulary [., a, e, i, n, v]; visible context .n; predict the next token.

. n → ?

OUTPUT

What six numbers must the model return?

CHOICE

Will greedy and sampling always pick the same token?

LEARNING

Which embedding rows can this example update?

Commit before the model reveals anything

HELD CASE · OPENING COMMITMENT Vocabulary [., a, e, i, n, v]; visible context .n; predict the next token.

. n → ?

OUTPUT

What six numbers must the model return?

CHOICE

Will greedy and sampling always pick the same token?

LEARNING

Which embedding rows can this example update?

Commit silently. We will derive every answer from one set of weights and revisit all three.

The 80-minute route protects one numerical spine

CORE · 55 MIN

boundary + causal
windows → embedding
lookup → fixed-context
MLP → stable softmax →
CE → $p - y$ → token NLL /
PPL

SHOULD · 15 MIN

autoregressive loop →
greedy / sampling →
temperature → suffix-
blind diagnostic →
opening revisit

OPTIONAL · 10 MIN

full gradient matrices →
learned representation →
Colab and primary
provenance

The final RNN handoff is CORE even when optional material is skipped.

Three clocks reuse the distribution but change the question

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

Three clocks reuse the distribution but change the question

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

Three clocks reuse the distribution but change the question

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

Three clocks reuse the distribution but change the question

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

Always ask: are we updating weights, choosing a token, or scoring held-out truth?

By the end, one sentence should be operational

V VISUAL

A language model maps an observed prefix to a distribution over the next token.

A language model maps an observed prefix to a distribution over the next token.

REPRESENT

turn categorical ids into
learned vectors without
inventing numeric order

PREDICT

produce V logits,
normalize, and train with
token cross-entropy

DIAGNOSE

separate decoding
choices from an
architecture that cannot
see enough history

— observed context / TRAIN — model score or probability / INFER — target or chosen token — correct / retained — loss or error — EVAL

One tiny corpus supplies every visible token

TRAIN NAMES

naveen · navin · neev · neena · vani
naina · avni · anvi · veena · navi

HELD-OUT NAMES

navina · naveena · veenavi · anvina

Toy convention: . is one shared boundary symbol. Repeat it $k = 2$ times as left padding; use one . target to end a name. It is not punctuation.

Tokenization chooses the units the model must predict

unit	example	vocabulary	sequence
character	learning → l e a r n i n g	small	long
word	deep learning → deep learning	very large	short
subword	unpredictable → un predict able	moderate	moderate

Characters keep every row visible. Production usually separates <BOS>, <EOS>, <PAD>, and <UNK>; this toy merges <BOS> and <EOS> as ..

The vocabulary fixes ids and output coordinates

.	a	e	i	n	v
0	1	2	3	4	5

The vocabulary fixes ids and output coordinates

.	a	e	i	n	v
0	1	2	3	4	5

$V = 6$ means six logits, six probabilities, and six possible next tokens.

HELD CASE · VOCABULARY The target a is coordinate 1; context .n reads ids (0, 4).

Sliding a two-character window creates seven examples

..	.n	na	av	ve	ee	en
n	a	v	e	e	n	.

Sliding a two-character window creates seven examples

..	.n	na	av	ve	ee	en
n	a	v	e	e	n	.

top = visible context; bottom = next-token target

HELD CASE · TRAINING PAIRS .n → a is the second example and remains our hand calculation.

Causal training shifts the same sequence by one site

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

input	..	.n	na	av	ve	ee	en
target	n	a	v	e	e	n	.

Causal training shifts the same sequence by one site

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

input	..	.n	na	av	ve	ee	en
target	n	a	v	e	e	n	.

At position t , the target token is not part of the input context.

The chain rule turns sequence probability into next-token factors

Let $B = (., .)$ denote the fixed left boundary padding and let $x_7 = .$ be the end target.

$$p(x_{1:7} \mid B) = \prod_{t=1}^7 p(x_t \mid B, x_{<t}).$$

The chain rule turns sequence probability into next-token factors

D DERIVATION

Let $B = (., .)$ denote the fixed left boundary padding and let $x_7 = .$ be the end target.

$$p(x_{1:7} \mid B) = \prod_{t=1}^7 p(x_t \mid B, x_{<t}).$$

This is an exact identity. The fixed-window approximation comes next.

One ledger will cross probability, loss, and evaluation

HELD CASE · SEQUENCE LEDGER Correct targets for naveen. under contexts . ., .n, na, av, ve, ee, en.

context	. .	. n	na	av	ve	ee	en
target	n	a	v	e	e	n	.
<i>p</i>(target)	.30	derive	.70	.60	.50	.40	.80

One ledger will cross probability, loss, and evaluation

HELD CASE · SEQUENCE LEDGER Correct targets for naveen. under contexts . ., .n, na, av, ve, ee, en.

context	. .	.n	na	av	ve	ee	en
target	n	a	v	e	e	n	.
<i>p(target)</i>	.30	derive	.70	.60	.50	.40	.80

Six probabilities are held illustrative. The orange .n → a entry will come from the exact weights.

Checkpoint: what did the chain rule assume?

QUESTION

A

tokens are independent

B

the sequence has fixed length

C

nothing extra; it is an identity

D

future tokens are visible

Checkpoint: what did the chain rule assume?

QUESTION

A

tokens are independent

B

the sequence has fixed length

C

nothing extra; it is an identity

D

future tokens are visible

C. Each exact factor conditions on the complete preceding history and boundary.

The approximation begins only when the model discards all but the latest k tokens.

A fixed-context MLP deliberately truncates the history

$$p_{\theta}(x_t \mid B, x_{<t}) \approx p_{\theta}(x_t \mid x_{t-2:t-1}).$$

A fixed-context MLP deliberately truncates the history

$$p_{\theta}(x_t \mid B, x_{<t}) \approx p_{\theta}(x_t \mid x_{t-2:t-1}).$$

EXACT CONDITIONAL

boundary plus every earlier token
variable-length input



TOY MODEL

only two latest tokens
fixed-size input for an MLP

This model learns local spelling patterns; it cannot recover a token that has scrolled out.

A token id is a label, not a magnitude

n has id 4 and v has id 5, but

$5 - 4 = 1$ **does not mean** $v - n = 1$.

A token id is a label, not a magnitude

n has id 4 and v has id 5, but

$5 - 4 = 1$ **does not mean** $v - n = 1$.

Feeding raw ids as scalar features invents ordering and distance. Ids are categorical addresses.

The representation must preserve identity without pretending that adjacent ids are similar.

One-hot vectors give each token an independent axis

With vocabulary order [., a, e, i, n, v]:

$$o_{.} = (1 \ 0 \ 0 \ 0 \ 0 \ 0)^T, \quad o_n = (0 \ 0 \ 0 \ 0 \ 1 \ 0)^T.$$

One-hot vectors give each token an independent axis

With vocabulary order $[\cdot, a, e, i, n, v]$:

$$o_{\cdot} = (1 \ 0 \ 0 \ 0 \ 0 \ 0)^T, \quad o_n = (0 \ 0 \ 0 \ 0 \ 1 \ 0)^T.$$

For different tokens $i \neq j$, $\|o_i - o_j\|^2 = 2$.

One-hot is faithful but sparse, V -dimensional, and has no learned similarity structure.

An embedding is a learned row selected by an id

token	feature 1	feature 2
.	1	0
a	0	1
e	1	1
i	-1	1
n	0	2
v	2	0

An embedding is a learned row selected by an id

token	feature 1	feature 2
.	1	0
a	0	1
e	1	1
i	-1	1
n	0	2
v	2	0

$$E \in \mathbb{R}^{6 \times 2}, \quad e_i = E[i] = E^T o_i.$$

The id addresses a row; gradient descent decides what the two coordinates become.

Context .n reads exactly two rows

HELD CASE · LOOKUP ids $(0, 4)$ select rows for . and n.

$$e_{.} = E[0] = (1, 0), \quad e_n = E[4] = (0, 2).$$

Context .n reads exactly two rows

HELD CASE · LOOKUP ids (0, 4) select rows for . and n.

$$e_{.} = E[0] = (1, 0), \quad e_n = E[4] = (0, 2).$$



The target a affects the output loss, but its embedding row is not read on this example.

Commit: concatenate or average the two rows?

QUESTION

A • AVERAGE

$$c = \frac{e. + e_n}{2} = (.5, 1)$$

B • CONCATENATE

$$c = [e.; e_n] = (1, 0, 0, 2)$$

Commit: concatenate or average the two rows?

QUESTION

A · AVERAGE

$$c = \frac{e_{\cdot} + e_n}{2} = (.5, 1)$$

B · CONCATENATE

$$c = [e_{\cdot}; e_n] = (1, 0, 0, 2)$$

B. Concatenation gives the first and second context positions different coordinate blocks.

Commit: concatenate or average the two rows?

A • AVERAGE

$$c = \frac{e_{\cdot} + e_n}{2} = (.5, 1)$$


B • CONCATENATE

$$c = [e_{\cdot}; e_n] = (1, 0, 0, 2)$$


B. Concatenation gives the first and second context positions different coordinate blocks.

$$c_{\cdot n} = (1, 0, 0, 2) \neq (0, 2, 1, 0) = c_{n \cdot}$$

One position is an ordinary six-class classifier



One position is an ordinary six-class classifier



two ids $\rightarrow$ two 2-D rows $\rightarrow$ ordered 4-D context $\rightarrow$ two hidden units $\rightarrow$ six logits

Language modeling changes how examples are constructed; the local head is familiar classification.

The tensor ledger follows the pipeline

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$c = [E[x_{t-2}]; E[x_{t-1}]] \in \mathbb{R}^4,$$

The tensor ledger follows the pipeline

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$c = [E[x_{t-2}]; E[x_{t-1}]] \in \mathbb{R}^4,$$

$$u = W_h c + b_h, \quad h = \text{ReLU}(u) \in \mathbb{R}^2,$$

The tensor ledger follows the pipeline

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$c = [E[x_{t-2}]; E[x_{t-1}]] \in \mathbb{R}^4,$$

$$u = W_h c + b_h, \quad h = \text{ReLU}(u) \in \mathbb{R}^2,$$

$$z = W_o h + b_o \in \mathbb{R}^6, \quad p = \text{softmax}(z).$$

$W_h \in \mathbb{R}^{2 \times 4}$; $W_o \in \mathbb{R}^{6 \times 2}$; rows of W_o follow vocabulary order.

Count parameters once, then keep the shapes visible

block	shape	count
embedding E	$V \times d$	$Vd = 12$
hidden W_h, b_h	$H \times kd, H$	$Hkd + H =$ 10
output W_o, b_o	$V \times H, V$	$VH + V =$ 18

Count parameters once, then keep the shapes visible

block	shape	count
embedding E	$V \times d$	$Vd = 12$
hidden W_h, b_h	$H \times kd, H$	$Hkd + H = 10$
output W_o, b_o	$V \times H, V$	$VH + V = 18$

$12 + 10 + 18 = 40$ parameters.

Increasing context length k grows the hidden block Hkd even before we discuss what the model can remember.

Every weight for $.n \rightarrow a$ is now explicit

HELD CASE · EXACT MODEL input $x = (., n)$; target $y = a$; $c = (1, 0, 0, 2)^T$.

$$W_h = \begin{pmatrix} 0.5 & 0 & 0 & 0.25 \\ 0 & 1 & 1 & 0 \end{pmatrix}, \quad b_h = \begin{pmatrix} 0 \\ -1 \end{pmatrix},$$

Every weight for $.n \rightarrow a$ is now explicit


HELD CASE · EXACT MODEL input $x = (., n)$; target $y = a$; $c = (1, 0, 0, 2)^T$.

$$W_h = \begin{pmatrix} 0.5 & 0 & 0 & 0.25 \\ 0 & 1 & 1 & 0 \end{pmatrix}, \quad b_h = \begin{pmatrix} 0 \\ -1 \end{pmatrix},$$


$$W_o = \begin{pmatrix} -1 & 0 \\ 2 & 1 \\ 0 & -1 \\ -1 & 1 \\ 0 & 0 \\ 1 & -1 \end{pmatrix}, \quad b_o = 0.$$

Rows of W_o are $[., a, e, i, n, v]$. Keep that order for every vector below.

Forward 1: lookup and concatenate




Forward 1: lookup and concatenate



$$E[0] = (1, 0), \quad E[4] = (0, 2),$$

Forward 1: lookup and concatenate

D DERIVATION



$$E[0] = (1, 0), \quad E[4] = (0, 2),$$

$$c = [E[0]; E[4]] = (1, 0, 0, 2)^T.$$

HELD CASE · FORWARD · CONTEXT The ordered context occupies four coordinates; no averaging occurs.

Forward 2: one hidden unit turns on

$$u_1 = .5(1) + 0(0) + 0(0) + .25(2) + 0 = 1,$$

Forward 2: one hidden unit turns on

$$u_1 = .5(1) + 0(0) + 0(0) + .25(2) + 0 = 1,$$

$$u_2 = 0(1) + 1(0) + 1(0) + 0(2) - 1 = -1.$$

Forward 2: one hidden unit turns on

$$u_1 = .5(1) + 0(0) + 0(0) + .25(2) + 0 = 1,$$

$$u_2 = 0(1) + 1(0) + 1(0) + 0(2) - 1 = -1.$$

$$h = (\max(0, 1), \max(0, -1)) = (1, 0).$$

ACTIVE

$h_1 = 1$; gradient may pass

OFF

$h_2 = 0$; ReLU blocks its gradient here

Commit: compute two logits before seeing all six

With $h = (1, 0)$, calculate the rows for target a and alternative v :

TARGET a

$$z_a = (2, 1) \cdot (1, 0) = ?$$

ALTERNATIVE v

$$z_v = (1, -1) \cdot (1, 0) = ?$$

Commit: compute two logits before seeing all six

With $h = (1, 0)$, calculate the rows for target a and alternative v :

TARGET a

$$z_a = (2, 1) \cdot (1, 0) = ?$$

ALTERNATIVE v

$$z_v = (1, -1) \cdot (1, 0) = ?$$

$$z_a = 2 \quad z_v = 1$$

Commit: compute two logits before seeing all six

With $h = (1, 0)$, calculate the rows for target a and alternative v :

TARGET a

$$z_a = (2, 1) \cdot (1, 0) = ?$$

ALTERNATIVE v

$$z_v = (1, -1) \cdot (1, 0) = ?$$

$$z_a = 2 \quad z_v = 1$$

A logit is an unconstrained score, not yet a probability.

Forward 3: every output row supplies one logit

D DERIVATION

HELD CASE · FORWARD · LOGITS vocabulary order [., a, e, i, n, v].

$$z = W_o h = (-1, 2, 0, -1, 0, 1)^T.$$

Forward 3: every output row supplies one logit

HELD CASE · FORWARD · LOGITS vocabulary order [., a, e, i, n, v].

$$z = W_o h = (-1, 2, 0, -1, 0, 1)^T.$$

.	a	e	i	n	v
-1	2	0	-1	0	1

Softmax will preserve this ranking while converting scores into a normalized distribution.

Stable softmax subtracts a constant before exponentiating

D DERIVATION

Naively exponentiating large logits can overflow. Softmax is unchanged by a shared shift:

Stable softmax subtracts a constant before exponentiating

Naively exponentiating large logits can overflow. Softmax is unchanged by a shared shift:

$$\text{softmax}(z) = \text{softmax}(z - m), \quad m = \max_j z_j = 2.$$

Stable softmax subtracts a constant before exponentiating

Naively exponentiating large logits can overflow. Softmax is unchanged by a shared shift:

$$\text{softmax}(z) = \text{softmax}(z - m), \quad m = \max_j z_j = 2.$$

$$z - m = (-3, 0, -2, -3, -2, -1).$$

Subtracting the maximum is an implementation safeguard, not a model change.

Forward 4: exponentiate and normalize

$$\exp(z - m) = (.0498, 1, .1353, .0498, .1353, .3679),$$

Forward 4: exponentiate and normalize

$$\exp(z - m) = (.0498, 1, .1353, .0498, .1353, .3679),$$

$$Z = \sum_j \exp(z_j - m) = 1.7381,$$

Forward 4: exponentiate and normalize

$$\exp(z - m) = (.0498, 1, .1353, .0498, .1353, .3679),$$

$$Z = \sum_j \exp(z_j - m) = 1.7381,$$

$$p = \frac{\exp(z - m)}{Z} = (.028644, .575333, .077863, .028644, .077863, .211653).$$


HELD CASE · FORWARD · PROBABILITY The exact orange ledger entry is now $p(a|.n) = .575333$.

The six bars share a true 0-to-1 probability axis

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL




The six bars share a true 0-to-1 probability axis

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL



Direct labels follow [., a, e, i, n, v]; all six probabilities sum to 1.

a is most likely, but alternatives retain nonzero probability.

Forward 5: cross-entropy reads the target coordinate

D DERIVATION

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

The target is a, so $y = (0, 1, 0, 0, 0, 0)$.

Forward 5: cross-entropy reads the target coordinate

D DERIVATION

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

The target is a, so $y = (0, 1, 0, 0, 0, 0)$.

$$\mathcal{L} = - \sum_j y_j \log p_j = -\log(.575333) = .552806 \text{ nats.}$$

IF p_a RISES

the loss falls

IF $p_a \rightarrow 1$

$\mathcal{L} \rightarrow 0$

The sequence ledger now contains no second version

context	..	.n	na	av	ve	ee	en
target	n	a	v	e	e	n	.
<i>p(target)</i>	.30	.575333	.70	.60	.50	.40	.80

The sequence ledger now contains no second version

context	..	.n	na	av	ve	ee	en
target	n	a	v	e	e	n	.
$p(\text{target})$	.30	.575333	.70	.60	.50	.40	.80

$$p_{\theta(\text{naveen}|\text{..})} \approx .30(.575333)(.70)(.60)(.50)(.40)(.80) = .0115987.$$

The sequence probability is small because all seven next-token decisions must agree.

Logs turn the same product into NLL and perplexity

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$-\log(.0115987) = 4.45686 \text{ nats,}$$

Logs turn the same product into NLL and perplexity

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$-\log(.0115987) = 4.45686 \text{ nats},$$

$$\text{mean NLL} = \frac{4.45686}{7} = .636694 \text{ nats/token},$$

Logs turn the same product into NLL and perplexity

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$-\log(.0115987) = 4.45686 \text{ nats},$$

$$\text{mean NLL} = \frac{4.45686}{7} = .636694 \text{ nats/token},$$

$$\text{PPL} = \exp(.636694) = 1.89022.$$

| This is one illustrative name ledger, not the trained notebook's validation score.

Commit: what sign must the target-logit gradient have?

QUESTION

A • POSITIVE

gradient descent will lower z_a

B • NEGATIVE

gradient descent will raise z_a

Commit: what sign must the target-logit gradient have?

QUESTION

A • POSITIVE

gradient descent will lower z_a

B • NEGATIVE

gradient descent will raise z_a

B. To reduce $-\log p_a$, the update must increase the target logit relative to alternatives.

Commit: what sign must the target-logit gradient have?

QUESTION

A • POSITIVE

gradient descent will lower z_a

B • NEGATIVE

gradient descent will raise z_a

B. To reduce $-\log p_a$, the update must increase the target logit relative to alternatives.

Softmax plus cross-entropy makes the exact sign visible in one vector.

Backward 1: softmax plus CE gives $p - y$

D DERIVATION

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$\nabla_z \mathcal{L} = p - y$$

Backward 1: softmax plus CE gives $p - y$

D DERIVATION

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$\nabla_z \mathcal{L} = p - y$$

$$= (.028644, -.424667, .077863, .028644, .077863, .211653).$$

- the target coordinate is negative, so gradient descent raises z_a ;
- every alternative coordinate is positive, so gradient descent lowers it.

Backward 2: output gradients are outer products

Because $z = W_o h + b_o$,

$$\nabla_{W_o} \mathcal{L} = (p - y)h^T, \quad \nabla_{b_o} \mathcal{L} = p - y.$$

Backward 2: output gradients are outer products

Because $z = W_o h + b_o$,

$$\nabla_{W_o} \mathcal{L} = (p - y)h^T, \quad \nabla_{b_o} \mathcal{L} = p - y.$$

With $h = (1, 0)$, the first column equals $p - y$ and the second column is zero.

An inactive feature emits no output-weight gradient on this example.

Backward 3: ReLU chooses which hidden signal survives

D DERIVATION

$$\nabla_h \mathcal{L} = W_o^T(p-y) = (-.694969, -.685539)^T.$$

Backward 3: ReLU chooses which hidden signal survives

D DERIVATION

$$\nabla_h \mathcal{L} = W_o^T(p-y) = (-.694969, -.685539)^T.$$

The forward pre-activation was $u = (1, -1)$, so the ReLU mask is $(1, 0)$.

Backward 3: ReLU chooses which hidden signal survives

$$\nabla_h \mathcal{L} = W_o^T(p-y) = (-.694969, -.685539)^T.$$

The forward pre-activation was $u = (1, -1)$, so the ReLU mask is $(1, 0)$.

$$\nabla_u \mathcal{L} = \nabla_h \mathcal{L} \odot (1, 0) = (-.694969, 0)^T.$$

The second unit had a nonzero upstream signal, but its negative pre-activation blocks that signal.

Backward 4: return to the ordered context

Since $u = W_h c + b_h$,

Since $u = W_h c + b_h$,

$$\nabla_c \mathcal{L} = W_h^T \nabla_u \mathcal{L} = (-.347485, 0, 0, -.173742)^T.$$

Backward 4: return to the ordered context

Since $u = W_h c + b_h$,

$$\nabla_c \mathcal{L} = W_h^T \nabla_u \mathcal{L} = (-.347485, 0, 0, -.173742)^T.$$

block	$e.[1]$	$e.[2]$	$e_{n[1]}$	$e_{n[2]}$
∇_c	$-.347485$	0	0	$-.173742$

Concatenation makes the split back into two embedding rows unambiguous.

$$\nabla_{e_{\cdot}} \mathcal{L} = (-.347485, 0), \quad \nabla_{e_n} \mathcal{L} = (0, -.173742).$$

$$\nabla_{e_1} \mathcal{L} = (-.347485, 0), \quad \nabla_{e_n} \mathcal{L} = (0, -.173742).$$

Scatter these into rows 0 and 4 of E ; every other row receives zero from this lookup.

Backward 5: split, scatter, and update only rows read

$$\nabla_{e_{\cdot}} \mathcal{L} = (-.347485, 0), \quad \nabla_{e_n} \mathcal{L} = (0, -.173742).$$

Scatter these into rows 0 and 4 of E ; every other row receives zero from this lookup. With $\eta = .1$:

$$e_{\cdot} \leftarrow (1.03475, 0), \quad e_n \leftarrow (0, 2.01737).$$

HELD CASE · OPENING · LEARNING The rows for $\cdot$ and n move. Target row a was not read and does not move through the lookup.

One scalar loss teaches every used block



One scalar loss teaches every used block



- output rows learn which hidden features support each token;
- hidden weights learn combinations of ordered context coordinates;
- read embedding rows learn features that improve future predictions.

The path is ordinary backpropagation; language structure entered when we built context-target pairs.

TRAIN repeats the calculation over a declared batch

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$\mathcal{L}_B = -\frac{1}{|B|} \sum_{i \in B} \log p_{\theta}(y_i | c_i).$$


TRAIN repeats the calculation over a declared batch

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$\mathcal{L}_B = -\frac{1}{|B|} \sum_{i \in B} \log p_{\theta}(y_i | c_i).$$



The notebook uses the entire training set as B on each Adam step: full-batch training, not a minibatch claim.

Split whole names before creating windows

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN

10 whole names $\rightarrow$ 56 context-target pairs
weights receive updates

VALIDATION

4 different whole names $\rightarrow$ 30 pairs
no updates

Split whole names before creating windows

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN

10 whole names $\rightarrow$ 56 context-target pairs
weights receive updates

VALIDATION

4 different whole names $\rightarrow$ 30 pairs
no updates

Do not randomly split windows from the same name. Nearby windows share almost all their evidence.

**The split unit must match the object whose generalization we claim:
here, a whole name.**

EVAL declares tokenizer, targets, and reduction

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$\text{NLL} = -\frac{1}{N} \sum_{i=1}^N \log p_{\theta}(y_i | c_i), \quad \text{PPL} = \exp(\text{NLL}).$$

EVAL declares tokenizer, targets, and reduction

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$\text{NLL} = -\frac{1}{N} \sum_{i=1}^N \log p_{\theta}(y_i | c_i), \quad \text{PPL} = \exp(\text{NLL}).$$

TOKENIZER

same character vocabulary and
boundary policy

TARGETS

N non-padding next-token targets

DATA

same held-out corpus and token-
weighted reduction


Perplexities are comparable only when all three contracts agree.

The executed notebook shows the full learning trajectory

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL



EXECUTED · SEED 11

300 full-batch Adam updates
constructed 10+4-name course corpus

COMPUTED · UPDATE 300


train · NLL 0.645308 · PPL 1.90657
held-out · NLL 0.677965 · PPL 1.96986

The executed notebook shows the full learning trajectory

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL



EXECUTED · SEED 11

300 full-batch Adam updates
constructed 10+4-name course corpus

COMPUTED · UPDATE 300

train · NLL 0.645308 · PPL 1.90657
held-out · NLL 0.677965 · PPL 1.96986

Every point evaluates the same post-update parameter snapshot; no smoothing. The train set has 56 targets and the four held-out names have 30 targets.

Optimization + in-protocol transfer on a constructed course corpus; not broad name quality.

Uniform prediction is the first evaluation sanity check

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

For $V = 6$, a uniform model has $p(y) = \frac{1}{6}$:

Uniform prediction is the first evaluation sanity check

TRAIN true context $\rightarrow \underline{p}$; CE $\rightarrow$ update

INFER current context $\rightarrow \underline{p}$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

For $V = 6$, a uniform model has $p(y) = \frac{1}{6}$:

$$\text{NLL} = -\log\left(\frac{1}{6}\right) = \log 6 = 1.79176, \quad \text{PPL} = 6.$$

Uniform prediction is the first evaluation sanity check

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

For $V = 6$, a uniform model has $p(y) = \frac{1}{6}$:

$$\text{NLL} = -\log\left(\frac{1}{6}\right) = \log 6 = 1.79176, \quad \text{PPL} = 6.$$

UNIFORM

1.79176 NLL
6.00 PPL

VALIDATION

.677965 NLL
1.96986 PPL

Uniform prediction is the first evaluation sanity check

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

For $V = 6$, a uniform model has $p(y) = \frac{1}{6}$:

$$\text{NLL} = -\log\left(\frac{1}{6}\right) = \log 6 = 1.79176, \quad \text{PPL} = 6.$$

UNIFORM

1.79176 NLL
6.00 PPL

VALIDATION

.677965 NLL
1.96986 PPL

The trained model predicts held-out next characters far better than uniform guessing under this protocol.

The same p now crosses from TRAIN to INFER

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

The same p now crosses from TRAIN to INFER

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

The same p now crosses from TRAIN to INFER

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN

true context $\rightarrow$ distribution $\rightarrow$ compare with true target
 $\rightarrow$ update

INFER

current context $\rightarrow$ distribution $\rightarrow$ choose token $\rightarrow$
append

The same p now crosses from TRAIN to INFER

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

TRAIN


true context $\rightarrow$ distribution $\rightarrow$ compare with true target
 $\rightarrow$ update

INFER


current context $\rightarrow$ distribution $\rightarrow$ choose token $\rightarrow$
append

Weights stay fixed during generation; the chosen token changes the next input context.

Autoregressive generation closes a feedback loop




Autoregressive generation closes a feedback loop



Start from . . ; stop when the model emits . or reaches a declared maximum length.

Autoregressive generation closes a feedback loop




Start from . . ; stop when the model emits . or reaches a declared maximum length.

Autoregressive means the model conditions on tokens already present, including its own earlier choices.


Commit: must greedy and sampling choose the same token?

QUESTION



Commit: must greedy and sampling choose the same token?

QUESTION



GREEDY


$$\hat{x} = \arg \max_v p(v|.n)$$

SAMPLING

$$\hat{x} \sim \text{Categorical}(p)$$

Commit: must greedy and sampling choose the same token?

QUESTION



GREEDY

$$\hat{x} = \arg \max_v p(v|.n)$$

SAMPLING

$$\hat{x} \sim \text{Categorical}(p)$$

No. Greedy always chooses a; sampling can choose any token with nonzero mass.

One uniform draw makes categorical sampling concrete

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

token	.	a	e	i	n	v
<i>p</i>	.029	.575	.078	.029	.078	.212
CDF	.029	.604	.682	.710	.788	1.000

One uniform draw makes categorical sampling concrete

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

token	.	a	e	i	n	v
p	.029	.575	.078	.029	.078	.212
CDF	.029	.604	.682	.710	.788	1.000

Draw $u = .650$ from $\text{Uniform}(0, 1)$.

One uniform draw makes categorical sampling concrete

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

token	.	a	e	i	n	v
<i>p</i>	.029	.575	.078	.029	.078	.212
CDF	.029	.604	.682	.710	.788	1.000

Draw $u = .650$ from $\text{Uniform}(0, 1)$.

$.604 < .650 \leq .682 \Rightarrow$ **sample e.**

One uniform draw makes categorical sampling concrete

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

token	.	a	e	i	n	v
p	.029	.575	.078	.029	.078	.212
CDF	.029	.604	.682	.710	.788	1.000

Draw $u = .650$ from $\text{Uniform}(0, 1)$.

$.604 < .650 \leq .682 \Rightarrow$ **sample e.**

The model supplies probabilities; the decoder supplies the decision rule and randomness.

Temperature rescales logits before the same softmax

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$p_{i(\tau)} = \text{softmax}\left(\frac{z}{\tau}\right)_i, \quad \tau > 0.$$

Temperature rescales logits before the same softmax

D DERIVATION

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$p_{i(\tau)} = \text{softmax}\left(\frac{z}{\tau}\right)_i, \quad \tau > 0.$$

$$\tau = .5$$

logit gaps double
distribution concentrates

$$\tau = 1$$

learned distribution
no rescaling

$$\tau = 2$$

logit gaps halve
distribution flattens

Temperature rescales logits before the same softmax

D DERIVATION

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

$$p_{i(\tau)} = \text{softmax}\left(\frac{z}{\tau}\right)_i, \quad \tau > 0.$$

$$\tau = .5$$

logit gaps double
distribution concentrates

$$\tau = 1$$




learned distribution
no rescaling

$$\tau = 2$$




logit gaps halve
distribution flattens

Temperature changes neither weights nor token ranking for positive τ .

The same six logits produce three exact distributions






The same six logits produce three exact distributions



τ	.	a	e	i	n	v
.5	.002	.850	.016	.002	.016	.115
1	.029	.575	.078	.029	.078	.212
2	.080	.359	.132	.080	.132	.218

The same six logits produce three exact distributions



τ	.	a	e	i	n	v
.5	.002	.850	.016	.002	.016	.115
1	.029	.575	.078	.029	.078	.212
2	.080	.359	.132	.080	.132	.218

Concentration is measurable; “safer” or “more correct” is not guaranteed.

Choose temperature by behavior, not a moral label

QUESTION

SYMPTOM

samples collapse to the same spelling

TEST

hold logits and seed fixed; compare token frequencies across

$\mathcal{T}$

Choose temperature by behavior, not a moral label

QUESTION

SYMPTOM

samples collapse to the same spelling

TEST

hold logits and seed fixed; compare token frequencies across

τ

LOWER τ

more mass on current top tokens; less variation

HIGHER τ

more mass on alternatives; more variation

Choose temperature by behavior, not a moral label

QUESTION

SYMPTOM

samples collapse to the same spelling

TEST

hold logits and seed fixed; compare token frequencies across

τ

LOWER τ

more mass on current top tokens; less variation

HIGHER τ

more mass on alternatives; more variation

If the model ranks a bad token highest, lowering temperature can make that error more consistent.

The trained MLP is blind to prefixes with the same suffix

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

PREFIX A

na $\rightarrow$ visible context na

$\equiv$

PREFIX B

veena $\rightarrow$ visible context na

The trained MLP is blind to prefixes with the same suffix

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

PREFIX A

na $\rightarrow$ visible context na

$\equiv$

PREFIX B

veena $\rightarrow$ visible context na

$$p = (.418671, 9.91 \times 10^{-7}, 3.77 \times 10^{-6}, .142728, 3.17 \times 10^{-5}, .438565)$$

The trained MLP is blind to prefixes with the same suffix

TRAIN true context $\rightarrow p$; CE $\rightarrow$ update

INFER current context $\rightarrow p$; choose $\rightarrow$
append

EVAL held-out truth $\rightarrow$ mean NLL / PPL

PREFIX A

na $\rightarrow$ visible context na

$\equiv$

PREFIX B

veena $\rightarrow$ visible context na

$$p = (.418671, 9.91 \times 10^{-7}, 3.77 \times 10^{-6}, .142728, 3.17 \times 10^{-5}, .438565)$$

The two distributions are exactly identical because the model receives exactly the same two ids.

Diagnose the architecture before tuning the optimizer

SYMPTOM

different long prefixes produce identical
next-token distributions

SUSPECT

the useful distinction lies outside the
fixed window

TEST

hold the suffix fixed, change only earlier
tokens, compare p exactly

SYMPTOM

different long prefixes produce identical
next-token distributions

SUSPECT

the useful distinction lies outside the
fixed window

TEST

hold the suffix fixed, change only earlier
tokens, compare p exactly

If p cannot change, more training cannot make this architecture use the hidden prefix.

SYMPTOM

different long prefixes produce identical
next-token distributions

SUSPECT

the useful distinction lies outside the
fixed window

TEST

hold the suffix fixed, change only earlier
tokens, compare p exactly

If p cannot change, more training cannot make this architecture use the hidden prefix.

This is an information-access failure, not evidence that Adam, embeddings, or temperature are broken.

A larger window delays the failure but does not remove it

D DERIVATION

$W_h \in \mathbb{R}^{H \times kd} \Rightarrow Hkd \text{ weights.}$

A larger window delays the failure but does not remove it

$W_h \in \mathbb{R}^{H \times kd} \Rightarrow Hkd \text{ weights.}$

- doubling k doubles this block;
- the same token in slot 1 and slot 2 meets different columns;
- the maximum usable history remains fixed at model construction.

A larger window delays the failure but does not remove it

$W_h \in \mathbb{R}^{H \times kd} \Rightarrow Hkd \text{ weights.}$

- doubling k doubles this block;
- the same token in slot 1 and slot 2 meets different columns;
- the maximum usable history remains fixed at model construction.

Concatenation preserves order, but the MLP does not reuse one position-independent update across time.

Revisit: every opening commitment now has evidence

QUESTION

HELD CASE · OPENING REVISIT context .n, exact six-token model.

OUTPUT

$$p = (.029, .575, .078, .029, .078, .212)$$

CHOICE

greedy $\rightarrow$ a; sample with $u = .65$
 $\rightarrow$ e

LEARNING

lookup updates rows . and n

Revisit: every opening commitment now has evidence

QUESTION

HELD CASE · OPENING REVISIT context .n, exact six-token model.

OUTPUT

$$p = (.029, .575, .078, .029, .078, .212)$$

CHOICE

greedy $\rightarrow$ a; sample with $u = .65$
 $\rightarrow$ e

LEARNING

lookup updates rows . and n

One distribution supports training, decoding, and evaluation only after the clock and policy are declared.

OPTIONAL: expose the full output-gradient matrix

With $h = (1, 0)$,

$$\nabla_{w_o} \mathcal{L} = \begin{pmatrix} .028644 & 0 \\ -.424667 & 0 \\ .077863 & 0 \\ .028644 & 0 \\ .077863 & 0 \\ .211653 & 0 \end{pmatrix}.$$

OPTIONAL: expose the full output-gradient matrix

With $h = (1, 0)$,

$$\nabla_{w_o} \mathcal{L} = \begin{pmatrix} .028644 & 0 \\ -.424667 & 0 \\ .077863 & 0 \\ .028644 & 0 \\ .077863 & 0 \\ .211653 & 0 \end{pmatrix}.$$

Each vocabulary row receives its own logit error; the inactive hidden feature leaves the second column zero.

OPTIONAL: use the notebook as a prediction protocol

I INTERACTIVE

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L11/01_char_mlp_language_model.ipynb

Before running: predict the seven windows, active ReLU unit, sign of the target gradient, two embedding rows that move, and whether `na` and `veena` can differ.

OPTIONAL: use the notebook as a prediction protocol

I INTERACTIVE

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L11/01_char_mlp_language_model.ipynb

Before running: predict the seven windows, active ReLU unit, sign of the target gradient, two embedding rows that move, and whether `na` and `veena` can differ.

VERIFY

exact forward/backward arrays; train/validation NLL;
generated names

EXPLAIN

why temperature changes diversity and why a shared suffix
forces identical p

OPTIONAL: use the notebook as a prediction protocol

I INTERACTIVE

INTERACTIVE

↗ colab.research.google.com/github/nipunbatra/dl-teaching/blob/master/notebooks/L11/01_char_mlp_language_model.ipynb

Before running: predict the seven windows, active ReLU unit, sign of the target gradient, two embedding rows that move, and whether `na` and `veena` can differ.

VERIFY

exact forward/backward arrays; train/validation NLL;
generated names

EXPLAIN

why temperature changes diversity and why a shared suffix
forces identical p

Prediction before execution turns the notebook from a demo into evidence.

PRIMARY SOURCE

Bengio et al. (2003)
distributed representations plus a feedforward neural
probabilistic language model

FOLLOW-ALONG

Exact ES 667 Colab
deterministic NumPy forward/backward, full-batch training,
generation, and suffix test

PRIMARY SOURCE

Bengio et al. (2003)
distributed representations plus a feedforward neural
probabilistic language model

FOLLOW-ALONG

Exact ES 667 Colab
deterministic NumPy forward/backward, full-batch training,
generation, and suffix test

The paper scales the central idea; our six-token model keeps its mechanics auditable by hand.

Next: replace a fixed window with a shared recurrent state

THIS LECTURE

$c_t = [e_{t-2}; e_{t-1}]$
fixed access; position-specific columns



PUBLIC L13

$h_t = f_{\theta}(h_{t-1}, e_t)$
shared update; variable-length stream

Next: replace a fixed window with a shared recurrent state

THIS LECTURE

$$c_t = [e_{t-2}; e_{t-1}]$$

fixed access; position-specific columns



PUBLIC L13

$$h_t = f_{\theta}(h_{t-1}, e_t)$$

shared update; variable-length stream

The next-token objective stays; the mechanism for carrying context changes.

A language model estimates a distribution over the **next token**.

Next: **RNNs, LSTMs, and GRUs** - learn what recurrent state remembers, forgets, and struggles to train.