

Sequence Models I: RNNs, BPTT, LSTMs and GRUs

Remembering with a hidden state — and gating what to keep

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

From fixed windows to recurrent state

Last lecture built a language model over a **fixed window**: predict the next token from the previous k .

$$p_{\theta}(x_{t+1} \mid x_{<t}) \approx p_{\theta}(x_{t+1} \mid x_{t-k+1:t})$$

This lecture keeps the same **task** — next-token prediction — but changes **how the model remembers the past**: from a fixed window of inputs to a **running hidden state**.

A language model factorizes the joint probability by the chain rule:

$$p(x_{1:T}) = \prod_{t=1}^T p(x_t \mid x_{<t})$$

and we train it by minimizing the negative log-likelihood — token-level cross-entropy:

$$\mathcal{L} = - \sum_{t=1}^T \log p_{\theta}(x_t \mid x_{<t})$$

each factor is one V -way classification: “which token comes next?”

The fixed-window MLP, and its limits

The MLP-LM concatenates the last k embeddings and runs one MLP:

fixed context — anything before x_{t-k} is invisible; long-range dependencies are lost.

params grow with k — $W_h \in \mathbb{R}^{H \times kd}$; doubling context doubles the layer.

No shared state across time. The same token at position 1 vs 2 hits **different** columns of W_h — computation is not reused, and the input length is frozen at k .

What we actually want

A better sequence model should:

- read **one token at a time**, for a sequence of **any length**;
- keep a **memory** h_t — a summary of everything seen so far, $x_{1:t}$;
- **predict** from that memory, $\hat{y}_t = g(h_t)$;
- **share** the same parameters at every step.

a running state updated by **one** function: $h_t = f_\theta(h_{t-1}, x_t)$

The recurrent idea

Replace the fixed window with a **state** that is updated at every step:

$$\underbrace{p(x_{t+1} \mid x_{t-k:t})}_{\text{MLP: fixed window}} \longrightarrow \underbrace{h_t = f_\theta(h_{t-1}, x_t)}_{\text{RNN: recurrent state}}$$


One **shared** function f_θ applied over and over gives us, for free:

Property	Why it follows
variable length	same f_θ runs for as many steps as there are tokens
parameter sharing	one f_θ , not one weight block per position
learned state	h_t is a trainable summary of $x_{1:t}$
online	update as each token arrives — no need to see the whole sequence


Fixed window vs recurrent state

v

window $k = 3$ only



fixed window (MLP-LM)



state carries all of the past

recurrent (RNN)

the MLP forgets past the window; the RNN **carries** the past forward in h_t

The vanilla (Elman) RNN

At each step t : look up the token, update the state, read out a prediction.


$$e_t = E[x_t] \quad (\text{embedding lookup})$$

$$a_t = W_{xh} e_t + W_{hh} h_{t-1} + b_h, \quad h_t = \varphi(a_t)$$

$$z_t = W_{hy} h_t + b_y, \quad p_t = \text{softmax}(z_t)$$

$\varphi = \tanh$ usually; one set of weights W_{xh}, W_{hh}, W_{hy} reused at **every** step (Elman)

One shared cell: the past state h_{t-1} and the new input x_t meet, then φ gives h_t .




$$h_t = \varphi(W_{xh}e_t + W_{hh}h_{t-1} + b_h) - \text{the dashed edge is the recurrence}$$

Unrolling the RNN in time

v

The **same** cell, drawn once per step — this is what backprop actually sees:



W_{xh}, W_{hh}, W_{hy} are **shared** across all steps — the diagram repeats one cell

Unroll the recurrence and the whole past is nested inside h_t :

$$h_t = f(h_{t-1}, x_t) = f(f(h_{t-2}, x_{t-1}), x_t) = \dots$$

$$h_t = f(f(\dots f(h_0, x_1) \dots, x_{t-1}), x_t)$$

h_t is a **learned, fixed-size summary** of the entire prefix $x_{1:t}$ — no matter how long. The window is gone; the state remembers.

Feed a sentence in; the **target at each step is the next token** (teacher forcing):

step t	1	2	3	4
input x_t	deep	learning	is	fun
state	h_1	h_2	h_3	h_4
target x_{t+1}	learning	is	fun	<eos>

loss = $\sum_t \text{cross-entropy}(p_t, x_{t+1})$ – every position supervises the step before it

The same cell, wired differently, covers most sequence tasks:

Pattern	Shape	Example
many-to-one	$x_{1:T} \rightarrow y$	sentiment: read all, predict at h_T
many-to-many (aligned)	$x_{1:T} \rightarrow y_{1:T}$	POS tag each token
language model	$x_{1:T} \rightarrow x_{2:T+1}$	next-token at every step
seq2seq (encoder–decoder)	$x_{1:T} \rightarrow y_{1:U}$	translation — different length

Tensor shapes

Track a batch of B sequences, length T , embedding d , hidden m , vocab V :

Tensor	Shape	Meaning
token ids X	$B \times T$	one id per position
embeddings $E[X]$	$B \times T \times d$	one d -vector per token
hidden states H	$B \times T \times m$	the state at every step
logits Z	$B \times T \times V$	a next-token score at every step

time is now an **axis**, not a fixed slot — the model handles any T

A fully-worked RNN forward pass

Strip it to **one** hidden unit — everything is a scalar so we can do it by hand:

$$h_t = \tanh(w_x x_t + w_h h_{t-1} + b)$$

$$w_x = 1, \quad w_h = 0.5, \quad b = 0, \quad h_0 = 0, \quad x = (1, 2, -1)$$

three steps, three states — watch the memory build

$$a_1 = w_x x_1 + w_h h_0 + b = 1 \cdot 1 + 0.5 \cdot 0 = 1$$

$$h_1 = \tanh(1) \approx 0.762$$

with no past state ($h_0 = 0$), the first step is just the input through $\tanh$

Now the past enters through $w_h h_1$:

$$a_2 = w_x x_2 + w_h h_1 = 1 \cdot 2 + 0.5 \cdot 0.762 = 2.381$$

$$h_2 = \tanh(2.381) \approx 0.983$$

the state is being pushed toward the +1 saturation of $\tanh$

A negative input, but the state still remembers the positive past:

$$a_3 = w_x x_3 + w_h h_2 = 1 \cdot (-1) + 0.5 \cdot 0.983 = -0.509$$


$$h_3 = \tanh(-0.509) \approx -0.469$$

the $+0.49$ carried from h_2 **softened** the -1 input — that is memory at work

The forward trace

v

t	x_t	$a_t = x_t + 0.5h_{t-1}$	$h_t = \tanh(a_t)$
1	+1	1.000	0.762
2	+2	2.381	0.983
3	-1	-0.509	-0.469



Substitute backwards

h_3 depends on **every** earlier input:

$$h_3 = \tanh \left(x_3 + 0.5 \underbrace{\tanh \left(x_2 + 0.5 \underbrace{\tanh(x_1)}_{h_1} \right)}_{h_2} \right)$$

x_1 reaches h_3 through **two** nested $\tanh$ and two factors of $w_h = 0.5$. Information persists — but each step **attenuates** it by w_h . Hold that thought (Part V).

INTERACTIVE

(to build) — set w_x, w_h, b and type a short sequence; watch h_t update step by step and the state trajectory trace out, exactly like the scalar example above.

Backpropagation through time

Unrolled, the RNN is just a **very deep** feedforward net whose layers **share weights**:

$$h_0 \rightarrow h_1 \rightarrow h_2 \rightarrow \dots \rightarrow h_T$$

training it = backprop on the unrolled graph = backpropagation through time (BPTT)

same chain rule as Lecture 3 — only now the “depth” is the sequence length T

Every step contributes a term; the total loss is their sum:

$$\mathcal{L} = \sum_{t=1}^T \mathcal{L}_t, \quad \mathcal{L}_t = -\log p_{\theta}(x_{t+1} \mid x_{\leq t})$$

Because the weights are **shared**, each step's gradient lands on the **same** parameters:

$$\nabla_{W_{hh}} \mathcal{L} = \sum_{t=1}^T \nabla_{W_{hh}} \mathcal{L}_t$$

this is exactly Lecture 3's rule: a shared parameter **accumulates** gradient from every use

Write $\bar{u} = \partial \mathcal{L} / \partial u$. Walking one step back through the cell:

$$\bar{a}_t = \bar{h}_t \odot \varphi'(a_t) \quad (\text{through the tanh})$$

$$\nabla_{W_{xh}} += \bar{a}_t e_t^\top, \quad \nabla_{W_{hh}} += \bar{a}_t h_{t-1}^\top$$


$$\bar{h}_{t-1} += W_{hh}^\top \bar{a}_t \quad (\text{state gradient to the previous step})$$

The state h_t is used **twice** — so its gradient is a **sum** of two paths:

$$\bar{h}_t = \underbrace{\frac{\partial \mathcal{L}_t}{\partial h_t}}_{\text{output path: its own loss}} + \underbrace{\frac{\partial \mathcal{L}_{\geq t+1}}{\partial h_t}}_{\text{future path: } W_{(h_t)}^{\text{top}} \text{ overline{a}}_{(t+1)}}$$

Exactly the **branch rule**: h_t feeds both this step's readout **and** the next step's state, so gradients from both add. Miss the second term and long-range learning is impossible.

Forward state (teal) flows right; the adjoint (orange = own loss, red = future) flows **back**:



each $\bar{h}_t$ collects its own loss ($\downarrow$) plus the future via $W_{hh}^\top$ ($\leftarrow$)

Full BPTT over a length- 10^4 sequence is expensive and unstable. So **chunk** it:

- process the sequence in blocks of length K (say $K = 35$);
- backprop **within** a block; then **carry the state but detach it**:

$$h_{\text{start}}^{(\text{next block})} = \text{detach}\left(h_{\text{end}}^{(\text{this block})}\right)$$

the state keeps flowing forward; the **gradient** stops at the block boundary — bounded cost

INTERACTIVE

(to build) — step through the unrolled graph and watch each $\bar{h}_t$ accumulate its two contributions; toggle truncation length K and see where the gradient is cut.

Vanishing and exploding gradients

The long-range gradient is a product

D

To learn a dependency from step k to step T , the gradient runs the whole chain:

$$\frac{\partial h_T}{\partial h_k} = \prod_{j=k+1}^T \frac{\partial h_j}{\partial h_{j-1}}, \quad \frac{\partial h_j}{\partial h_{j-1}} = D_j W_{hh}$$

where $D_j = \text{diag}(\varphi'(a_j))$ is the activation slope at step j .

a product of $(T - k)$ matrices — it tends to **vanish or **explode****

The scalar recurrence lays it bare

D

Take one unit, $h_t = \tanh(w h_{t-1})$. Near 0, $\tanh'(0) = 1$, so:

$$\frac{\partial h_T}{\partial h_0} \approx w^T$$

$|w| < 1$: $w^T \rightarrow 0$ — **vanishing**. The distant past cannot send a gradient.

$|w| > 1$: $w^T \rightarrow \infty$ — **exploding**. Gradients blow up, training diverges.

Fifty steps back, with two nearby weights:

$$w = 0.9 : \quad 0.9^{50} \approx 0.005 \quad (\text{vanished})$$


$$w = 1.1 : \quad 1.1^{50} \approx 117 \quad (\text{exploded})$$

a 10% change in w flips the gradient from $\frac{1}{200}$ to $\times 117$ — over just 50 steps

The saturation problem too

v


$\tanh'(z) = 1 - \tanh^2(z) \leq 1$, and $\rightarrow 0$ once the unit saturates:



each $D_j \leq 1$ can only **shrink** the signal further — saturation makes vanishing worse

Gradient norm vs time-lag

v



vanish (decay), explode (blow up), or — with a **gate** holding $f_t \approx 1$ — stay **flat**

Explosion has a cheap fix: cap the global gradient **norm** before the update.

$$\text{if } \|g\| > c : \quad g \leftarrow c \frac{g}{\|g\|}$$

Rescale, don't truncate — the **direction** of g is preserved, only its **length** is capped at c (Pascanu, Mikolov & Bengio 2013).

Threshold $c = 5$, gradient $g = (6, 8)$:

$$\|g\| = \sqrt{6^2 + 8^2} = 10 > 5$$

$$g \leftarrow 5 \cdot \frac{6, 8}{10} = (3, 4), \quad \|(3, 4)\| = 5 \quad \checkmark$$

same direction $(3, 4) \parallel (6, 8)$, length now exactly 5

Clipping is only half the cure

Fixes exploding — caps the step size so a single huge gradient cannot blow up training.

Does nothing for vanishing — you cannot rescale a gradient that has already decayed to ≈ 0 .

to carry gradients across long lags we must change the **architecture → gates**

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/vanishing-gradients

Sweep the recurrent weight w and the lag T ; watch the long-range gradient collapse for $|w| < 1$, blow up for $|w| > 1$, and stay flat once a gate holds the memory open.

The LSTM

The vanilla RNN **overwrites** its state every step: $h_t = \varphi(W_{hh}h_{t-1} + \dots)$.

Every step multiplies the old state by W_{hh} and squashes it — memory is **rewritten**, and (Part V) the gradient is **multiplied** away. We want to **choose** what to keep.

let the network **decide**, per component: what to **forget**, what to **write**,
what to **expose**

Two states, one additive highway

D

The LSTM (Hochreiter & Schmidhuber, 1997) carries **two** vectors:

cell state c_t — the long-term memory (the “highway”).

hidden state h_t — the exposed output, a filtered view of c_t .

The cell updates by **addition**, not full overwrite:

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

keep a gated fraction of the old memory, **add** a gated bit of new — this is the whole idea

A **gate** is a vector in $[0, 1]$, computed by a sigmoid from the input and previous state:

$$\text{gate} = \sigma(W u_t + b) \in (0, 1)^m, \quad u_t = [h_{t-1} ; x_t]$$


0 = block the channel, 1 = let it pass — applied **elementwise** with $\odot$

Shared input $u_t = [h_{t-1}; x_t]$; three sigmoid gates and one tanh candidate:

$$f_t = \sigma(W_f u_t + b_f), \quad i_t = \sigma(W_i u_t + b_i), \quad o_t = \sigma(W_o u_t + b_o)$$

$$\tilde{c}_t = \tanh(W_c u_t + b_c) \quad (\text{candidate memory})$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t, \quad h_t = o_t \odot \tanh(c_t)$$



cell state runs straight across the top; gates tap in to forget, write, and read out

Gate	Operation	Meaning
forget f_t	$f_t \odot c_{t-1}$	1 = keep this memory, 0 = erase it
input i_t	$i_t \odot \tilde{c}_t$	how much of the new candidate to write
output o_t	$o_t \odot \tanh(c_t)$	how much of the memory to expose as h_t

candidate $\tilde{c}_t = \tanh(\dots)$ proposes **what** to write; the gates decide **how much**

The cell state's step-to-step Jacobian is just the **forget gate** — no weight matrix:

$$\frac{\partial c_t}{\partial c_{t-1}} = \text{diag}(f_t)$$

So over many steps the gradient along the cell highway is a product of **gates**:

$$\frac{\partial c_T}{\partial c_k} = \prod_{j=k+1}^T \text{diag}(f_j)$$

if $f_j \approx 1$, the gradient flows **undamped — no repeated $\times W_{hh}$**

the central LSTM advantage: an **open** forget gate is a gradient highway across time

One cell, old memory $c_{t-1} = 5$. **Retain** case — forget open, input nearly shut:

$$f_t = 0.8, i_t = 0.2, \tilde{c}_t = 3 : c_t = 0.8 \cdot 5 + 0.2 \cdot 3 = 4.6$$

Replace case — forget nearly shut, input open, candidate negative:

$$f_t = 0.1, i_t = 0.9, \tilde{c}_t = -2 : c_t = 0.1 \cdot 5 + 0.9 \cdot (-2) = -1.3$$

same equation: gates near $(1, 0)$ **hold** memory ($5 \rightarrow 4.6$); near $(0, 1)$ **overwrite** it ($5 \rightarrow -1.3$)

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/lstm-gates

Drag the forget, input and output gates and watch the cell state c_t retain, overwrite, or leak — and see the gradient highway open and close as f_t moves between 0 and 1.

The GRU

A simpler gated cell

The LSTM works but has **three** gates, a candidate, and **two** states. Can we do less?


The GRU (Cho et al. 2014, in the RNN encoder–decoder for translation) keeps the gating idea but uses **one** state and **two** gates.

Update gate z_t , reset gate r_t , candidate $\tilde{h}_t$, then a **convex combination**:

$$z_t = \sigma(W_z x_t + U_z h_{t-1} + b_z), \quad r_t = \sigma(W_r x_t + U_r h_{t-1} + b_r)$$

$$\tilde{h}_t = \tanh(W_h x_t + U_h (r_t \odot h_{t-1}) + b_h)$$

$$h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t$$



reset r_t gates the past **into** the candidate; update z_t mixes old state and candidate

update z_t : $h_t = (1 - z_t)h_{t-1} + z_t\tilde{h}_t$.
 $z = 0 \rightarrow$ copy old state; $z = 1 \rightarrow$ take candidate.

reset r_t : scales h_{t-1} **inside** the candidate.
 $r = 0 \rightarrow$ ignore the past, propose from x_t alone.

z_t plays the LSTM's forget + input as **one** knob; $z \approx 0$ is the gradient highway

One unit, old state $h_{t-1} = 4$, candidate $\tilde{h}_t = 0$. **Keep** — update nearly shut:

$$z_t = 0.25 : \quad h_t = 0.75 \cdot 4 + 0.25 \cdot 0 = 3$$

Replace — update wide open:

$$z_t = 0.9 : \quad h_t = 0.1 \cdot 4 + 0.9 \cdot 0 = 0.4$$

z small holds the state ($4 \rightarrow 3$); z large pulls it to the candidate ($4 \rightarrow 0.4$)

	LSTM	GRU
states	two: h_t and c_t	one: h_t
gates	forget, input, output	update, reset
memory path	additive cell c_t , Jacobian $\text{diag}(f_t)$	convex mix, carry $(1 - z_t)$
parameters	more (4 weight blocks)	fewer (3 weight blocks)

both fix vanishing gradients; GRU is lighter, LSTM sometimes stronger — **try both, it is task-dependent**

INTERACTIVE

(to build) — move the update and reset gates and watch the state interpolate between “keep the past” and “take the candidate”, side by side with the LSTM.

Sequence tasks with RNNs

Feed characters in, predict the **next** character at every step — the makemore idea, recurrent:

t	1	2	3	4
input	a	l	i	c
target	l	i	c	e

$p_t = \text{softmax}(W_o h_t + b_o)$; generate by **sampling** $\hat{x}_t \sim p_t$ and feeding it back

Many-to-one: read the whole sequence, predict **once**.

last state: $\hat{y} = \text{softmax}(W h_T + b)$ —
the final h_T summarizes everything.

mean pool: $\hat{y} = \text{softmax}(W \bar{h} + b)$,
 $\bar{h} = \frac{1}{T} \sum_t h_t$ — robust to length.

sentiment, topic, spam — one label from a variable-length input

For **non-causal** tasks, run two RNNs — forward and backward — and concatenate:

$$\vec{h}_t = \text{RNN}_{\rightarrow}(x_{1:t}), \quad \overleftarrow{h}_t = \text{RNN}_{\leftarrow}(x_{t:T}), \quad h_t = [\vec{h}_t ; \overleftarrow{h}_t]$$

Each h_t now sees the **whole** sequence — great for tagging or classification. But it peeks at the **future**, so it is **illegal for autoregressive generation**.

Batches mix lengths — pad to a common T , then **mask the padding out of the loss**:

$$\mathcal{L} = \frac{\sum_{i,t} m_{it} (-\log p_{it})}{\sum_{i,t} m_{it}}, \quad m_{it} \in \{0, 1\}$$

$m_{it} = 1$ for real tokens, 0 for <pad>. The mask makes padding invisible to the gradient — otherwise the model learns to predict <pad>.

Training RNNs in practice

The training loop

Detach for truncated BPTT, mask the padding, clip the gradient:

```
h = model.init_hidden(batch_size)
for x, y, mask in loader:
    h = h.detach()                # truncated BPTT: cut the graph
    logits, h = model(x, h)       # forward one chunk
    loss = masked_cross_entropy(logits, y, mask)  # ignore <pad>
    optimizer.zero_grad()
    loss.backward()               # BPTT within the chunk
    nn.utils.clip_grad_norm_(model.parameters(), 5.0)  # tame explosions
    optimizer.step()
```

the three RNN-specific lines: detach, masked loss, clip_grad_norm_

Symptom	Likely cause
loss goes NaN	exploding gradients — clip; lower the learning rate
predicts <pad> a lot	no loss mask on padded positions
memory grows each step	state not detached — full graph kept forever
no long-term dependence	vanilla RNN vanishing — use LSTM / GRU
generation repeats itself	greedy decoding — sample / temperature / top- k
training is slow	not packing / batching variable lengths

The RNN is inherently **sequential**: h_t needs h_{t-1} , so time cannot be parallelized.

RNN: $h_t = f(h_{t-1}, x_t)$ — a state carried step by step; $O(T)$ sequential.

TCN: stacked **dilated causal convolutions** — a fixed receptive field, fully **parallel** over time.

next lecture: convolutions over time, and then **attention** removes the recurrence entirely

Summary

Idea	Takeaway
recurrent state	$h_t = f_\theta(h_{t-1}, x_t)$ — shared f_θ , any length, memory in h_t
BPTT	unroll in time; shared-weight gradients accumulate over steps
vanilla RNN	$\partial h_T / \partial h_k \approx w^T$ — vanishing / exploding gradients
clipping	caps exploding norm; does nothing for vanishing
LSTM	cell $c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$; forget / input / output gates
GRU	one state, update + reset gates; simpler, often as good

RNNs **remember** by carrying a hidden state.

LSTMs and GRUs **learn when** to remember, forget, and update.

a gate is a controllable path — for information **and** for gradients

We gave the model a **memory** and gates to control it.

Next: **attention and seq2seq** — let every step look **directly** at every other, no recurrence required.