

Transformers I: Self-Attention, Positional Encoding and Transformer Blocks

Replacing recurrence with attention between all tokens

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

Why Transformers?

Where we are

Last lectures built sequence models on **recurrence** — RNNs, LSTMs, and seq2seq with **attention**:

$$h_t = f(h_{t-1}, x_t) \quad (\text{encoder: a state carried step by step})$$

$$c_u = \sum_{t=1}^T \alpha_{ut} h_t \quad (\text{decoder attends over encoder states } h_t)$$

Attention already lets the decoder look **directly** at any source position. This lecture asks: what if attention were the **whole** model — and we dropped recurrence entirely?

Attention fixed the **bottleneck**, but the encoder states h_t are still built **sequentially**:

$$h_1 \rightarrow h_2 \rightarrow h_3 \rightarrow \dots \rightarrow h_T \quad (\text{each } h_t \text{ waits for } h_{t-1})$$

$h_t = f(h_{t-1}, x_t)$ is **inherently serial**: you cannot compute step t before step $t - 1$. On modern parallel hardware (GPUs) this wastes the machine — training time grows with sequence length T .

we want the representation of every token computed in parallel

Replace the recurrent update with a function of **all** tokens at once:

$$\underbrace{h_t = f(h_{t-1}, x_t)}_{\text{RNN: depends on the previous state}} \longrightarrow \underbrace{z_i = f(x_1, \dots, x_T)}_{\text{Transformer: attends to all tokens}}$$

Each token builds its new representation by **directly attending** to every other token:

Self-attention relates positions **within the same sequence**. No state to carry — every z_i is computed independently and in parallel (Vaswani et al. 2017, “Attention Is All You Need”).


Three ways to mix information

	RNN	CNN	Transformer
context mechanism	hidden state h_t	receptive field	attention
reach per layer	all past (decayed)	fixed kernel k	all positions
parallel over positions?	no — serial	yes	yes
content-dependent mixing?	partly	no — fixed	yes

the Transformer keeps the CNN's parallelism **and** gains content-dependent, full-range mixing


Recurrence vs attention, side by side

path $i \rightarrow j$: $O(|i - j|)$ steps



RNN — sequential

path $i \rightarrow j$: $O(1)$ — one layer



self-attention — all-to-all

the RNN routes information along a **chain**; self-attention wires **every token to every token** in one layer

Attention as a soft lookup

Attention is a **differentiable dictionary lookup**. Given a query q and T key–value pairs:

$$s_i = q^\top k_i \quad (\text{score: how well the query matches key } i)$$

$$\alpha = \text{softmax}(s) \quad o = \sum_{i=1}^T \alpha_i v_i$$


a **soft** lookup: instead of picking one entry, return a weighted blend of **all** values

Three learned roles, one per vector:

Vector	Question it answers	Role
query q	what am I looking for?	probe
key k_i	what do I contain?	address
value v_i	what do I pass on?	content

match q against every k_i to decide **how much** of each v_i to retrieve — all differentiable

Attention as a soft dictionary



compare query to keys $\rightarrow$ scores $\rightarrow$ softmax weights $\alpha \rightarrow$ blend the values into o

$q = (1 \ 0)$, three keys, three values:

$$k_1 = (1 \ 0), \quad k_2 = (0 \ 1), \quad k_3 = (1 \ 1)$$

$$v_1 = (10 \ 0), \quad v_2 = (0 \ 10), \quad v_3 = (10 \ 10)$$

Scores are dot products of the query with each key:

$$s = (q^\top k_1, q^\top k_2, q^\top k_3) = (1, 0, 1)$$

Softmax the scores $s = (1, 0, 1)$:

$$\alpha = \text{softmax}(1, 0, 1) = \frac{e^1, e^0, e^1}{e^1 + e^0 + e^1} \approx (0.422, 0.155, 0.422)$$

Blend the values with these weights:

$$o = 0.422 v_1 + 0.155 v_2 + 0.422 v_3$$

$$o \approx (8.45, 5.78)$$

keys 1 and 3 match the query, so o leans toward v_1 and v_3 — a soft blend, not a hard pick

Stack the queries, keys and values into matrices and do it all at once:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V$$

this single expression is the central operation of the Transformer

$QK^\top$ is every query dotted with every key; softmax normalizes each row; the result weights V

Why divide by $\sqrt{d_k}$?

Suppose query and key entries are independent with mean 0, variance 1. Then a single score

$$q^\top k = \sum_{i=1}^{d_k} q_i k_i \quad \Rightarrow \quad \text{Var}(q^\top k) = d_k$$

Large $d_k \rightarrow$ scores of magnitude $\sim \sqrt{d_k} \rightarrow$ softmax saturates into a near one-hot spike $\rightarrow$ gradients through it vanish. Learning stalls.

divide by $\sqrt{d_k}$ to hold the score variance at ≈ 1 , whatever d_k is

One self-attention layer, sequence length T , model width d , head width d_k :

Tensor	Shape	How
input X	$T \times d$	one row per token
W_Q, W_K, W_V	$d \times d_k$	learned projections
Q, K, V	$T \times d_k$	$Q = XW_Q$, etc.
scores $A = \text{softmax}(QK^\top / \sqrt{d_k})$	$T \times T$	token-to-token weights
output $O = AV$	$T \times d_k$	attended representations

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/attention

Drag the query and watch the scores $q^\top k_i$, the softmax weights α_i , and the blended output $o = \sum_i \alpha_i v_i$ update — the worked example above, made live.

Self-attention

In **self**-attention, all three come from the **same** input X — each token is a query, a key, **and** a value:


$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

Each token builds a **new representation** by mixing information from the other tokens it finds relevant — a query looks around, keys advertise, values are gathered.

same operation as before, but the sequence attends **to itself**

The attention matrix

A_{ij} = how much token i (query) attends to token j (key); every **row sums to 1**:



a $T \times T$ matrix — each row is a distribution over which tokens to gather from

Take $W_Q = W_K = W_V = I$, so $Q = K = V = X$ with $X = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{pmatrix}$. The scores $QK^\top = XX^\top$:

	x_1	key k_j x_2	x_3
x_1	1	0	1
query q_i x_2	0	1	1
x_3	1	1	2

the diagonal grows with each token's norm; scale by $\sqrt{d_k} = \sqrt{2}$ before the softmax next

Scale by $\sqrt{2}$; row 1 (token 1 as query) before softmax is $(1 \ 0 \ 1)/\sqrt{2} = (0.71 \ 0 \ 0.71)$:

$$\alpha_1 = \text{softmax}(0.71, 0, 0.71) \approx (0.40, 0.20, 0.40)$$

$$z_1 = 0.40 x_1 + 0.20 x_2 + 0.40 x_3 \approx (0.80 \ 0.60)$$

token 1 attends to tokens 1 and 3 more than 2 — its output is a weighted average of their values

Self-attention is content-dependent

Compare how a CNN and self-attention decide **which positions** interact:

CNN: the receptive field is **fixed by the kernel** — position i always mixes with $i - 1, i, i + 1$, regardless of content.

Self-attention: A_{ij} depends on the **token representations** — each token **decides** which others matter, per input.

same weights, different wiring for every sentence — the mixing pattern is learned and dynamic

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/attention

Edit a short sentence and watch the $T \times T$ attention matrix light up — each row a softmax distribution over which tokens that position pulls information from.

Masked (causal) self-attention

The future-leakage problem

For a **language model**, position t must predict x_{t+1} from **only** the past:

$$p(x_{t+1} \mid x_{\leq t}) \quad (\text{it may look at } x_1, \dots, x_t \text{ — never ahead})$$

Plain self-attention lets token t attend to **every** position, including $x_{t+1}, x_{t+2}, \dots$ — it would peek at the answer. That is **illegal** for autoregressive generation.

Compute the scores, then set the **future** entries to $-\infty$ before the softmax:

$$S = \frac{QK^\top}{\sqrt{d_k}}, \quad S_{ij} \leftarrow -\infty \text{ if } j > i$$

Softmax turns $-\infty$ into a weight of 0:

$$A_{ij} = 0 \quad \text{for all } j > i$$

each row now normalizes over the **past and present only** — the future contributes nothing

The causal mask, drawn

$T = 5$: allowed entries (lower triangle) vs masked future ($-\infty$, upper triangle):

		key position j				
		1	2	3	4	5
query position i	1	✓	$-\infty$	$-\infty$	$-\infty$	$-\infty$
	2	✓	✓	$-\infty$	$-\infty$	$-\infty$
	3	✓	✓	✓	$-\infty$	$-\infty$
	4	✓	✓	✓	✓	$-\infty$
	5	✓	✓	✓	✓	✓

query i may attend to key j only when $j \leq i$ — a lower-triangular attention pattern

The **same** self-attention op, two masking regimes:

Setting	Mask	Sees
encoder (BERT-style)	none	bidirectional — the whole sequence
decoder (GPT-style)	causal	only the past $x_{\leq t}$

bidirectional is great for **understanding**; causal is required for **generation**

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/attention

Toggle the causal mask on a short sequence and watch the upper triangle of the attention matrix go dark — the future becomes unreachable, one row at a time.

Multi-head attention

Why more than one head?

A single attention pattern can track **one** kind of relation at a time. Use H of them in parallel:

$$\text{head}_i = \text{Attention}(XW_i^Q, XW_i^K, XW_i^V)$$

$$\text{MHA}(X) = \begin{pmatrix} \text{head}_1 \\ \vdots \\ \text{head}_H \end{pmatrix} W_O \quad (\text{concatenate, then mix with } W_O)$$

each head projects into its **own** subspace and learns a **different** relation

The classic setup: $d_{\text{model}} = 512$, $H = 8$ heads:

$$d_k = d_v = \frac{d_{\text{model}}}{H} = \frac{512}{8} = 64$$

Per head	Shape	Note
W_i^Q, W_i^K, W_i^V	512×64	projects into a 64-dim subspace
head_i	$T \times 64$	attention within that subspace
concat of 8 heads	$T \times 512$	back to full width
W_O	512×512	mixes the heads

Different heads specialize — some interpretable patterns observed in trained models:

- attend to the **previous** word
- **subject-verb** agreement
- matching **brackets** / quotes

- local **phrase** structure
- **positional** patterns
- coarse **semantic** similarity

Caution. Individual-head interpretation is **noisy** and post-hoc — heads are not cleanly labeled functions, and behavior is distributed across many of them.

Combining all H heads, the per-projection weights stack into full $d \times d$ matrices.
With $d = 512$:


$$W_Q, W_K, W_V, W_O \in \mathbb{R}^{512 \times 512}$$

$$4 \times 512^2 = 4 \times 262,144 =$$

1,048,576 parameters

≈ 1.05 M — this is one attention sublayer; the FFN typically adds a few times more

The multi-head block



H attentions run in parallel on the same X , concatenate, then W_O mixes them back to width d

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/attention

Switch between heads and see each one attend to a **different** relation on the same sentence — then watch their outputs concatenate and mix through W_O .

Positional information

Attention is a **weighted sum** — it depends on **content**, not position. Shuffle the tokens and:

$$\text{SelfAttention} \left(\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} \right) \text{ is a permutation of } \text{SelfAttention} \left(\begin{pmatrix} x_2 \\ x_1 \\ x_3 \end{pmatrix} \right)$$

Self-attention is **permutation-equivariant**: with no position signal, the model cannot tell “dog bites man” from “man bites dog”. Order is invisible.

we must inject position information into the token representations

Give every position t a vector p_t and **add** it to the token embedding:

$$r_t = e_t + p_t \quad (\text{token content} + \text{where it sits})$$

Now two identical tokens at different positions get **different** inputs r_t — attention can use the position signal, and permuting the sequence changes the output.

The simplest choice: a **lookup table** of position vectors, trained like any embedding:

$$P \in \mathbb{R}^{T_{\max} \times d}, \quad p_t = P_{t,:}$$

Pro — simple, flexible, learns whatever the data needs (GPT / BERT use this).

Con — undefined for $t > T_{\max}$; cannot extrapolate to sequences longer than seen in training.

The original Transformer uses a **fixed**, deterministic pattern of sines and cosines:


$$\text{PE}_{(\text{pos}, 2i)} = \sin\left(\frac{\text{pos}}{10000^{2i/d}}\right)$$

$$\text{PE}_{(\text{pos}, 2i+1)} = \cos\left(\frac{\text{pos}}{10000^{2i/d}}\right)$$

dimension $2i$ oscillates at wavelength $2\pi \cdot 10000^{2i/d}$ — no parameters, defined for **any** position

Many frequencies, many scales

Low dimensions oscillate **fast**, high dimensions **slow** — position encoded at every scale:



left: sinusoids at different frequencies · right: PP^T — nearby positions have **similar** encodings

Take $d = 4$, so $i \in \{0, 1\}$ and the exponents are $10000^0 = 1$ and $10000^{1/2} = 100$:

$$\text{PE}(\text{pos}) = \left(\sin\left(\frac{\text{pos}}{1}\right) \quad \cos\left(\frac{\text{pos}}{1}\right) \quad \sin\left(\frac{\text{pos}}{100}\right) \quad \cos\left(\frac{\text{pos}}{100}\right) \right)$$

At $\text{pos} = 0$ every sine is 0 and every cosine is 1:

$$\text{PE}(0) = (0 \ 1 \ 0 \ 1)$$

a unique, smooth signature per position — added onto the token embedding

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/positional-encoding

Slide the position and watch each sinusoidal dimension turn; compare learned vs sinusoidal encodings and see how nearby positions stay similar while distant ones drift apart.

The Transformer block

One block stacks two sublayers, each wrapped in a residual connection and a normalization:

multi-head self-attention $\rightarrow$ residual + norm $\rightarrow$ FFN $\rightarrow$ residual + norm

Attention mixes information **across positions**; the **FFN** then processes **each position** on its own. Communication, then computation.

Applied **independently and identically** to every position — a two-layer MLP:

$$\text{FFN}(x) = W_2 \varphi(W_1 x + b_1) + b_2$$

- same weights W_1, W_2 at **every** position (a 1×1 transform over the sequence);
- original sizing: $d_{\text{model}} = 512$ expands to $d_{\text{ff}} = 2048$, then back to 512.

φ is usually ReLU or GELU; the hidden layer is $4 \times$ wider than the model

Why an FFN **after** attention? The two sublayers do **complementary** jobs:

Attention = **communication**. Tokens exchange information — each position gathers from others.

FFN = **computation**. Each token, now informed, is transformed on its own — no mixing.

attention decides **what to gather**; the FFN decides **what to do** with it, tokenwise

Each sublayer is wrapped so gradients flow and activations stay well-scaled:


post-norm : $x' = \text{LN}(x + \text{MHA}(x)), \quad y = \text{LN}(x' + \text{FFN}(x'))$

pre-norm : $x' = x + \text{MHA}(\text{LN}(x)), \quad y = x' + \text{FFN}(\text{LN}(x'))$

pre-norm keeps a clean residual highway — more stable for **deep** stacks (the modern default)

The Transformer encoder block

the horizontal spine is the **residual stream** (LN = LayerNorm)



the residual stream runs straight through; each sublayer **branches off, computes, and adds back**


A full encoder–decoder decoder layer has **three** sublayers:

1. **masked** self-attention over the output so far (causal);
2. **cross-attention** — queries from the decoder, keys and values from the **encoder**;
3. a position-wise FFN.

cross-attention is how the decoder reads the source — we unpack encoder–decoder next lecture

The decoder-only (GPT-style) block

Drop the encoder and cross-attention: **causal self-attention + FFN**, then an LM head:



output $p(x_{t+1} \mid x_{\leq t}) = \text{softmax}(W_o h_t)$ – this is the GPT family

INTERACTIVE

(to build) — send a sequence through one block and watch the residual stream: attention mixes across positions, the FFN transforms each, and the norms keep it stable.

Complexity and comparison

The attention matrix is $T \times T$ — quadratic in sequence length:

$$QK^\top : O(T^2 d) \text{ compute}, \quad A \in \mathbb{R}^{T \times T} : O(T^2) \text{ memory}$$

Every token attends to every token, so cost grows with T^2 . This is the **long-context bottleneck** — the price of full-range, content-dependent mixing.

fine for $T \sim 10^3$; the target of much later work (efficient / sparse attention) for very long T

Path length: how far is information?

How many steps must a signal travel between positions i and j ?

Model	Path length $i \rightarrow j$
RNN	$O(i - j)$ — along the chain
dilated CNN	$O(\log i - j)$ — up the dilation tree
self-attention	$O(1)$ — one hop, any distance

short paths = easy gradient flow = long-range dependencies are learnable

	RNN	TCN (CNN)	Transformer
parallel over time	low	high	high
longest path $i \rightarrow j$	$O(T)$	$O(\log T)$	$O(1)$
compute per layer	$O(Td^2)$	$O(Tkd^2)$	$O(T^2d)$
context	state (decays)	fixed field k	content-dependent, full


the Transformer trades $O(T^2)$ cost for **parallelism** and **constant** path length

- **parallel training** — no sequential dependency, the whole sequence at once;
- **dynamic context** — content-dependent attention, not a fixed window;
- **short paths** — $O(1)$ between any two positions, so long-range signals survive;
- **stackable** — residual + norm blocks compose into very deep, scalable models.

the same architecture scales from a toy LM to the largest models trained today

A tiny decoder-only language model

A minimal GPT: turn <BOS> I like into a prediction of the next token:



token ids $\rightarrow$ embeddings + position $\rightarrow$ causal self-attention $\rightarrow$ FFN $\rightarrow$ logits $\rightarrow$ softmax

The logits at position t predict the token at $t + 1$ — the **same next-token objective** as before:

$$\mathcal{L} = - \sum_{t=1}^{T-1} \log p_{\theta}(x_{t+1} \mid x_{\leq t})$$

Nothing new in the loss — token-level cross-entropy. What changed is **how** p_{θ} is computed: causal self-attention over the prefix, not a recurrent state.

Feed the sequence in; the target is the **same sequence shifted left by one**:

position t	1	2	3	4
input x_t	<BOS>	I	like	cats
target x_{t+1}	I	like	cats	<EOS>

the **causal mask** is what makes this safe — position t cannot see its own target x_{t+1}

INTERACTIVE

(to build) — type a short prefix and watch it flow through embeddings, causal attention and the FFN to a next-token distribution; sample a token and feed it back.

Summary

What changed from RNNs

RNN	Transformer
$h_t = f(h_{t-1}, x_t)$	$H' = \text{SelfAttention}(H)$
memory carried through state	interactions computed directly
serial over time	parallel over positions
path $i \rightarrow j$ is $O(i - j)$	path $i \rightarrow j$ is $O(1)$

recurrence **routes** information step by step; attention **relates** all positions at once

Step	Equation
project	$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$
attend	$A = \text{softmax}(QK^\top / \sqrt{d_k} + \text{mask})$
gather	$O = AV$
multi-head	$\text{MHA} = \begin{pmatrix} \text{head}_1 \\ \vdots \\ \text{head}_H \end{pmatrix} W_O$
block	$\text{MHSA} \rightarrow \text{residual} + \text{norm} \rightarrow \text{FFN} \rightarrow \text{residual} + \text{norm}$

Final mental model — one idea

Self-attention: every token asks the others for information.

Multi-head: several relation types, in parallel.

Positional encoding: puts the order back in.

a Transformer block = attention + tokenwise MLP + residual / norm

We replaced recurrence with **attention between all tokens.**

Next — **Transformers II**: encoder-only (BERT), decoder-only (GPT), encoder–decoder; pretraining, finetuning and generation.