

Transformers II: BERT, GPT, Pretraining, Finetuning and Generation

Same block, different mask + objective + head $\Rightarrow$ different model family

Prof. Nipun Batra

ES 667 · Deep Learning · IIT Gandhinagar

One block, three families

Where we are — recall Transformers I

Last lecture built the **Transformer block**: self-attention over all tokens, no recurrence.

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^\top}{\sqrt{d}} + \text{mask} \right) V$$

One block = multi-head self-attention $\rightarrow$ residual + norm $\rightarrow$ FFN $\rightarrow$ residual + norm. This lecture: the **same block** becomes three different model **families**.

The whole lecture in one line

We keep the Transformer block **fixed** and change three knobs:

mask + **objective** + **head** $\Rightarrow$ **model family**

Knob	Controls	Choices
attention mask	what a token may see	bidirectional / causal / cross
pretraining objective	what the model learns to do	masked / next-token / denoise
output head	how we read it out	[CLS] / LM head / decoder

architecture controls information flow; objective controls what it learns


The three families v

	Encoder-only	Decoder-only	Encoder-decoder
mask	bidirectional	causal	enc: bidir · dec: causal
objective	masked-token	next-token	seq2seq / denoise
reads	both sides	left context	source → target
best at	understanding	generation	transduction
example	BERT	GPT	T5 · BART


three columns, one underlying block — the differences are the **mask**, the **objective** and the **head**

One picture, three families **v**

encoder-only



decoder-only




enc-decoder



encoder-only reads a whole sequence into **contextual reps**; decoder-only **predicts the next token**; encoder-decoder maps a **source** to a **target**

The masks that define them v

The mask is the single biggest lever — it decides who may attend to whom:



bidirectional (BERT): all-to-all · causal (GPT): only the past · cross-attention: decoder reads encoder

Interactive: model-family masks I

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/attention

Toggle between **bidirectional**, **causal** and **cross-attention** masks on a short sequence and watch which entries of the $T \times T$ attention matrix light up — the same operation, three information-flow patterns.

Encoder-only: BERT

The encoder-only Transformer

Feed all tokens in at once; every token attends **every** non-padding token

The encoder-only Transformer

no causal mask:

$$x_1, \dots, x_T \xrightarrow{\text{Encoder}} h_1, \dots, h_T$$

- each h_t is a **contextual representation** of token t — informed by the whole sequence;
- no left-to-right restriction: position t may look **left and right**.

the output is not a prediction — it is a **representation per token** to feed a task head

What BERT is good for **v**

Encoder-only models shine at **understanding** — mapping text to a fixed target:

Task	Shape
sentence classification / sentiment	$x_{1:T} \rightarrow y$
natural language inference (NLI)	pair $\rightarrow \{\text{entail, contradict, neutral}\}$
extractive QA	passage $\rightarrow$ answer span
token classification / NER	$x_{1:T} \rightarrow y_{1:T}$
retrieval / embeddings	text $\rightarrow$ vector

whenever the output is a **label** or a **span**, both-side context helps

Why not next-token LM? D

A causal LM only ever sees the **left**:

$$p(x_t \mid x_{<t}) \quad (\text{left-to-right only})$$

But **understanding** a word often needs **both sides**:

“The **bank** approved the loan” vs “The **bank** was flooded”

The word after bank disambiguates it. A left-only model has not seen it yet when it encodes bank. Bidirectional context lets the representation use **the whole sentence**.

Masked language modeling (MLM)

BERT's trick: **corrupt** some tokens, then predict the originals from **both sides**:

the cat sat on the [MASK] → predict mat

$$p_{\theta}(x_m \mid x_{\setminus m}) \quad (\text{recover masked token } m \text{ from all the others})$$

Loss is applied **only on the masked positions** M :

$$\mathcal{L}_{\text{MLM}} = - \sum_{m \in M} \log p_{\theta}(x_m \mid \tilde{x})$$

$\tilde{x}$ is the corrupted sequence · unmasked tokens get **no** loss, they are only context

MLM vs next-token, side by side v

MLM (BERT): predict [MASK] from both sides



CLM (GPT): predict next token from the left only



MLM predicts a hidden token from **both** neighbours; a causal LM predicts the **next** token from the **left** only

MLM worked example

Sequence deep learning is useful; mask position 2:

```
deep [MASK] is useful  target: learning
```

The encoder sees deep, is, useful (both sides) and scores the vocabulary at position 2:

$$\mathcal{L} = -\log p_{\theta}(\text{learning} \mid \tilde{x})$$

a single V -way classification at the masked slot — using **right** context (is useful) too

BERT input format

Two sentences packed into one sequence with special tokens:

[CLS] sentence A [SEP] sentence B [SEP]

Each token's input embedding is a **sum of three** embeddings:

$$r_t = e_t + p_t + s_t$$

Component	Symbol	Encodes
token	e_t	which word / subword
position	p_t	where it sits
segment	s_t	sentence A vs B

The [CLS] representation

A special [CLS] token sits at the front; its final state summarizes the **whole** sequence:

$$h_{[\text{CLS}]} \xrightarrow{\text{classifier}} p(y \mid x) = \text{softmax}(W h_{[\text{CLS}]} + b)$$

Attention lets [CLS] gather from every token, so its representation is a natural **pooled** summary. A small head on top of $h_{[\text{CLS}]}$ gives sentence-level predictions.

classification, NLI, sentiment — all read out from this one vector

Token classification

Some tasks want a label at **every** position — put the head on **each** h_t :

$$p(y_t \mid x) = \text{softmax}(W h_t + b) \quad (\text{one prediction per token})$$

- named-entity recognition, part-of-speech tagging;
- a **padding mask** keeps padded positions out of the loss.

bidirectional context helps: tagging a word can use the words **after** it

Answer = a **span** in the passage. Learn a **start** and an **end** pointer over tokens:

$$p_{\text{start}(t)} = \text{softmax}(w_s^\top h_t), \quad p_{\text{end}(t)} = \text{softmax}(w_e^\top h_t)$$

Train to point at the true start s^* and end e^* :

$$\mathcal{L} = -\log p_{\text{start}(s^*)} - \log p_{\text{end}(e^*)}$$

two little heads over the same contextual states — no text is generated, just **located**

The BERT finetuning pattern

Every downstream task is the **same** recipe:

pretrained encoder → task head → supervised loss

Start from the MLM-pretrained weights; bolt on a small task head ($h_{[CLS]}$ or per-token); train on labelled data. Typically **all** parameters update.

one pretrained backbone, many cheap heads — the pretrain-then-finetune paradigm

RoBERTa: the recipe matters

Same architecture and MLM objective as BERT — but a **better training recipe**:

Changes — more data, longer training, bigger batches, dynamic masking; dropped the next-sentence task.

Result — a large jump on the same benchmarks, with **no** architectural change.

the objective matters, but so does the **recipe: data, scale, and training choices**

Decoder-only: GPT

The decoder-only Transformer

Causal self-attention: token t may attend only to $x_{\leq t}$, and predicts the next token:

$$h_t = f(x_{\leq t}), \quad p(x_{t+1} \mid x_{\leq t}) = \text{softmax}(W_o h_t)$$

Drop the encoder and cross-attention from the Transformer: what remains is a stack of **causal self-attention + FFN** blocks, capped by an LM head.

this is the **same** next-token model of Lecture 11 — now with attention instead of a fixed window

The language-modeling objective

Factorize the sequence by the chain rule, then take negative log-likelihood:

$$p(x_{1:T}) = \prod_{t=1}^T p(x_t \mid x_{<t})$$

$$\mathcal{L}_{\text{LM}} = - \sum_{t=1}^T \log p_{\theta}(x_t \mid x_{<t})$$

token-level cross-entropy — **every** position is a training example, computed in parallel

The causal mask

Compute scores, then forbid attending to the future **before** the softmax:

$$A_{ij} = 0 \quad \text{for} \quad j > i$$

Set $S_{ij} = -\infty$ when $j > i$; softmax turns those into 0. Each row normalizes over the **past and present only** — no token can peek at its own target.

this lower-triangular pattern is **the** difference between GPT and BERT attention

Shifted labels

Feed the sequence in; the target is the **same sequence shifted left by one**:

position t	1	2	3	4
input x_t	<BOS>	deep	learning	is
target x_{t+1}	deep	learning	is	useful

the causal mask makes this safe — position t predicts $t + 1$ without seeing it

GPT-style pretraining D

Pretrain the decoder to predict the **next token** on a large corpus. Then adapt:

Adaptation	What it means
prompting	condition on text; read the continuation
finetuning	continue training on task data
instruction tuning	finetune on (instruction → response) pairs
RLHF (later)	align to human preferences

GPT-1 introduced **generative pretrain + finetune**; the objective never changes — only the adaptation

GPT-2: LM as unsupervised multitask D

Train a **large** decoder-only model on broad web text — then ask it to do tasks with **no** finetuning:

A big enough next-token model performs many tasks **zero-shot**, purely from a prompt. Language modeling on diverse text is **implicitly** multitask learning (Radford et al.).

scale + broad data $\Rightarrow$ a single LM that generalizes to unseen tasks from prompts alone

Prompting as conditioning

Prompting is **not** an architecture change — it is just **conditioning the LM**:

$$p(\text{continuation} \mid \text{prompt})$$

Translate to French: cat => **(the model continues)** chat

the prompt sets the left context; the **same** next-token machinery produces the answer

Few-shot prompting

Put a few **demonstrations** directly in the context, then the query:

```
great movie => positive  
awful plot => negative  
loved every minute => ?
```

No weights change. The model **reads the pattern** from the examples in its context window and continues it — “in-context learning”.

Interactive: in-context learning I

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/in-context-learning

Add demonstrations to a prompt and watch a decoder-only model shift its next-token distribution toward the pattern you set — few-shot prompting with **no** gradient steps.

Classification via prompting

Recast a label task as **next-token scoring**:

Review: a dull, lifeless film. Sentiment: $\rightarrow$ negative


Compare the LM's probability of each candidate label token:

$$\hat{y} = \arg \max_{y \in \{\text{positive}, \text{negative}\}} p_{\theta}(y \mid \text{prompt})$$

the classifier **is** the language model — pick the label it finds most probable

The generation loop **v**

Decoder-only generation is one loop, run token by token:



tokenize → forward → logits → **decode rule** → append → repeat (the rule is Part VII)

Encoder–decoder

The architecture

Two stacks: an encoder reads the **source**, a decoder generates the **target**:

$$H = \text{Encoder}(x), \quad p(y_u \mid y_{<u}, x) = \text{softmax}(W_o s_u)$$

The decoder layer has **three** sublayers:

1. **causal** self-attention over the target prefix $y_{<u}$;
2. **cross-attention** into the encoder states H ;
3. a position-wise FFN.

self-attention reads what's generated so far; cross-attention reads the **source**

Cross-attention D


The bridge between the stacks: **queries from the decoder, keys and values from the encoder**:

$$Q = SW_Q \text{ (decoder)}, \quad K = HW_K, V = HW_V \text{ (encoder)}$$

$$\text{CrossAttn}(S, H) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right) V$$

each target position **asks the source** which parts to attend to — this is how translation aligns

Cross-attention, drawn $\mathbf{v}$



Q comes from the decoder state S ; K and V come from the encoder states H – the decoder reads the source

Good for sequence-to-sequence

Encoder–decoder fits tasks where output length $U \neq$ input length T :

Task	Source $\rightarrow$ target
translation	English sentence $\rightarrow$ German sentence
summarization	long document $\rightarrow$ short summary
question answering	question + context $\rightarrow$ answer
data-to-text	table / triples $\rightarrow$ sentence

a full re-encoding of the source, then free-length generation of the target

T5: text-to-text

Cast **every** task as text → text, so one model and one loss handle all of them:

```
translate English to German: That is good. → Das ist gut.
```

```
sst2 sentence: a dull film → negative
```

Classification, translation, summarization — all become “read this text, write that text”. A task prefix tells the model what to do (Raffel et al.).

BART: denoising seq2seq

Corrupt the input, then reconstruct the **original** — a bidirectional encoder feeds a causal decoder:

$$\tilde{x} = C(x) \xrightarrow{\text{Encoder-Decoder}} x$$

Encoder — bidirectional, reads the corrupted text (like BERT).

Decoder — left-to-right, reconstructs the clean text (like GPT).

BART unifies BERT-style understanding and GPT-style generation in one seq2seq model (Lewis et al.)

Pretraining objectives compared

Three objectives v

Each family is defined by **what it must predict**:

Objective	Input → predict	Architecture	Family
causal LM	prefix → next token	decoder-only	GPT
masked LM	corrupted → masked tokens	encoder-only	BERT
denoising seq2seq	corrupted → original	encoder-decoder	T5 · BART

same block; the **prediction target** is what differs

The three losses

causal LM : $\mathcal{L}_{\text{CLM}} = - \sum_t \log p(x_t \mid x_{<t})$

masked LM : $\mathcal{L}_{\text{MLM}} = - \sum_{m \in M} \log p(x_m \mid \tilde{x})$

denoising : $p(x \mid \tilde{x}) = \prod_t p(x_t \mid x_{<t}, \tilde{x}), \quad \tilde{x} = C(x)$

CLM $\rightarrow$ generation $\cdot$ MLM $\rightarrow$ understanding / bidirectional $\cdot$ denoising $\rightarrow$ reconstruct target from input

Span corruption (T5-style)

Instead of masking single tokens, mask **contiguous spans** and replace each with one sentinel:

input: The quick brown fox jumps over the lazy dog

corrupt: The quick <X> jumps over <Y> dog

target: <X> brown fox <Y> the lazy

the decoder emits the **dropped spans**, tagged by sentinel — shorter targets, harder task

Which objective for which task v


You want	Pick	Why
classify / tag / retrieve	MLM (encoder)	both-side context
open-ended generation	causal LM (decoder)	native next-token
translate / summarize	seq2seq (enc-dec)	source → target
one model for everything	text-to-text (T5)	unified interface

match the **objective** to the **shape** of your task

Finetuning and adaptation

Pretrain, then finetune v

The dominant paradigm: pretrain once on a huge corpus, then specialize cheaply:



$$\theta_0 \text{ (pretrained)} \rightarrow \theta^* = \arg \min_{\theta} \mathcal{L}_{\text{task}}(\theta)$$

the pretrained weights θ_0 are the **starting point**; finetuning nudges them to the task

Full finetuning

Update **all** parameters on the task data:

Pro — best adaptation; the whole network can reshape to the task.

Con — memory-heavy; can **overfit** small data; a **separate copy** of the model per task.

the strong default when you have enough labelled data and compute

Feature extraction (frozen backbone)

Freeze the pretrained encoder; train **only** the task head:

$$\underbrace{\theta_{\text{backbone}}}_{\text{frozen}} + \underbrace{\theta_{\text{head}}}_{\text{trained}}$$

- fast, cheap, tiny memory footprint;
- good for **small data** or a **quick baseline**;
- usually weaker than full finetuning.

the backbone is a fixed feature extractor; only a small head learns the task

Between frozen and full: update a **tiny** set of extra parameters, keep the backbone fixed.

Method	Idea
adapters	small bottleneck modules inserted per layer
prefix / prompt tuning	learn virtual tokens prepended to the input
LoRA	learn low-rank updates $\Delta W = BA$ to weights

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/lora-adapter

Watch a low-rank update $\Delta W = BA$ adapt a frozen weight matrix — a full lecture on **parameter-efficient** finetuning comes later.

Classification-head example

Finetune BERT for sentiment — one small head over [CLS]:

[CLS] the movie was wonderful [SEP]

$$z = W h_{[\text{CLS}]} + b, \quad \mathcal{L} = -\log \text{softmax}(z)_{y^*}$$

pretrained encoder + a linear head + cross-entropy — the whole finetuning setup

Generation finetuning

For a generation task (e.g. summarization), finetune with the **seq2seq** loss:

$$\mathcal{L} = - \sum_{u=1}^U \log p_{\theta}(y_u \mid y_{<u}, x)$$

- x is the document, y the reference summary;
- teacher forcing: feed the true $y_{<u}$ during training.

same next-token loss, now conditioned on the source x — the model learns to **transduce**

Decoding and generation

The decoding problem

Given the next-token distribution p_t , **how do we choose** the next token?

Training gives us $p_\theta(x_{t+1} \mid x_{\leq t})$. Decoding is a **separate** choice made at generation time — no weights change, but it hugely shapes the output.

the same model can sound repetitive or creative depending only on the **decoding rule**

Greedy decoding

Always take the single most probable token:

$$\hat{x}_{t+1} = \arg \max_v p_t(v)$$

Pro — simple, deterministic, fast; the model's local best guess.

Con — repetitive and bland; locally-best $\neq$ globally-best sequence.

Beam search

Keep the B highest-scoring **partial sequences** at each step, by cumulative log-probability:

$$\text{score}(y_{1:u}) = \sum_{k=1}^u \log p(y_k \mid y_{<k})$$

- explores several hypotheses in parallel, not just one;
- **good** for translation / summarization (a single high-likelihood answer);
- **poor** for open-ended text — tends to be generic and repetitive.

$B = 1$ is greedy; larger B searches wider but costs more

Rescale the logits by τ **before** the softmax to control sharpness:

$$p_i(\tau) = \text{softmax}(z/\tau)_i$$

τ	Effect
$\tau < 1$	sharper — bolder, more deterministic
$\tau = 1$	the model's own distribution
$\tau > 1$	flatter — more random, more diverse

$\tau \rightarrow 0$ recovers greedy; $\tau \rightarrow \infty$ approaches uniform

Top-k and nucleus (top-p)

Truncate the tail **before** sampling, so rare tokens can't derail generation:

Top- k — keep the k most probable tokens, renormalize, sample.

Nucleus (top- p) — keep the smallest set S with $\sum_{i \in S} p_i \geq p$, then sample.

top- p **adapts** the candidate set to the model's confidence — narrow when sure, wide when unsure

Repetition penalties

Explicit fixes for the degenerate repetition that sampling and greedy can produce:

- **no-repeat n-gram** — forbid emitting an n-gram already seen;
- **repetition / frequency penalty** — down-weight tokens already generated;
- **presence penalty** — discourage reusing any token that has appeared.

cheap, decoding-time nudges away from loops — no retraining

Interactive: the decoding lab I

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/softmax-temperature

Sweep **temperature**, **top- k** and **top- p** on a fixed set of logits and watch the next-token bar chart sharpen or flatten — then sample repeatedly and see the diversity of outputs change.

Worked examples

BERT MLM: the loss

Sentence I like [MASK]; candidates and BERT's predicted probabilities:

token	cats	blue	running
p	0.80	0.05	0.15

Target is cats; cross-entropy picks out its probability:

$$\mathcal{L} = -\log(0.80) \approx 0.223$$

a confident, correct prediction — small loss; a perfect one would give 0

GPT next-token: the loss D

Prompt I like; the decoder's next-token distribution:

token	cats	dogs	blue
p	0.6	0.3	0.1

Target is cats:

$$\mathcal{L} = -\log(0.6) \approx 0.511$$

same cross-entropy loss as MLM — the difference is **causal** context, not left-and-right

T5 text-to-text: the loss D

Prompt the encoder, generate the target with the decoder:

```
translate English to Hindi: I like cats → target sequence
```

Sum the per-token negative log-likelihood over the target, conditioned on the source:

$$\mathcal{L} = - \sum_{u=1}^U \log p_{\theta}(y_u \mid y_{<u}, x)$$

exactly the seq2seq loss — the **source** x conditions every target-token prediction

Choosing a model in practice

Classification D

Three valid routes — pick by data and infrastructure:

Approach	How
encoder-only	[CLS] → classifier head
decoder-only	prompt → score label tokens
encoder-decoder	input → generate label text

encoder-only is the classic strong, cheap default for a fixed label set

Need	Choice
open-ended text	decoder-only, sampling (top- p)
input-conditioned output	encoder-decoder
translation	beam search (high-likelihood answer)
creative / diverse	sampling with temperature

decode **deterministically** when there's one right answer; **sample** when diversity matters

Token tagging and small data D

Token tagging — encoder-only. Both-side context helps label each token:
 $h_t \rightarrow y_t$ (NER, POS).

Small data — start from a **pretrained** model; escalate frozen $\rightarrow$ PEFT $\rightarrow$ full only as needed.

always start pretrained; keep a **held-out validation set and watch for overfitting**

Limitations

Context length and cost

Self-attention is **quadratic** in sequence length, and context is **finite**:

$$O(T^2 d) \text{ compute,} \quad T \leq T_{\max}$$

Every token attends every token, so cost grows with T^2 . Models have a hard context limit $T_{\max}$; long inputs must be truncated or chunked.

efficient / long-context attention is a whole topic for a later lecture

Tokenization issues

The tokenizer silently shapes everything downstream:

- text is split into **subwords** (BPE) — not words or characters;
- **script differences** change how many tokens a language needs;
- sequence length in tokens depends on the tokenizer, not just the text.

Perplexity is not comparable across tokenizers — a character model and a subword model predict different-sized units. Never compare PPL across vocabularies.

Interactive: tokenization **I**

INTERACTIVE

↗ nipunbatra.github.io/interactive-articles/bpe-merges

Watch **byte-pair encoding** merge frequent character pairs into subword tokens — and see how the same sentence becomes a different number of tokens under different vocabularies.

Pretraining data matters

A pretrained model inherits its **corpus**:

- **biases** and toxicity present in the data;
- **coverage** — domains, languages, styles it did or didn't see;
- **factual staleness** — knowledge frozen at the training cutoff.

the model is a mirror of its data — know what it was trained on before you trust it

Evaluation depends on the task

No single number fits every model family:

Task type	Metrics
classification	accuracy · F1 · AUROC
generation	BLEU · ROUGE · BERTScore · human
language modeling	NLL · perplexity
retrieval	MRR · Recall@k · NDCG

choose the metric that matches the **task shape**, not the architecture

Summary

The model families, one table

	Encoder-only	Decoder-only	Encoder-decoder
mask	bidirectional	causal	bidir + causal
pretraining	masked LM	next-token	denoise / seq2seq
best fit	understanding	generation	transduction
example	BERT · RoBERTa	GPT	T5 · BART

mask × pretraining objective × head $\Rightarrow$ the family and its sweet spot

The objectives, in one breath

- **BERT** — predict the **missing tokens** using **both sides** (understanding);
- **GPT** — predict the **next token** from the **previous** ones (generation);
- **T5 / BART** — **reconstruct or generate** the target from a corrupted / source input (transduction).

one Transformer block; three ways to train it, three things it learns to do

Final mental model — three knobs

Architecture controls **information flow** — who attends to whom.

Objective controls **what it learns to do** — predict, mask, or reconstruct.

Finetuning / prompting controls **task adaptation** — how we specialize it.

mask + objective + head $\Rightarrow$ model family

Same block — three families, by **mask**, **objective** and **head**.

Next: **Vision Transformers & multimodal** — images as patches, ViT, and CLIP.